


Ch3-6

MultiTreading


3.11 in edition 6



Pipeline Hazards

LW r1, 0(r2)
LW r5, 12(r1)
ADDI r5, r5, #12
SW 12(r1), r5

- Each instruction may depend on the next
 - Without bypassing, need interlocks



- Bypassing cannot completely eliminate interlocks or delay slots

2



Multithreaded Software

□ Process

- Each process has its unique address space
- Can consist of several threads

□ Thread – each thread has its unique execution context:

- Its own PC + registers + stack
- All threads within a process share same address space
- Private heap is optional


□ Multithreaded app's: process broken into threads

- Increase concurrency
- Partial blocking
- Centralized (smarter) resource management (by process)

Multithreading

- How can we guarantee no dependencies between instructions in a pipeline?
 - One way is to interleave execution of instructions from different program threads on same pipeline

Interleave 4 threads, **T1-T4**, on non-bypassed 5-stage pipe



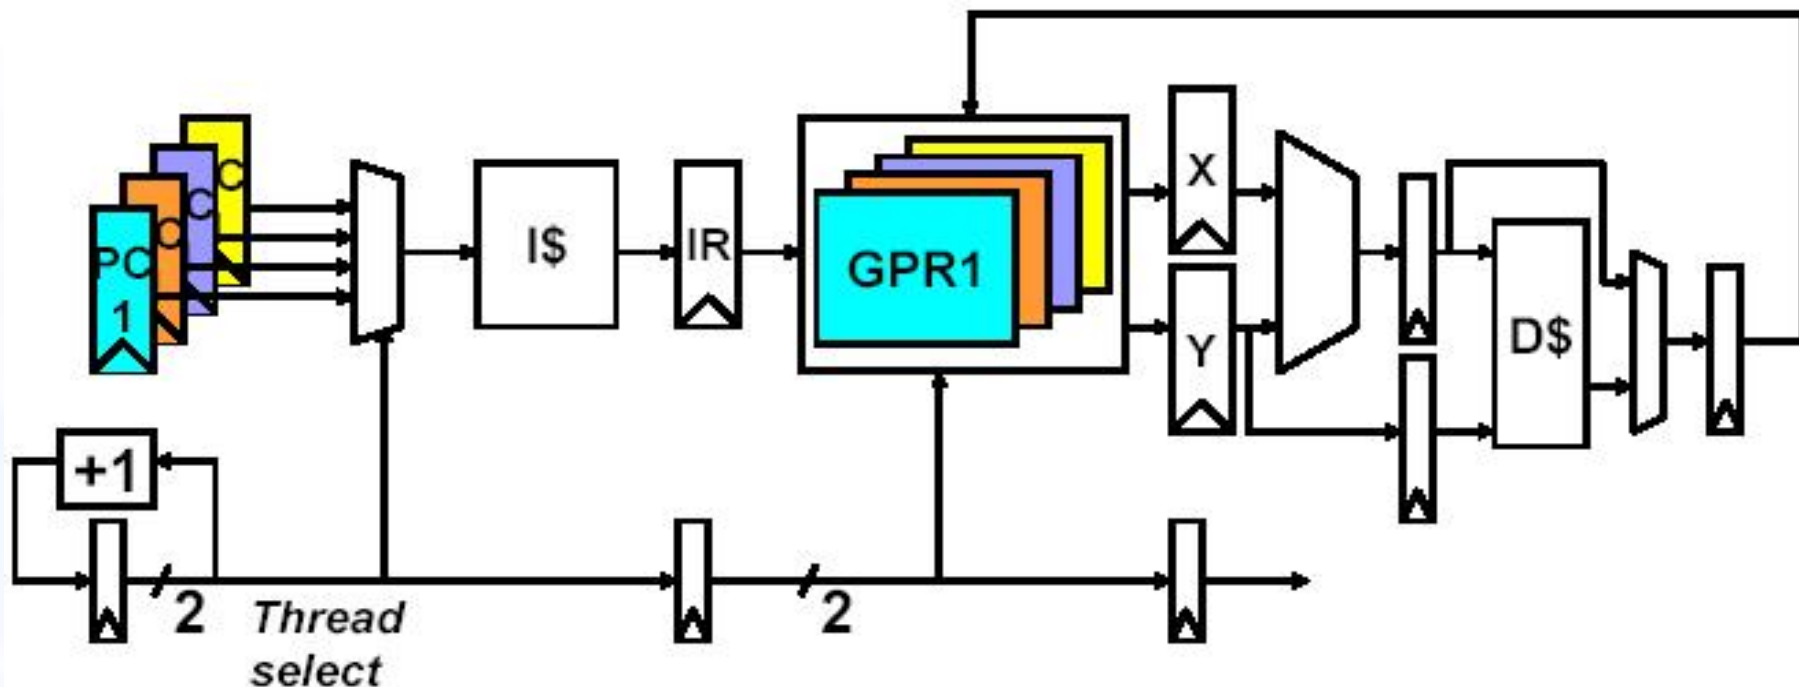
3



Multithreaded Architecture

- ❑ Processor capable of executing multiple software threads
 - Can execute “simultaneously”
 - Thread can be HW switched without OS control
 - Shared resource
 - Sharing -> better resource utilization -> better throughput
 - Can belong to the same process – but not have to !

Simple Multithreaded Pipeline



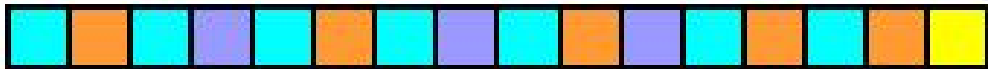
- Have to carry thread select down pipeline to ensure correct state bits read/written at each pipe stage

Multithreading Costs

- Appears to software (including OS) as multiple slower CPUs
- Each thread requires its own user state
 - GPRs
 - PC
- Also, needs own OS control state
 - virtual memory page table base register
 - exception handling registers
- *Other costs?*



Thread scheduling policies

- Fixed interleave (*CDC 6600 PPU's, 1965*)
 - each of N threads executes one instruction every N cycles
 - if thread not ready to go in its slot, insert pipeline bubble
 - Software-controlled interleave (*TI ASC PPU's, 1971*)
 - OS allocates S pipeline slots amongst N threads
 - hardware performs fixed interleave over S slots, executing whichever thread is in that slot
- 
- Hardware-controlled thread scheduling (*HEP, 1982*)
 - hardware keeps track of which threads are ready to go
 - picks next thread to execute based on hardware priority scheme



Denelcor HEP

(Burton Smith, 1982)

The Denelcor HEP was a uniform shared memory multiprocessor that used fine-grain multithreading to tolerate memory latency, synchronization latency, and even functional unit latency.

First commercial machine to use hardware threading in main CPU

- 120 threads per processor
- 10 MHz clock rate
- Up to 8 processors
- precursor to Tera MTA (Multithreaded Architecture)



Tera MTA (1990-97)

- Up to 256 processors
- Up to 128 active threads per processor
- Processors and memory modules populate a sparse 3D torus interconnection fabric
- Flat, shared main memory
 - No data cache
 - Sustains one main memory access per cycle per processor
- GaAs logic in prototype, 1KW/processor @ 260MHz
 - CMOS version, MTA-2, 50W/processor



MTA Architecture

- Each processor supports 128 active hardware threads
 - $1 \times 128 = 128$ stream status word (SSW) registers,
 - $8 \times 128 = 1024$ branch-target registers,
 - $32 \times 128 = 4096$ general-purpose registers
- Three operations packed into 64-bit instruction (short VLIW)
 - One memory operation,
 - One arithmetic operation, plus
 - One arithmetic or branch operation
- Thread creation and termination instructions
- Explicit 3-bit “lookahead” field in instruction gives number of subsequent instructions (0-7) that are independent of this one
 - c.f. instruction grouping in VLIW
 - allows fewer threads to fill machine pipeline
 - used for variable-sized branch delay slots



Coarse-Grain Multithreading

Tera MTA designed for supercomputing applications with large data sets and low locality

- No data cache
- Many parallel threads needed to hide large memory latency

Other applications are more cache friendly

- Few pipeline bubbles when cache getting hits
- Just add a few threads to hide occasional cache miss latencies
- Swap threads on cache misses



IBM Power RS64-IV (2000)

- Commercial coarse-grain multithreading CPU
- Based on PowerPC with quad-issue in-order five-stage pipeline
- Each physical CPU supports two virtual CPUs
- On L2 cache miss, pipeline is flushed and execution switches to second thread
 - short pipeline minimizes flush penalty (4 cycles), small compared to memory access latency
 - flush pipeline to simplify exception handling




Multithreading design choices

- Fine-grained multithreading
 - Context switch among threads every cycle
- Coarse-grained multithreading
 - Context switch among threads every few cycles, e.g., on:
 - » Function unit data hazard,
 - » L1 miss,
 - » L2 miss...
- Why choose one style over another?
- Choice depends on
 - Context-switch overhead
 - Number of threads supported (due to per-thread state)
 - Expected application-level parallelism...



Superscalar Machine Efficiency




- Why horizontal waste?
- Why vertical waste?

Cache misses




Vertical Multithreading



- What is the effect of cycle-by-cycle interleaving?
 - removes vertical waste, but leaves some horizontal waste



Chip Multiprocessing




- What is the effect of splitting into multiple processors?
 - eliminates horizontal waste,
 - leaves some vertical waste, and
 - caps peak throughput of each thread.



Ideal Superscalar Multithreading

[Tullsen, Eggers, Levy, UW, 1995]



- Interleave multiple threads to multiple issue slots with no restrictions



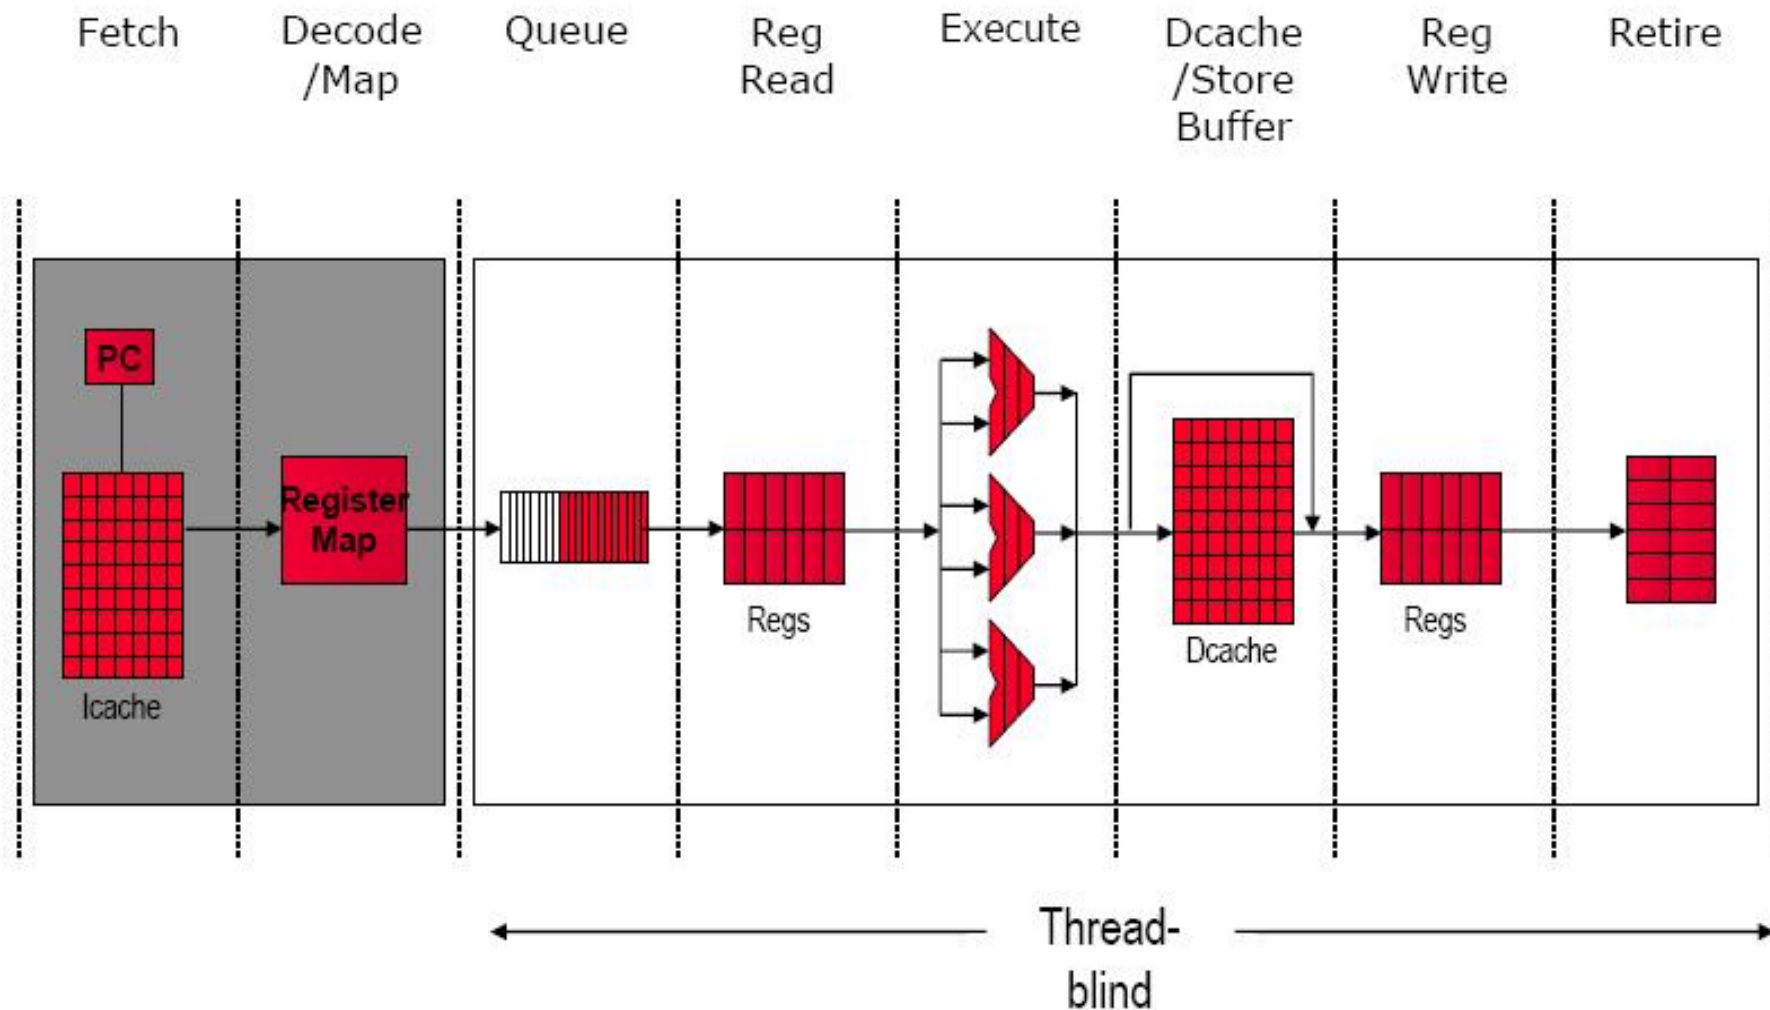
O-o-O Simultaneous Multithreading

[Tullsen, Eggers, Emer, Levy, Stamm, Lo, DEC/UW, 1996]

- Add multiple contexts and fetch engines and allow instructions fetched from different threads to issue simultaneously
- Utilize wide out-of-order superscalar processor issue queue to find instructions to issue from multiple threads
- OOO instruction window already has most of the circuitry required to schedule from multiple threads
- Any single thread can utilize whole machine

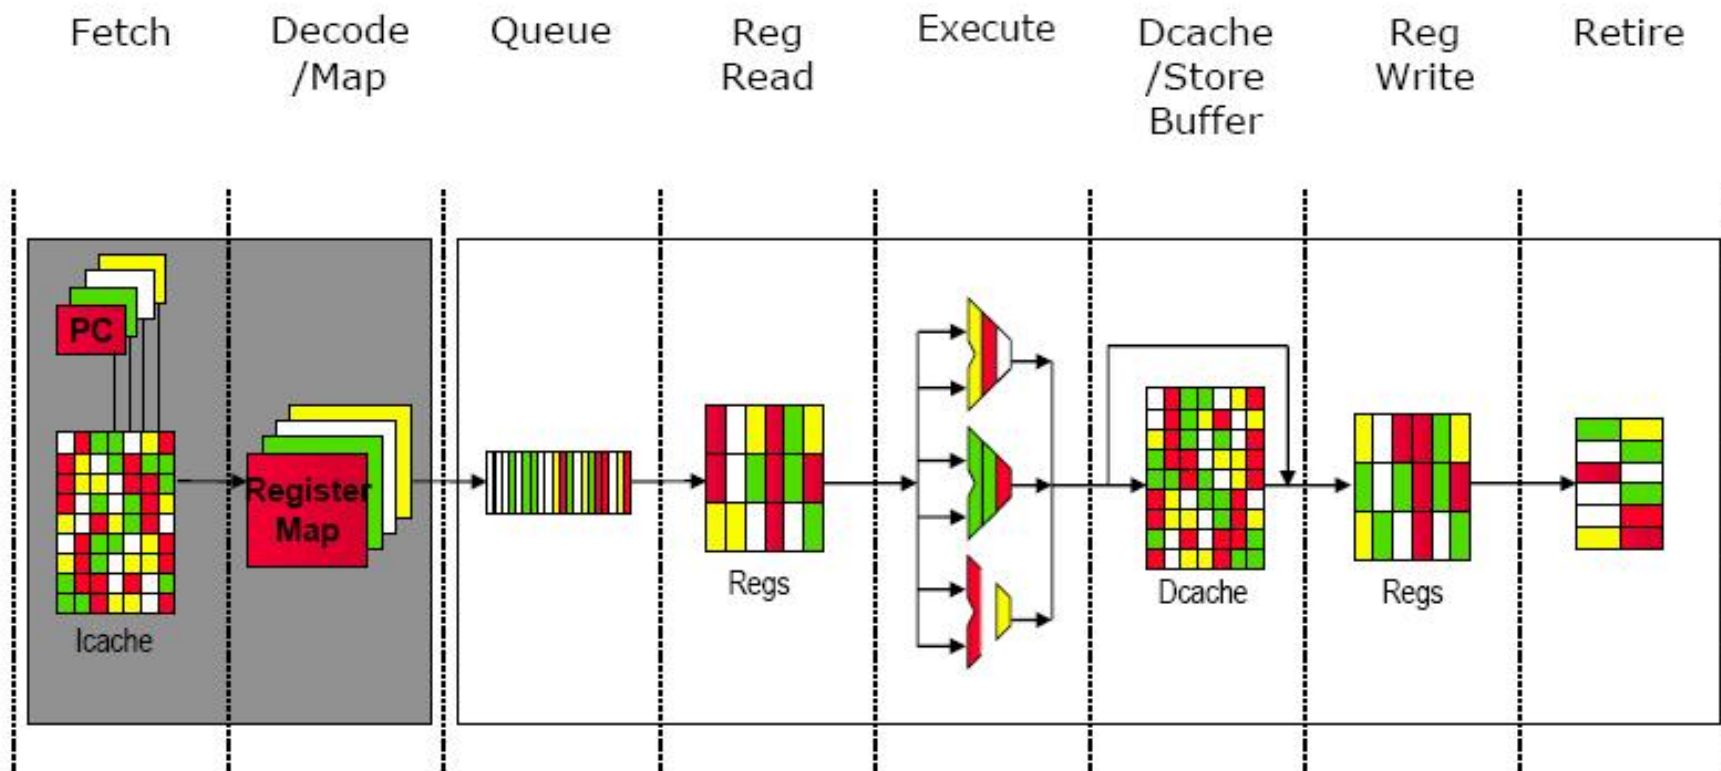


Basic Out-of-order Pipeline





SMT Pipeline





Design Challenges in SMT

- ❑ Trade-off between fine-grained implementation performance and single-thread performance.
 - preferred thread : might sacrifices throughput
 - Less likely to have a mix of instructions from several threads.
 - Maximize single-thread performance, should fetch as far ahead as possible, and have the fetch unit free when a branch is mispredicted or miss occur in prefetch buffer.



Design Challenges in SMT

- ❑ A larger register file to hold multiple context
- ❑ Not affecting the clock cycle, such as in instruction issue, in instruction completion.
- ❑ Ensuring that Cache and TLB conflicts do not cause performance degradation.
- ❑



Power 5 with SMT support vs Power 4

- ❑ Increasing the associativity of L1 icache and iTLB
- ❑ Adding per-thread load and store queues
- ❑ Increasing the size of L2 and L3
- ❑ Adding separate instruction prefetch and buffering
- ❑ Increasing virtual registers from 152 to 240
- ❑ Increasing size of several issue queues



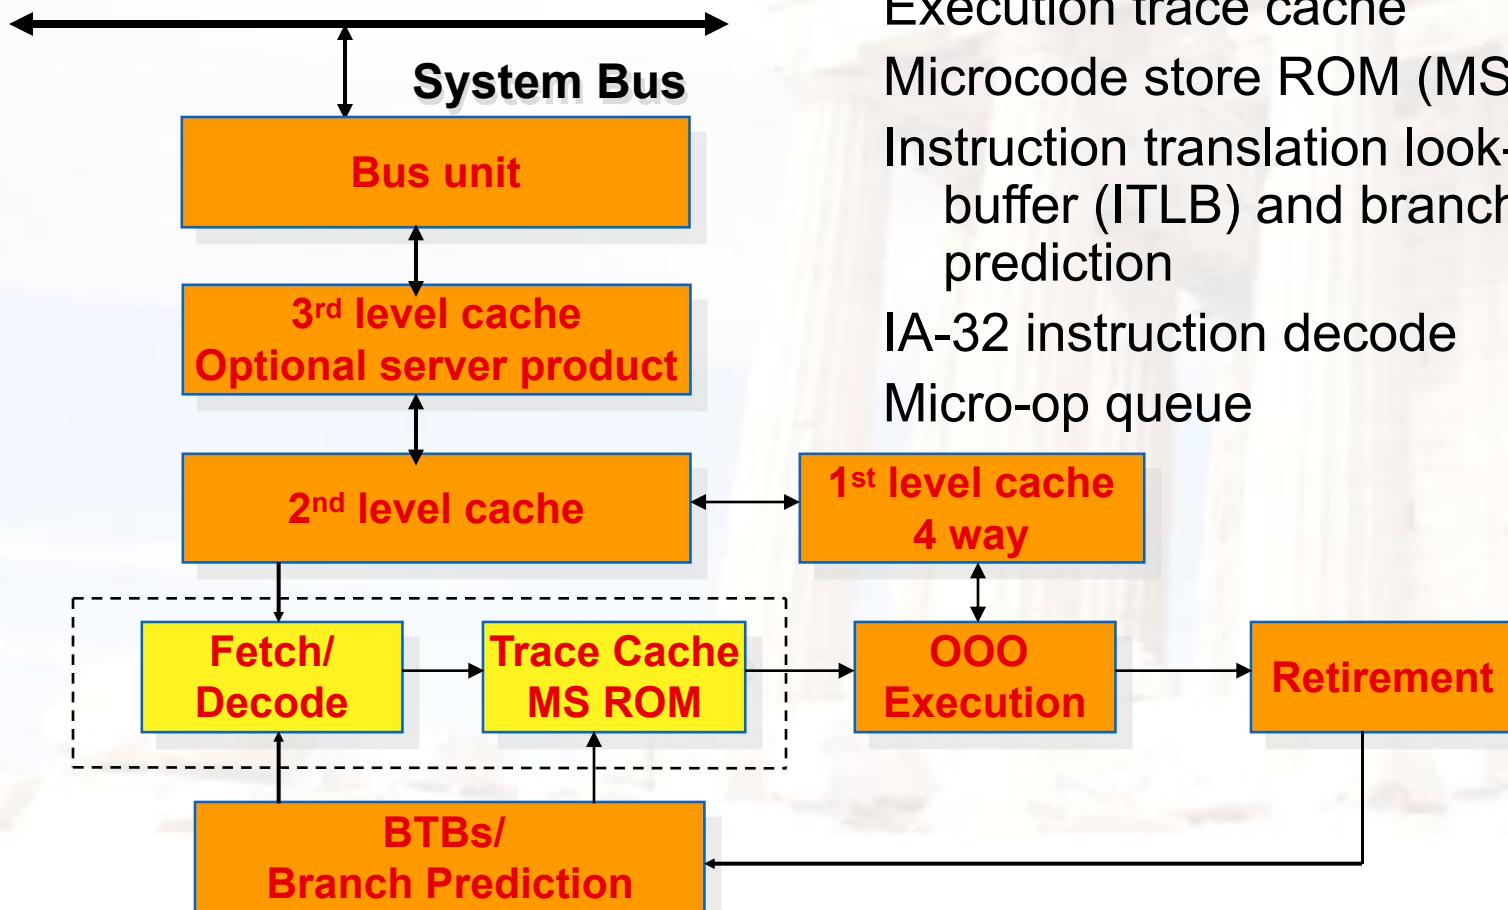
Pentium-4 Hyperthreading

- First commercial SMT design (2-way SMT)
 - Hyperthreading == SMT
- Logical processors share nearly all resources of the physical processor
 - Caches, execution units, branch predictors
- Die area overhead of hyperthreading $\sim 5\%$
- When one logical processor is stalled, the other can make progress
 - No logical processor can use all entries in queues when two threads are active
- Processor running only one active software thread runs at approximately same speed with or without hyperthreading



HT Technology Architecture

Front End






Front End

- ❑ Responsible for delivering instruction to the later pipe stages
- ❑ Trace cache hit
 - When the requested instruction trace is present in trace cache
- ❑ Trace cache miss
 - Requested instruction is brought in the trace cache from L2 cache




Trace Cache (TC) Hit Front End



- ❑ Two sets of instruction pointers
- ❑ Two logical processors arbitrates access to TC every clock cycle
- ❑ If one logical processor stalled, then other logical processor can use the full bandwidth of TC



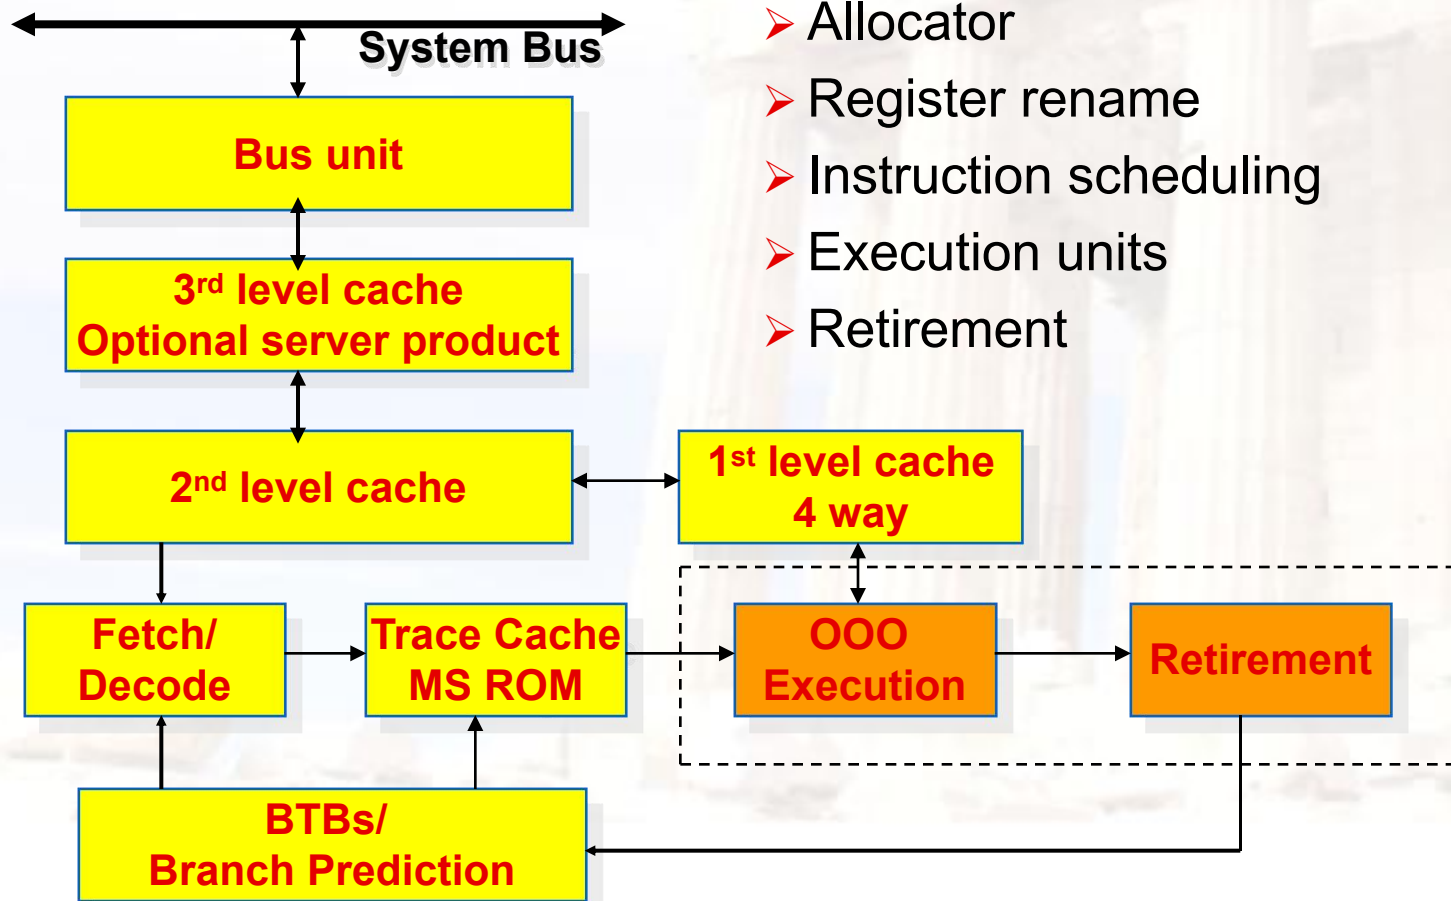
Trace Cache Miss Front End





HT Technology Architecture

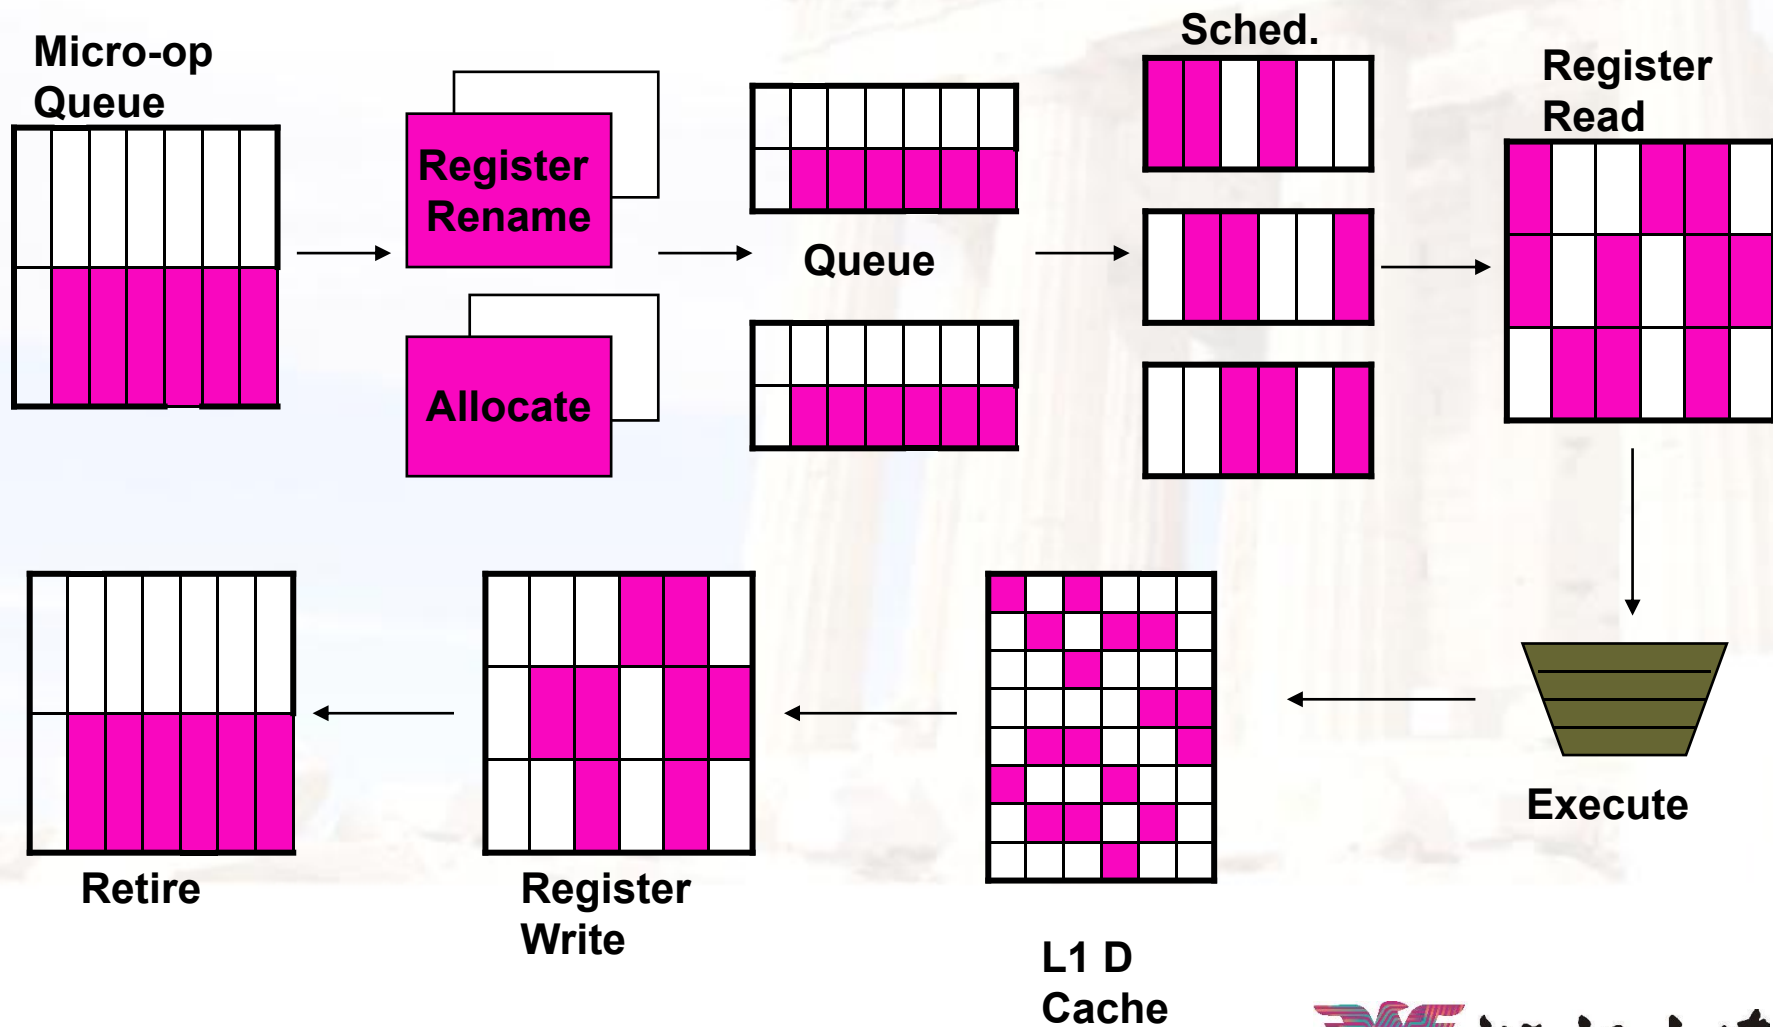
Back End



- Allocator
- Register rename
- Instruction scheduling
- Execution units
- Retirement

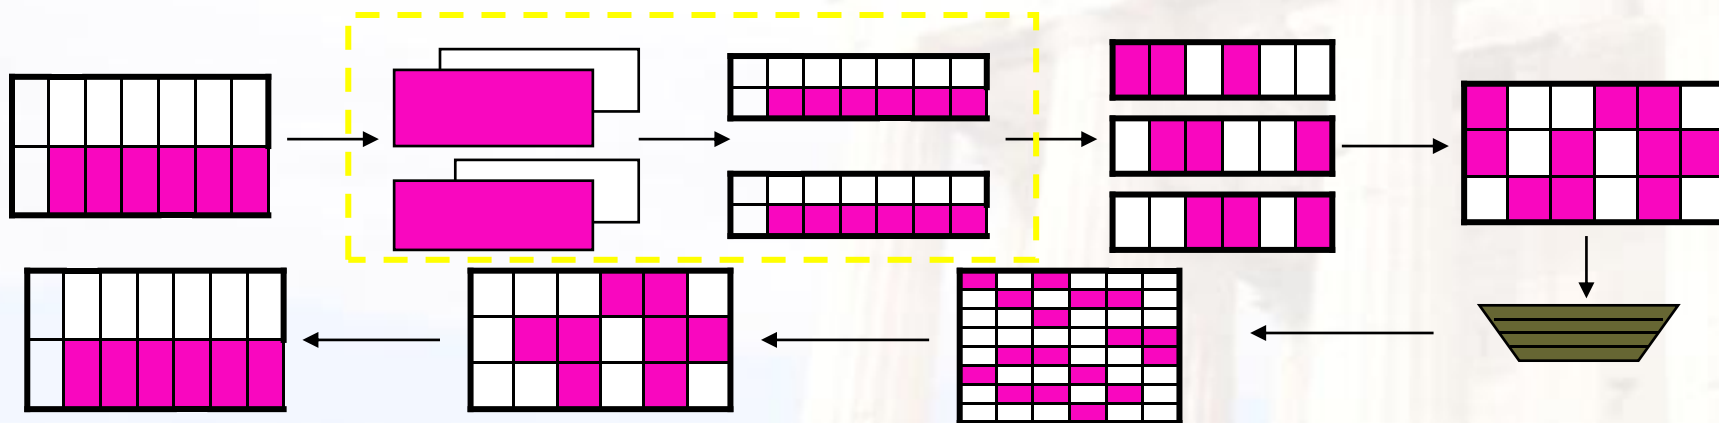


Detailed Pipeline





Allocator Back End



Allocates many of the key machine buffers including

- ❑ 126 re-order buffer entries
- ❑ 128 integer and floating point physical registers
- ❑ 48 load 24 store buffers
- ❑ When HT Technology is enabled, each logical processor can use maximum of 63 re-order buffer entries, 24 load buffers and 12 store buffers

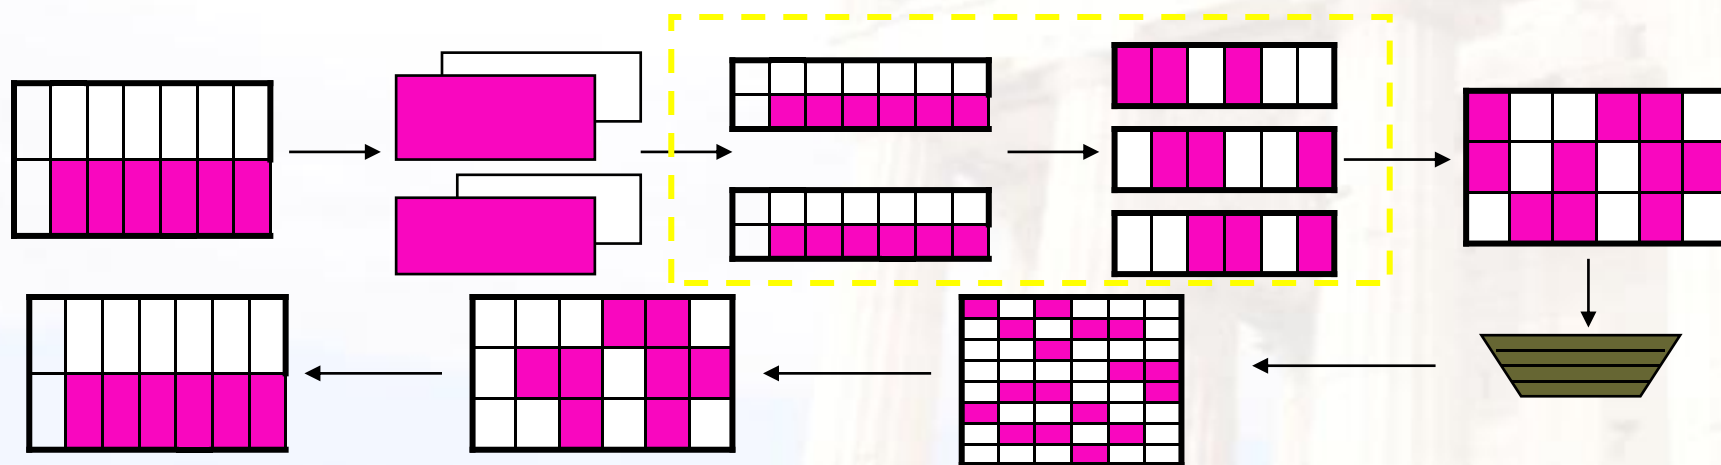


Register Rename Back End

- ❑ There are two Register Alias Tables (RAT) for each processor
- ❑ Rename done in parallel with allocator logic
- ❑ Once done, micro-ops are placed into two sets of queues: MIAQ(memory instruction address queue) and GIAQ(general instruction address queue)
- ❑ Queues are partitioned such that micro-ops from each logical processor can use at most half the entries



Instruction Scheduling Back End




- ❑ Five schedulers are used to schedule different types of micro-ops
- ❑ Both MIAQ and GIAQ send micro-op to five schedulers as fast as they can
- ❑ Each scheduler has its own scheduler queue
- ❑ Micro-ops are simply evaluated based on dependent inputs and availability of execution resources
- ❑ To avoid deadlock, there is a limit on number of active entries that a logical processor can have in scheduler's queue



Execution units


Back End



- ❑ After execution, micro-ops are placed in re-order buffer. It decouples execution stage from retirement stage
- ❑ Re-order buffer (ROB) is partitioned such that each logical processor can use half entries



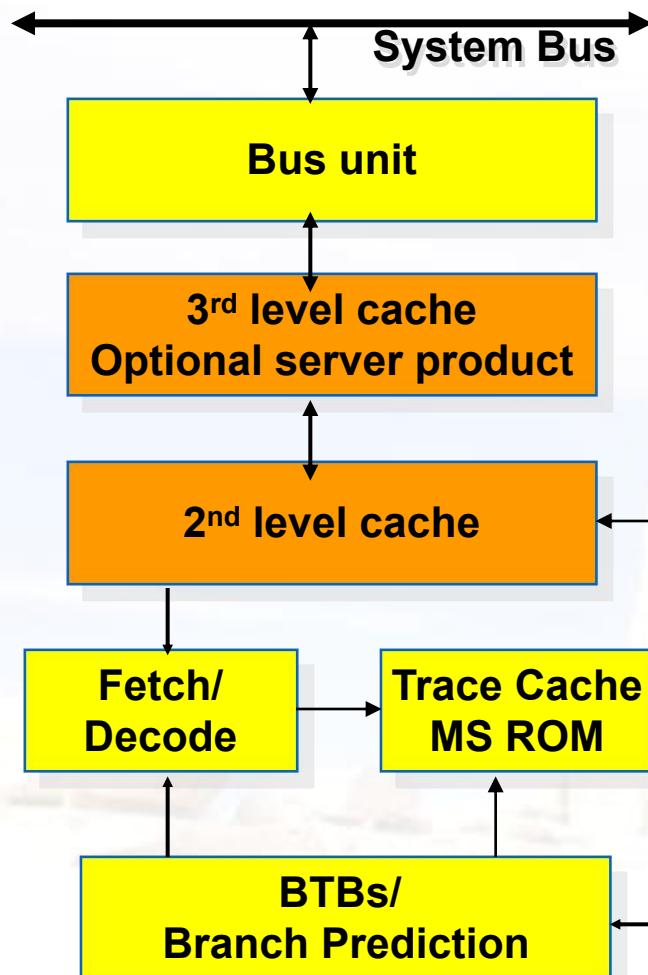
Retirement Back End



- ❑ Tracks when micro-ops from two logical processors are ready to be retired, then retires the micro-ops by alternating between two logical processors



HT Technology Architecture



Memory Subsystem

- DTLB
- L1 data cache
- L2 unified cache
- L3 unified cache



DTLB

Memory Subsystem

- ❑ Tagged resource
- ❑ Each entry in dual translation lookaside buffer (DTLB) includes a logical processor ID tag
- ❑ Each processor has reservation register to ensure fairness and forward progress in processing DTLB misses



Memory Subsystem

L1 Data Cache, L2 Cache, L3 Cache

- ❑ Both logical processors can share all entries in all three levels of cache



Bus

- ❑ Logical processors are treated on first-come first-serve basis. Priority is not given to any logical processor
- ❑ Distinction between requests from two logical processors are maintained
- ❑ Requests to APIC and interrupt delivery resources are unique and separate per logical processor



Sun T1 Multiprocessor


- ❑ Focus on Thread Level Parallelism
- ❑ Commercial application oriented
- ❑ Small single-issue core employ multithreading
- ❑ Multi core on chip



Single Thread Design

- ❑ L2 cache is getting larger with less performance contribution
- ❑ Power efficiency decrease
- ❑ Hot-spot
- ❑ Lower processor utilization

Single-Core Processor






Single Thread design multiple issue

Up to 85% Cycles Spent Waiting for Memory

Single Threaded Performance




Typical Utilization of Processor: 15–25%





Hardware Multi Threading

Utilization: Up to 85%*




* based on example of UltraSPARC T1

Multi-threaded
Performance






Chip MultiThreading




CMP
(Chip MultiProcessing,
a.k.a. “multicore”)

n cores per processor



HMT
(Hardware
Multithreading)

m threads per core




CMT
(Chip
MultiThreading)

$n \times m$ threads per processor



Why CMT works ?

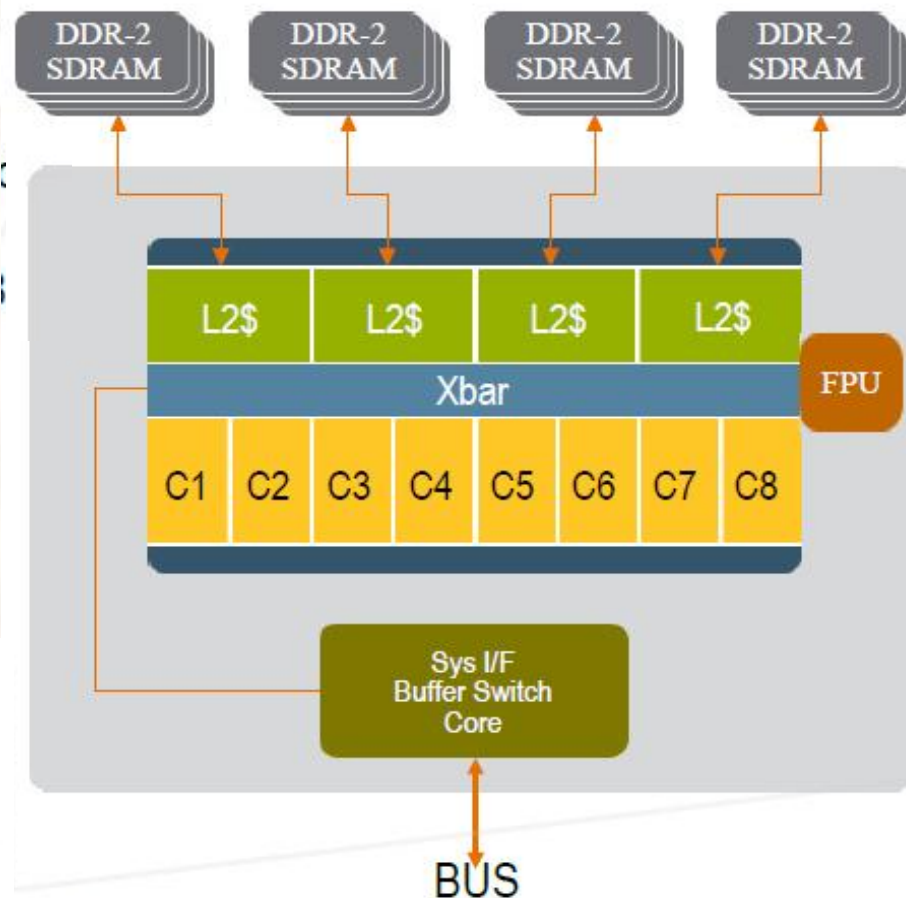
Goal: “100% Resource Utilization”
(given a fixed die size)






T1 Architecture

- ❑ 8 core/chip
- ❑ 4 Treads/core
- ❑ 1 FPU shared for 8 cores
- ❑ Core connected via crossbar switch(134.4GB/s)
- ❑ High-bandwidth, 12-way associative 3MB L2 cache
- ❑ 4 DDR2 channels (23GB/s)
- ❑ Power < 80W
- ❑ `300M transistors
- ❑ 378 mm² die




1 core





Chip Layout

- ❑ L1 cache: 16KB Icache, 32B line; Dcache, 8KB, 16B line; 4-way set associative, 64B block, L1 miss penalty=23 cycles,
- ❑ L2 cache: 4 banks, each 750KB, 64B block, L2 miss penalty = 110 cycles
- ❑ L1 cache coherence is enforced by directory associate with each L2 cache block
- ❑ Memory controller on chip to decrease miss penalty






Compare with SMP

32-thread Traditional SMP System


Example: Typical SMP Machine Configuration



BTU: british thermal Unit

32-thread OpenSPARC T1 Processor

One motherboard, *no* switch ASICs




Direct crossbar interconnect

- Lower cost
- better RAS
- lower BTUs,
- lower and uniform latency,
- greater and uniform bandwidth. . .



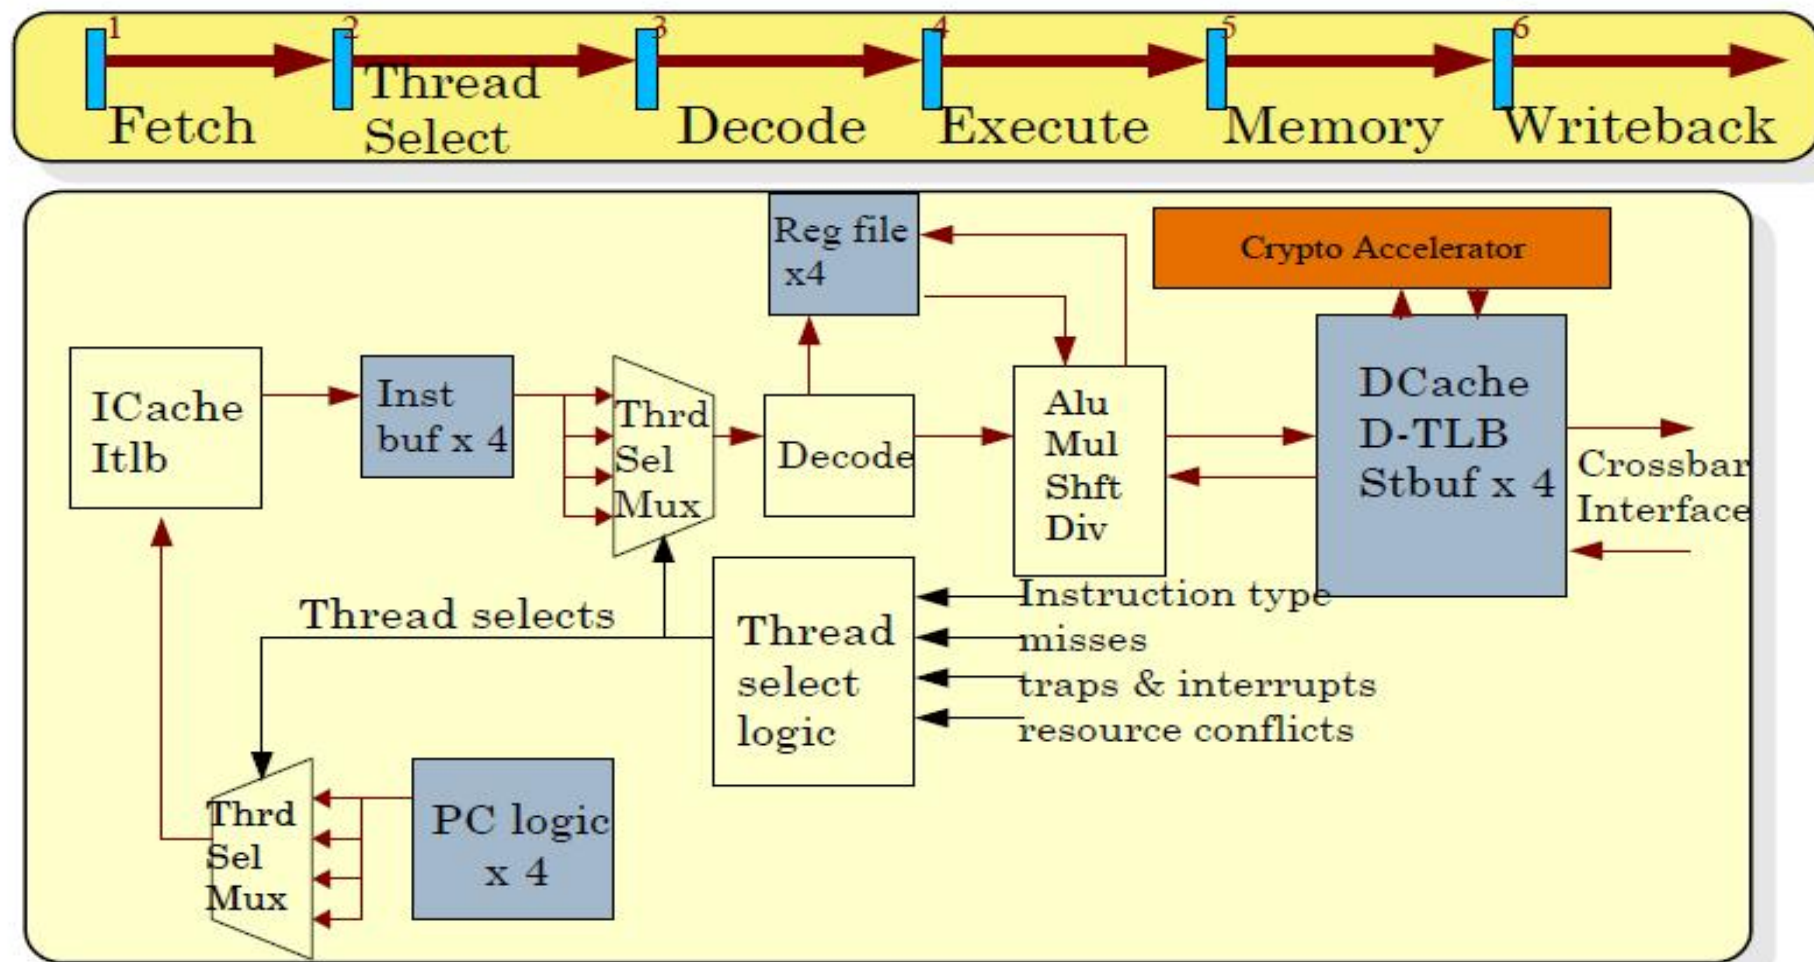
Core structure

- ❑ 4 threads /core
- ❑ Single issue
- ❑ 6 stage pipeline
- ❑ Unique resources/thread
 - Registers
 - Portions of I-fetch datapath
 - Store and Miss buffers
- ❑ Resource shared by 4 threads
 - Caches, TLBs, Exe Units
 - Pipeline registers and Data path
- ❑ Core area = 11mm² in 90nm
- ❑ MT adds ~20% area to core





T1 processor core pipeline





Thread Selection Policy

- ❑ Switch among available (ready to run) threads every cycle
 - Priority given to least-recently-executed thread

- ❑ Thread becomes not-ready-to-run due to:
 - Long latency operations like load, branch, mul, or div (load, branch incur 1 3-cycle delay)
 - Pipeline stall when cache miss, trap, or structure hazard



Instruction Fetch/Switch/Decode Unit (IFU)

❑ I-cache

- 16KB data, 4ways, 32B line
- Single-ported Instruction Tag
- Dual-ported(1R/1W) Valid-bit array
- Invalidate operations access Valid-bit array, not instruction Tags
- Pseudo-random replacement

❑ Fully associative Instruction TLB

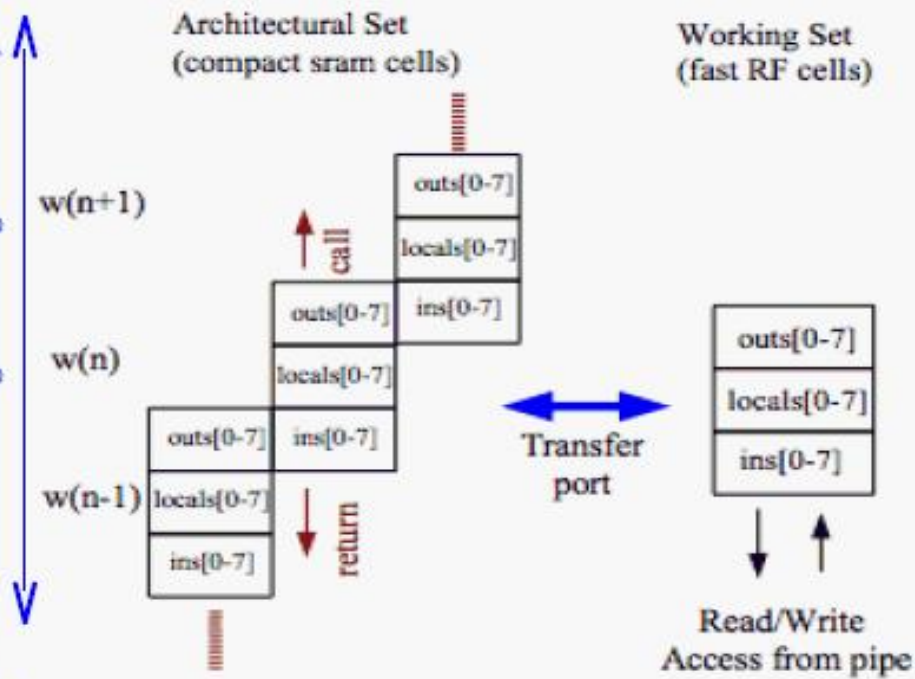
- 64 entries, Page sizes: 8K, 64K, 4M, 256M
- Multiple hits in TLB prevented by doing auto-demap on fill

❑ 2 instructions fetched each cycle



Windowed Int Register file

(16 reg x 8 windows + 8 global regs x 4 sets) x 4 threads



- ❑ 5KB 3R/2W/1T structure
 - >640 64b regs
- ❑ Only the 32 registers from current window is visible to thread
- ❑ 8 global + 24 window register
- ❑ Window changing in background under thread switch
- ❑ Single cycle R/W accwss



Execution Units

- ❑ Single ALU and shifter. ALU reused for branch address and virtual address calculation.
- ❑ Integer Multiplier
 - 5 clock latency, throughput of $\frac{1}{2}$ per cycle
 - Constrains accumulate function for Mod Arithmetic
 - 1 integer mul allowed outstanding per core
 - Multiplier shared between Core Pipe and Modular Arithmetic unit on a round robin basis
- ❑ Simple divider, with one divide outstanding per core
- ❑ Thread issuing a MUL/DIV will rollback and switch out if mul/div units is occupied.



Load Store Unit (LSU)

❑ D-Cache complex

- 8KB data, 4 ways, 16B line
- Single ported Data Tag
- Dual ported(1R/1W) Valid bit array
- Pseudo-random replacement
- Load are allocating, stores are non allocating

❑ 8 entry store buffer per thread, unified into single 32 entry array, with RAW bypassing



LSU(continued)

❑ Crossbar interface

- LSU priorities requests to the crossbar for FPOps, streaming ops, I and D misses, stores and interrupts.

❑ Request priority

- Imiss>ldmiss>stores, {fpu, strm, interrupt}
- Packed assembly for pcx.

❑ Handles return from crossbar and maintains order for cache updates and invalidates.

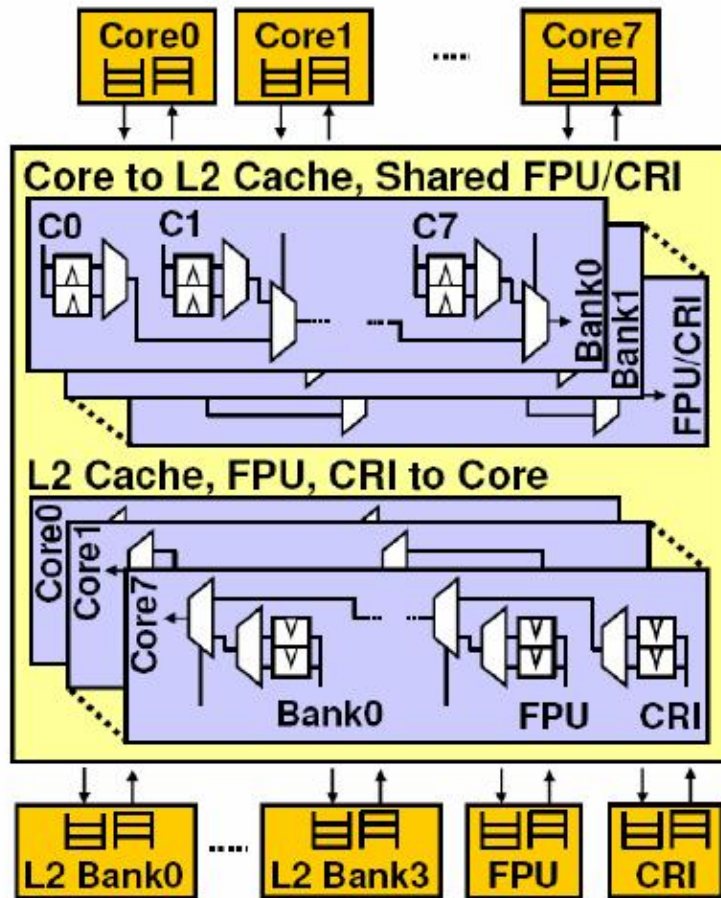


Other Functions

- ❑ Support for up to 64 pending interrupts per thread
- ❑ Support for 6 trap levels. Traps cause pipeline flush.
- ❑ Floating Point
 - The thread switch out when detecting ans FPop
 - Fpop is further decoded and FRF is read
 - FPop with operands are packetized and shipped over the crossbar to the FPU
 - FPU returned the result via crossbar
 - Write back to FRF and thread restart



Cross Bar



- ☐ Each requestor queues up to 2 packets per destination
- ☐ 3 stage pipeline: request, arbitrate and transmit
- ☐ Oldest requestor first
- ☐ Core to cache bus;
 - Address+double word store
- ☐ Cache to core bus
 - 16B line 32B lcache line fill delivered in 2 back to back clock cycles



L2 Cache

- ❑ 3MB, 4-way banked, 12-way associative, writeback
- ❑ 64B line size, 64B interleaved between banks
- ❑ Latency: 8 clocks for load, 9 clocks for I-miss, critical chunk first
- ❑ 16 outstanding misses per bank → 64 total
- ❑ Coherence maintained by shadowing L1 tags in an L2 directory structure
- ❑ L2 is global visibility. DMA from IO is serialised write traffic from cores in L2
- ❑ L2 has a 8-stage pipeline (C1-C8)



L2 Cache --Directory

□ Directory shadows L1 tags

- 2048 entries, half-half for Icache and Dcache
- L1 set index and L2 bank interleaving is such that $\frac{1}{4}$ of L1 entries come from each L2 bank
- On an L1 miss, the L1 replacement way and set index identify the physical location of tag
- On a store, write invalidate will be done
 - Invalidate operation is pointer to the physical location of L1. (no tag lookup in L1)
 - An acknowledgement packet is sent to the request CPU, an invalidation packet sent to all other core



L2 Cache Single Bank

- ❑ Arbiter (instruction from CCX, DMA instruction, stall signals from pipeline, instruction from FB and MB)
- ❑ L2 tag
- ❑ L2 VUAD array –Valid, Used, Allocated(set when a line is picked for replacement) Dirty
- ❑ L2 data
- ❑ Input Queue (16 entry FIFO) – packets arriving on the PCX
- ❑ Output Queue (16 entry FIFO) – operations waiting for access to CPX(cache-to-processor interface)
- ❑ Snoop input Queue (2-entry FIFO)
- ❑ Miss buffer (MB) (L2 misses, same line address with previous miss or write-back buffer, atomics and partial stores etc)
- ❑ Fill buffre (temp store data from DRAM)
- ❑ Write-back buffer
- ❑ Remote DMA write buffer (4 entry) for DMA write
- ❑ L2 cache directory –coherency management and maintain inclusive property



Coherence/consistency

- ❑ Load update directory & fill the L1 on return
- ❑ Stores are non-allocating in L1
 - TSO, PSO, RMO (one TSO store outstanding to L2 per thread)
 - Check directory and determine L1 hit
 - Directory send store ack/inv to core
 - Store update happens to Dcache on store ack
- ❑ Crossbar orders responses across cache banks
- ❑ 3 atomic instructions are processed by L2 cache
 - LDSTUB, SWAP, Compare and swap(CAS)
 - Two passes down to cache pipeline



On chip Mem controller

- ❑ 4 independent DDRII DRAM channels
- ❑ Upto 128GB memory size
- ❑ 25GB/s peak bandwidth
- ❑ Schedules across 8 rds + 8 writes
- ❑ 128+16b interface, chipkill support, correction, byte error detection
- ❑ Designed to work from 125-200Mhz



T1 Performance




Figure 4.34 Performance efficiency on SPECRate for four dual-core processors, normalized to the Pentium D metric (which is always 1).



Conclusion

- ❑ Sun T1 provides much better performance in terms of performance/watt, especially for TPC-C-like and SPECJBB05 benchmark.
- ❑ TLP approach is much more power efficient than an ILP-intensive approach for multithreaded applications.
- ❑ CMT may provide a way to increase performance in a power efficient fashion.