

# Chapter 4: Threads

# Chapter 4: Threads

- Overview
- Multithreading Models
- Threading Issues
- Pthreads
- Windows XP Threads
- Linux Threads
- Java Threads

## Suppose you are developing a Web browser

```
main() {  
    while(1) {  
        RetrieveData();    // Block for 1 second  
        DisplayData();     // Block for 1 second  
        GetInputEvents();  // Block for 1 second  
    }  
}
```

Now what if you want the program to be more responsive?

```
main() {  
    while(1) {  
        RetrieveALittleData();    // Block for 0.1 second  
        DisplayALittleData();     // Block for 0.1 second  
        GetAFewInputEvents();     // Block for 0.1 second  
    }  
}
```


# What if you want it to be even more responsive?

```
main() {  
    while(1) {  
        if (CheckData () == True) {  
            RetrieveALittleData ();    // 0.1 second  
            DisplayALittleData ();    // 0.1 second  
        }  
        if (CheckInputEvents () == True) {  
            GetAFewInputEvents ();    // 0.1 second  
        }  
    }  
}
```

Problem: A lot checks, not efficient. But still not responsive!

# To make it responsive enough, we need to

- Break the operations into very very small pieces;



These two requirements  
conflict with each other!

- However, to be efficient enough, we want to execute code in large pieces.
- More precisely, we want to **SCHEDULE** these operations in our own program code.
- Leave the tedious work to the OS which schedules them in *Threads*!

# Multi-thread version of the program

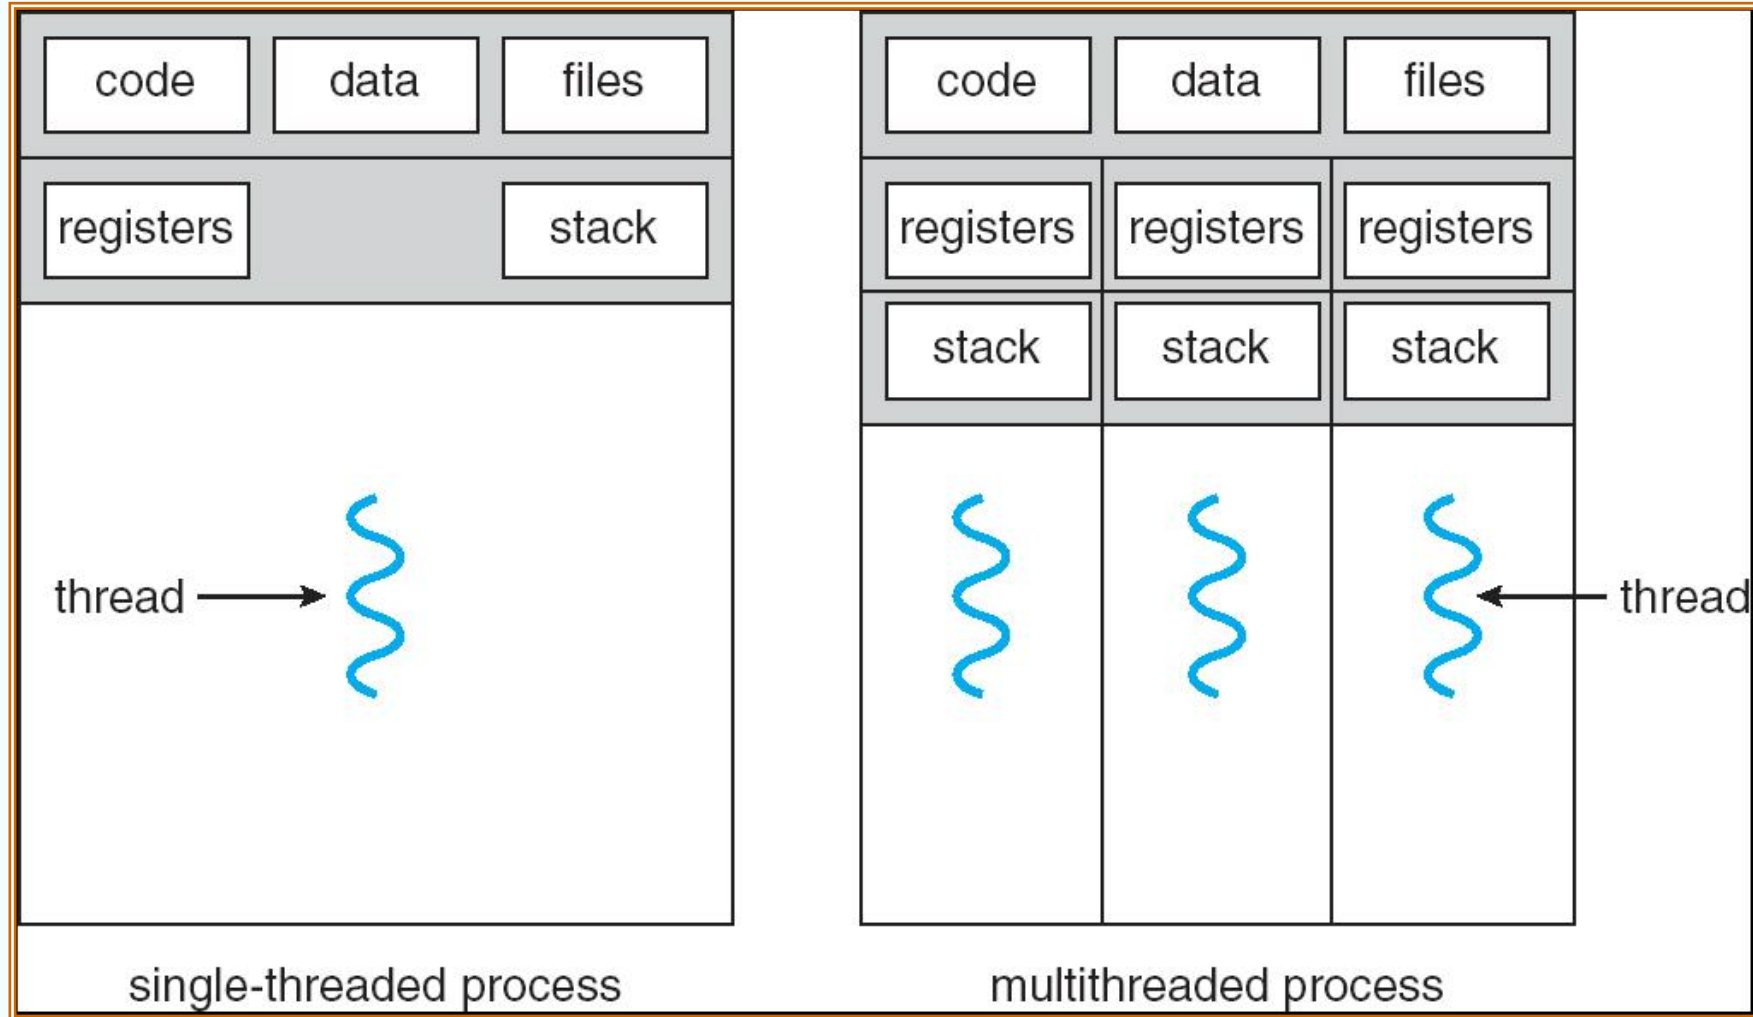
```
main() {  
    CreateThread (RetrieveData());  
    CreateThread (DisplayData());  
    CreateThread (GetInputEvents());  
    WaitForThreads();  
}
```

// Each thread routine enters a loop.

```
void RetrieveData() {  
    while(1) {  
        retrieveData();  
        ...  
    }  
}  
....
```

```
void DisplayData() {  
    while(1) {  
        displayData();  
        .....  
    }  
}
```

# Single and Multithreaded Processes




# Benefits


- **Responsiveness**  
interactive applications
- **Resource Sharing**  
memory for code and data can be shared.
- **Economy**  
creating processes are more expensive.
- **Utilization of MP Architectures**  
multi-threading increases concurrency.



# Concurrency vs. Parallelism



**Figure 4.3** Concurrent execution on a single-core system.



**Figure 4.4** Parallel execution on a multicore system.

# User Threads

- Thread management done by user-level threads library
- Three primary thread libraries:
  - POSIX Pthreads (can also be provided as system library)
  - Win32 threads
  - Java threads

# Kernel Threads

- Supported by the Kernel
- Almost all contemporary OS implements kernel threads. Examples
  - Windows XP/2000
  - Solaris
  - Linux
  - MacOS

# Multithreading Models


- **Many-to-One**  
thread management is efficient, but will block if making system call, kernel can schedule only one thread at a time
- **One-to-One**  
more concurrency, but creating thread is expensive
- **Many-to-Many**  
flexible

# Many-to-One

- Many user-level threads mapped to single kernel thread
- The scheduling is done completely by the thread library and the kernel itself is not aware of the multiple threads in user-space.
- Examples:
  - Solaris Green Threads
  - GNU Portable Threads


# Many-to-One Model

线程的分配，管理，  
调度在线程库的层次  
实现，实际上只对应  
一个kernel thread  
(local scheduling)




# One-to-One

- Each user-level thread maps to kernel thread
- Examples
  - Windows NT/XP/2000
  - Linux
  - Solaris 9 and later



线程的分配, 管理, 调度在 kernel 的层次实现, 由操作系统全权负责(global scheduling)

# One-to-one Model






# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- A program can have as many threads as are appropriate without making the process too heavy or burdensome. In this model, a user-level threads library provides sophisticated scheduling of user-level threads above kernel threads.
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package
- Go routines


# Many-to-Many Model



# Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to a kernel thread
- Examples
  - IRIX
  - HP-UX
  - Tru64 UNIX
  - Solaris 8 and earlier

# Two-level Model



# Threading Issues

- Semantics of **fork()** and **exec()** system calls
- Thread cancellation
- Signal handling
- Thread pools
- Thread-specific data
- Scheduler activations

# Semantics of `fork()` and `exec()`

- Does **`fork()`** duplicate only the calling thread or all threads?
- Some Unix systems have two versions of `fork()`, one that duplicates all threads and another that duplicates the thread that invokes `fork()`. It's not trivial though.
- `Exec()` will replace the entire process.

# Signal Handling

- Signals are used in UNIX systems to notify a process that a particular event has occurred
- A **signal handler** is used to process signals, either synchronous or asynchronous:
  1. Signal is generated by particular event
  2. Signal is delivered to a process
  3. Signal must be handled
- Options: (method of delivery depends on the type of signal)
  - Deliver the signal to the thread to which the signal applies
  - Deliver the signal to every thread in the process
  - Deliver the signal to certain threads in the process
  - Assign a specific thread to receive all signals for the process

# Thread Cancellation

- Terminating a thread before it has finished
- Two general approaches:
  - **Asynchronous cancellation** terminates the target thread immediately
  - **Deferred cancellation** allows the target thread to periodically check via a flag if it should be cancelled

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```



# Thread Pools

- Create a number of threads in a pool where they await work
- Advantages:
  - Usually slightly **faster** to service a request with an existing thread than create a new thread
  - Allows the number of threads in the application(s) to be bound to the size of the pool

# Thread Specific Data

- Thread-Local Storage (in 10<sup>th</sup> edition)
- Allows each thread to have its own copy of data
- In some ways similar to static data, but are unique to each thread
- Useful when you do not have control over the thread creation process (i.e., when using a thread pool)

# Scheduler Activations

- Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- LWP is a **virtual processor** attached to kernel thread
- Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the thread library
- Upcalls are handled by the thread library with an **upcall handler**
- This communication allows an application to **maintain the correct number** of kernel threads

*when an application thread is about to block, an upcall is triggered.*

# Pthreads

Q: Is it a user- or kernel-level library?

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems (Solaris, Linux, Mac OS X)

Can be either user- or kernel-level.

# Windows XP Threads

- Implements the one-to-one mapping
- Each thread contains
  - A thread id
  - Register set
  - Separate user and kernel stacks
  - Private data storage area
- The register set, stacks, and private storage area are known as the **context** of the threads
- The primary data structures of a thread include:
  - ETHREAD (executive thread block)
  - KTHREAD (kernel thread block)
  - TEB (thread environment block)

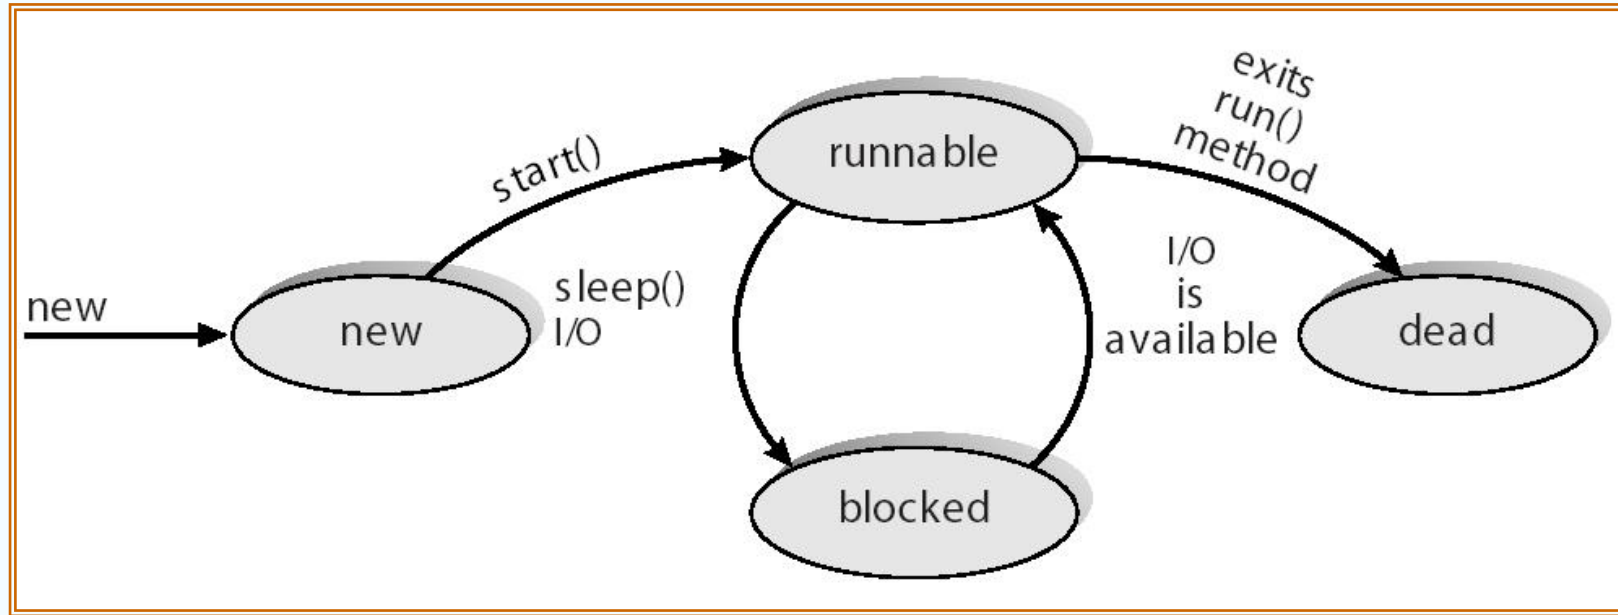
# Linux Threads

- Linux refers to them as *tasks* rather than *threads*
- Thread creation is done through **clone()** system call
- **clone()** allows a child task to share the address space of the parent task (process)

# Java Threads

- Java threads are managed by the JVM
- Java threads may be created by:
  - Extending Thread class
  - Implementing the Runnable interface

# Java Thread States





# End of Chapter 4