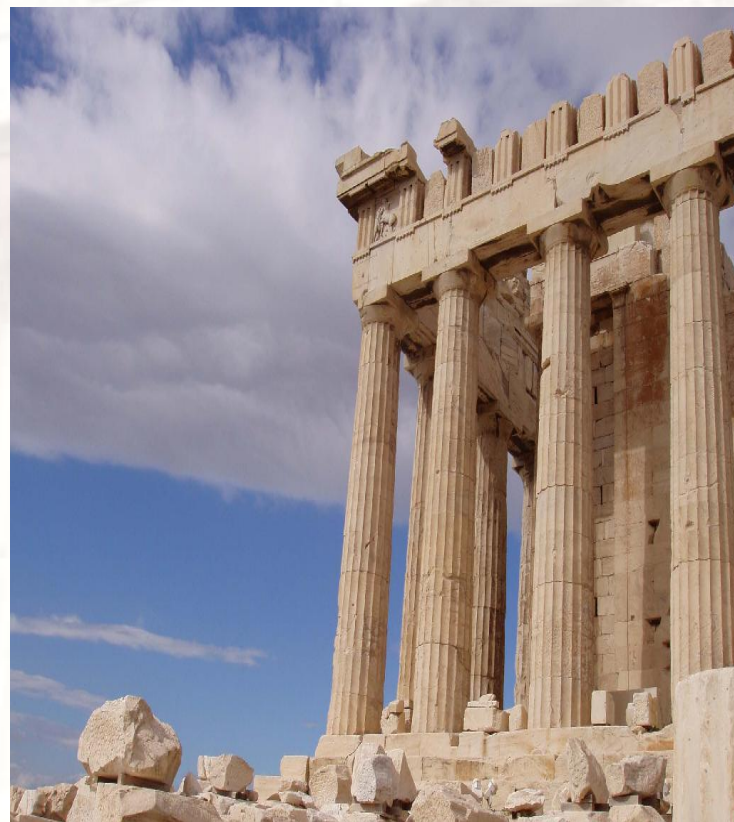


## Ch3-3

# ILP: Branch Prediction & Hardware-based speculation





# Review

## ❑ Structural Hazard

- Separate/replicate function unit, fully-pipelined function units

## ❑ Data Hazard

- RAW: double-bump, forwarding path, simple compiler scheduling
- WAR, WAW: implicit register renaming

## ❑ Control Hazard

- Resolve branch asap, Flushing, Predict-not-taken, Predict-taken, Delayed-Branch



# Control Hazard (example $P_{A-25}$ )

❑ **Flushing:** stall until the branch is resolved.

- Stall  $X$  cycles if delay slots= $x$

❑ **Predict-not-taken:**

- Hardware takes all branches as not-taken, compiler need to put frequent case in the not-taken branch
- Stall 0 cycles if not-taken, stall  $X$  cycles if taken

❑ **Predict-taken:**

- Hardware takes all branches as taken, can go forward only when get the branch target. Compiler need to put frequent case in the taken branch
- Stall  $y$  cycles waiting for branch target; stall  $X$  cycles if not taken.

❑ **Delayed Branch:**



# Static branch prediction

❑ Stall for control hazards

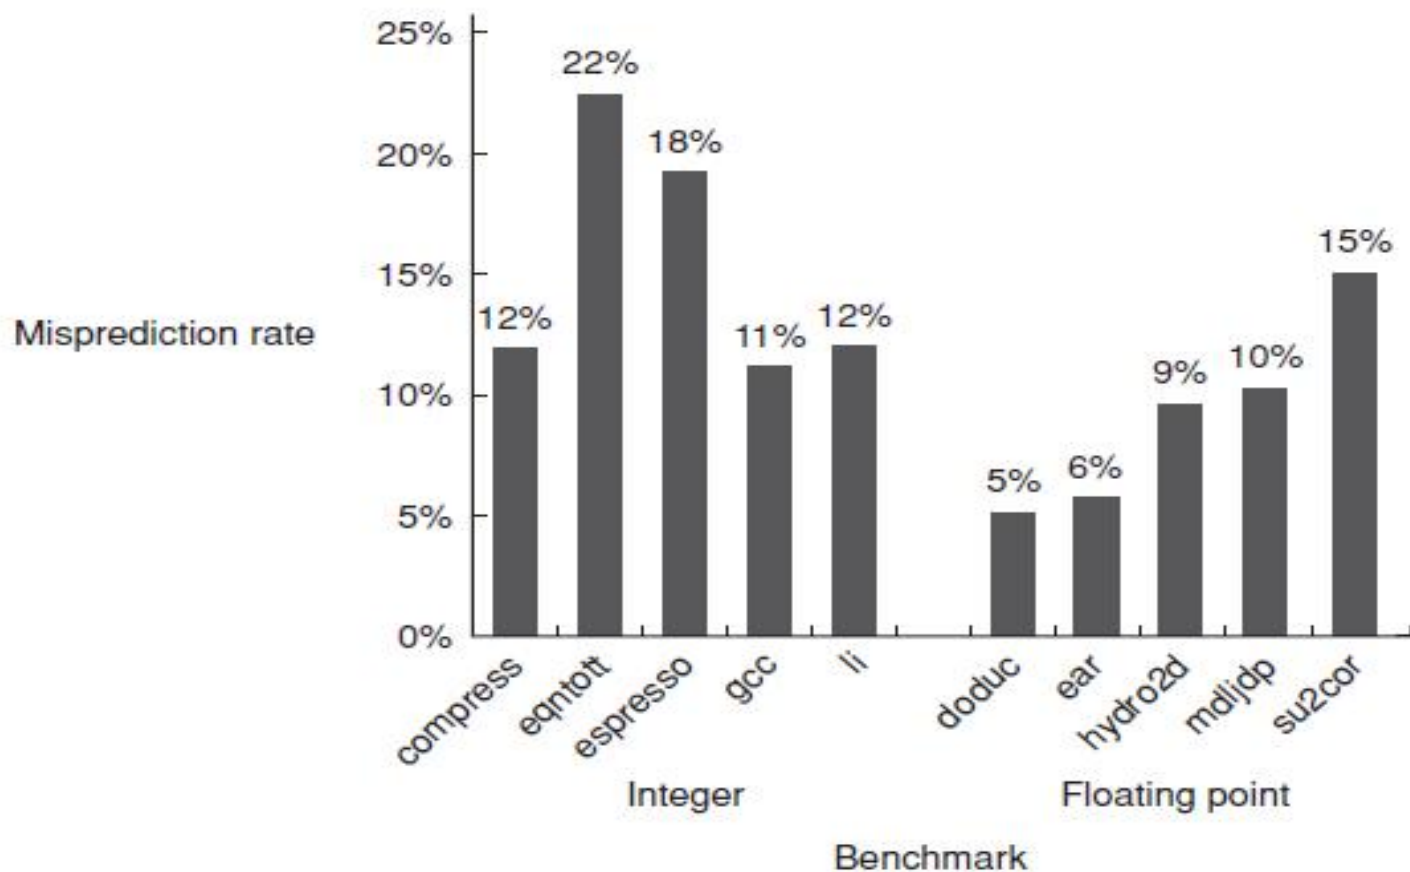
= f (branch frequency, prediction accuracy, misprediction penalty )

❑ Predict according to statistics: (predict taken)

➤ Misprediction: 9% ~ 59%, average 34%



# Predict based on profile information



Average misprediction rate 9% for FP programs,  
Average misprediction rate 15% for integer programs



# Dynamic Hardware Prediction

## --reducing branch costs

1. 1-bit Branch-Prediction Buffer
2. 2-bit Branch-Prediction Buffer
3. Correlating Branch Prediction Buffer
4. Tournament Branch Predictor
5. Branch Target Buffer
6. Integrated Instruction Fetch Units
7. Return Address Predictors





# 1-bit Branch-Prediction Buffer

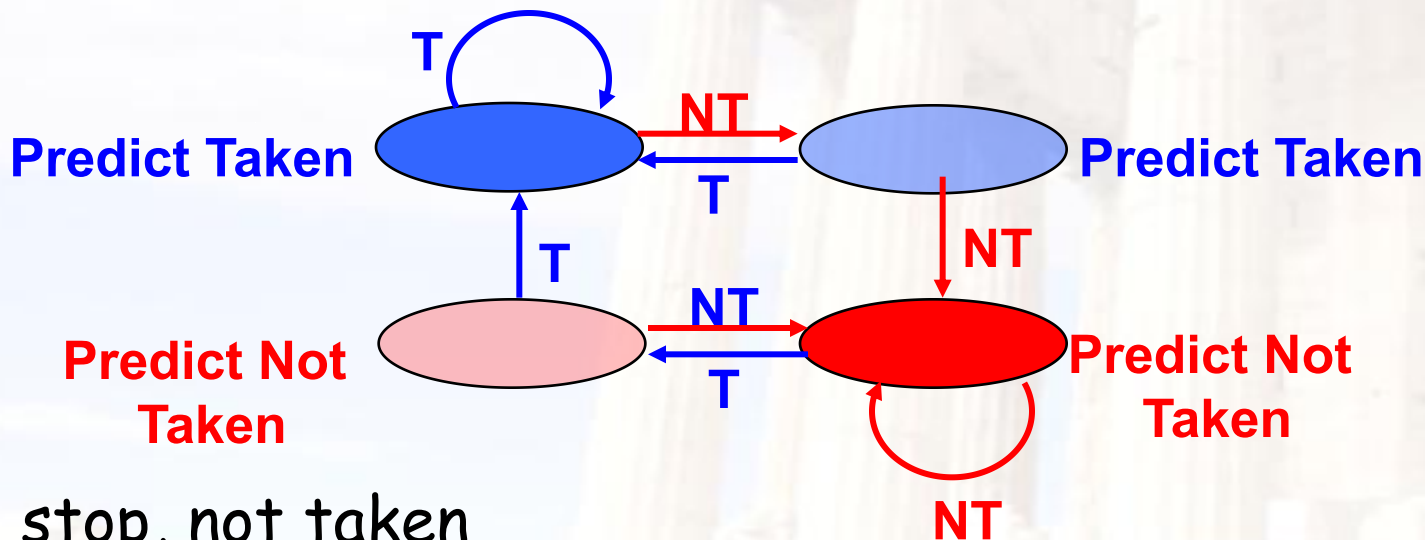
- ❑ Performance =  $f(\text{accuracy}, \text{cost of misprediction})$
- ❑ Branch History Table: Lower bits of PC address index table of 1-bit values
  - Says whether or not branch taken last time
  - No address check (saves HW, but may not be right branch)
- ❑ Problem: in a loop, 1-bit BHT will cause 2 mispredictions (avg  $2/n$  iterations before exit):
  - Only  $(n-2)/n$  accuracy even if loop  $(n-1)/n$  of the time
  - 0111.....1110111... ...1110111.....1110111...



# 2-bit Branch-Prediction Buffer

❑ Solution: 2-bit scheme where change prediction only if get misprediction *twice*:

➤ (Jim Smith, 1981, [A study of branch prediction strategies](#))



❑ Red: stop, not taken

❑ Blue: go, taken

❑ Adds *hysteresis* to decision making process

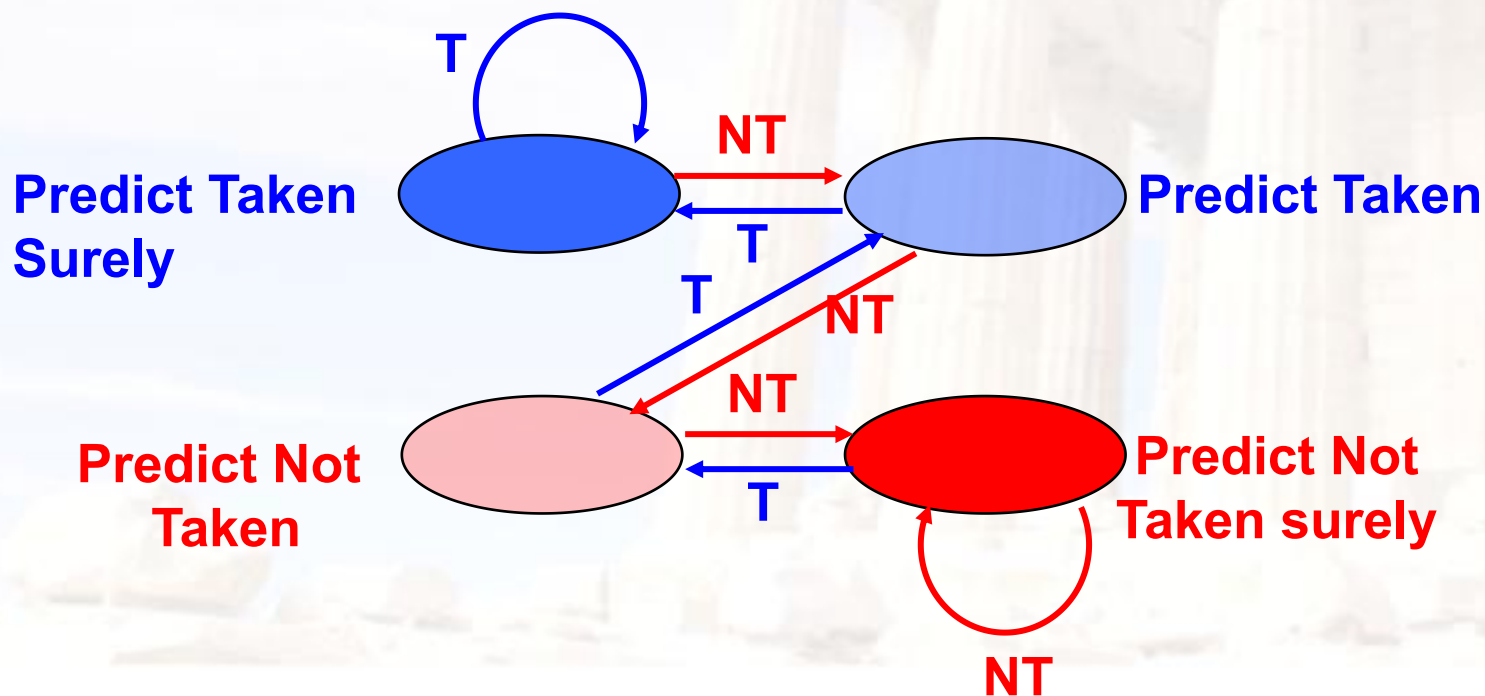
➤ 0111111.....11111011111...111011111





# An alternative 2-bit predictor

- It can be implemented using a 2-bit counter





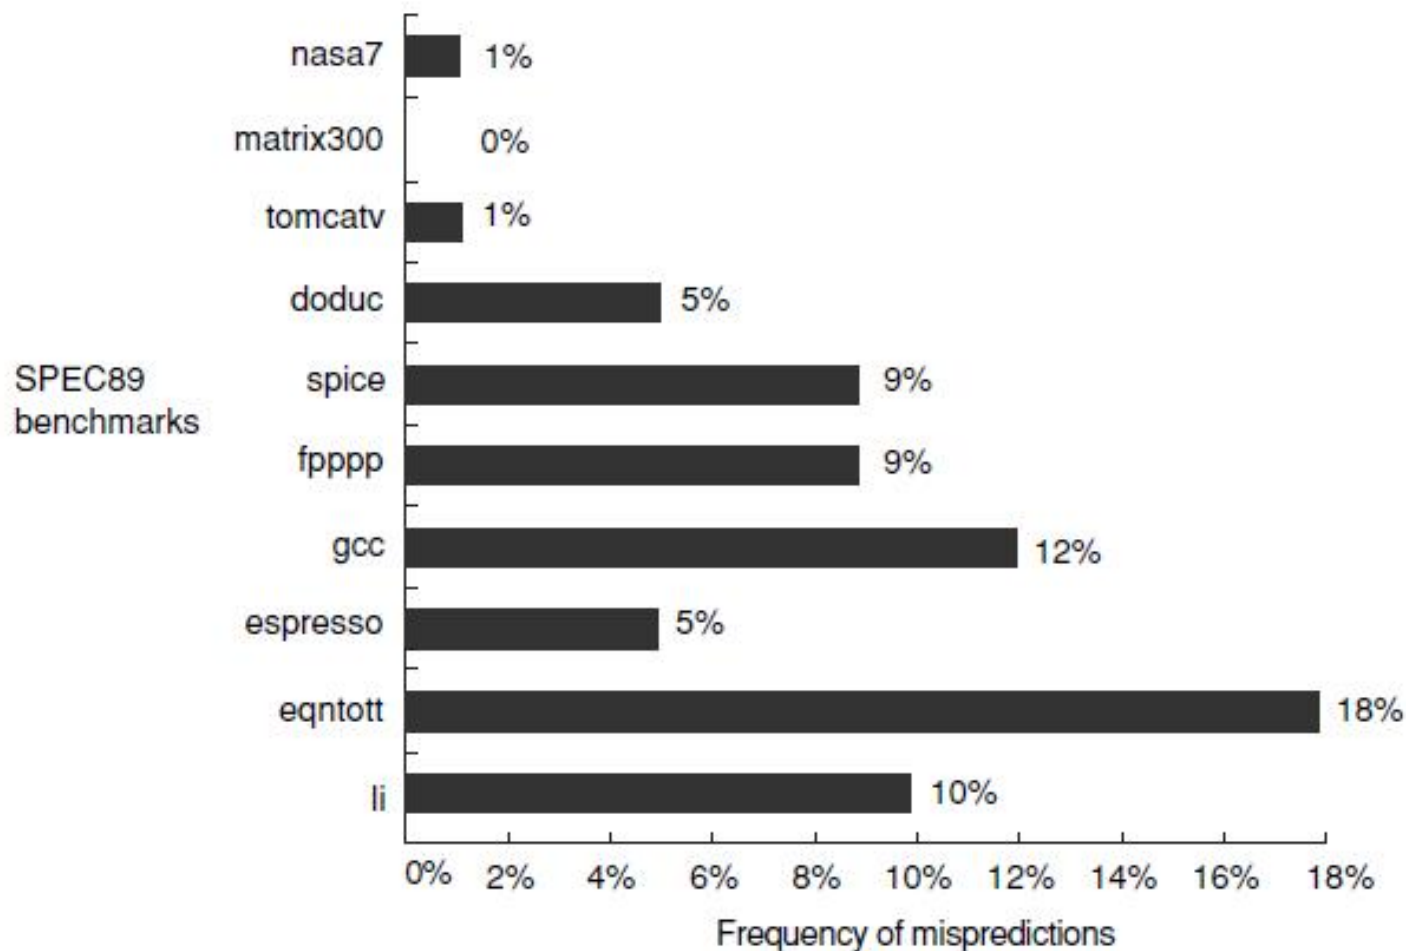
# Generalize the 2-bit predictor to n-bit predictor

- ❑ N-bit counter (  $0 \leq x < 2^{n-1}$  )
- ❑ Predict taken  $X \geq 2^{n-1}$
- ❑ Predict untaken  $X < 2^{n-1}$
- ❑ Counter + 1 when a taken branch
- ❑ Counter - 1 when a not taken branch
- ❑ Most system rely on 2-bit predictor rather than n-bit predictor.



# Accuracy of the 2-bit predictor

4096 entries, 2-bit predictor

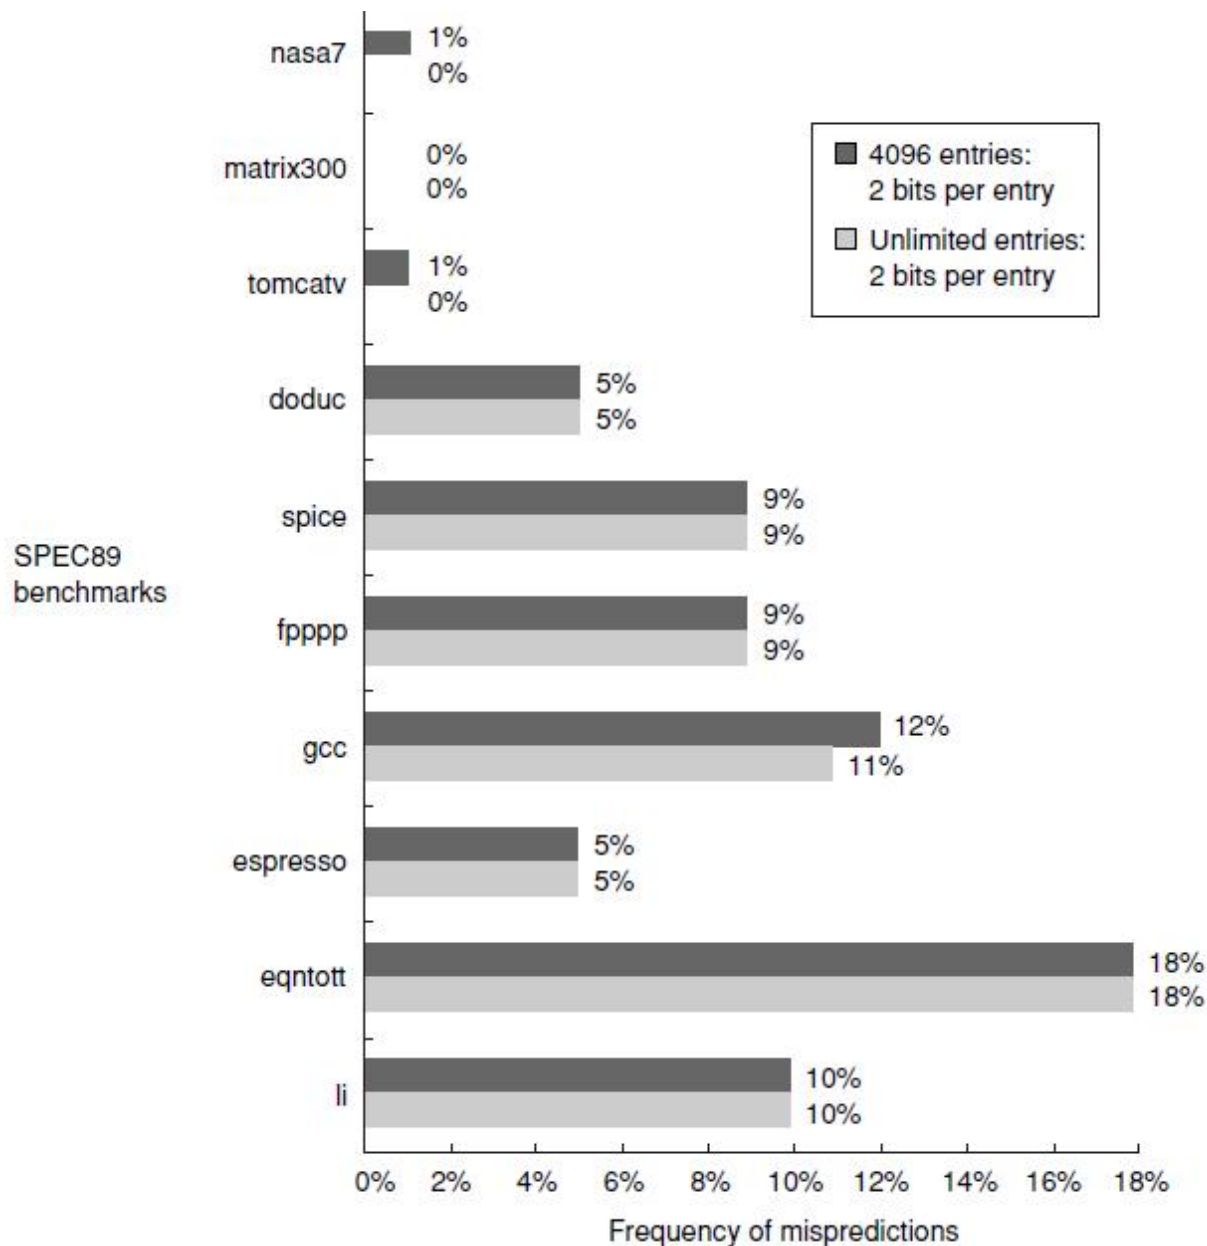


~82%

5~18%



# How about unlimited entries ?





# How to improve accuracy ?

□ Analyze the branch behavior

```
IF ( aa == 2)          SUBI   R3, R1, #2
    aa = 0;            BNEZ   R3, L1      ; br.b1
IF ( bb == 2)          (aa!=2)
    bb = 0;            ADD    R1, R0, R0 ; aa==0
IF (aa == bb ) L1: SUBI   R3, R2, #2
    {                  BNEZ   R3, L2      ; br.b2
    .....            (bb!=2)
    }                  ADD    R2, R0, R0 ; bb==0
L2: SUB    R3, R1, R2; R3=aa-bb
    BEQZ   R3, L3      ; br.b3
```



## Conclusion from the example

- ❑ B3 branch depends on b1 and b2. B3 is taken only when both b1 and b2 is untaken.
- ❑ It's impossible to predict B3 correctly only depending on its previous behavior.
- ❑ A branch's behavior not only rely on it's previous behavior but also on recent behavior of **other branches**.





## Let's check an example

```
If (d==0)
    d=1;
if (d==1){
    .....
}
```

□ Assume  $\text{Reg}[R1] = d$

```
BNEZ      R1, L1      ; br b1, (d!=0)
DADDIU    R1, R0, #1   ; d==0, so d=1
L1: DADDIU  R3, R1, #-1 ;
BNEZ      R3, L2      ; br b2, (d!=1)
    .....
L2:
```



# Feature of the code:

d的初值	d==0?	B1	在b2以前的d值	d==1?	b2
0	Yes	Not taken	1	Yes	Not taken
1	No	Taken	1	Yes	Not taken
2	No	Taken	2	no	Taken

Conclusion: b2 will be **not taken** if **b1 is not taken**

So, we can get accurate prediction if using correlate predictor.



# Always mispredict if using one-bit predictor

d=?	b1 预测	b1 动作	新的b1 预测	b2 预测	b2 动作	新的b2 预测
2	NT	T	T	NT	T	T
0	T	NT	NT	T	NT	NT
2	NT	T	T	NT	T	T
0	T	NT	NT	T	NT	NT

预测与实际动作  
总是相反

预测与实际动作  
总是相反



# Correlating predictor

❑ Correlating predictor is composed of two parts ( two bits ):

- 1st bit: prediction for previous br. is NT,
- 2nd bit: prediction for previous br. Is T

❑ Four combinations

预测组合	上一次Br为NT,预测本次位(看第一位)	上一次Br为T,预测本次为(看第二位)
NT/NT	NT	NT
NT/ T	NT	T
T /NT	T	NT
T / T	T	T

Note, previous Br might be a different br.



# Prediction with correlating predictor (1,1) with initial value = NT/NT.

d=?	b1 预测	b1 动作	新的 b1 预测	b2 预测	b2 动作	新的 b2 预测
2	NT/NT	T	T/NT	NT/NT	T	NT/T
0	T/NT	NT	T/NT	NT/T	NT	NT/T
2	T/NT	T	T/NT	NT/T	T	NT/T
0	T/NT	NT	T/NT	NT/T	NT	NT/T

粗体字表示本次预测



预测出错



b1动作与b2预测值选用关系



预测正确



b2动作与b1预测值选用关系



新预测值

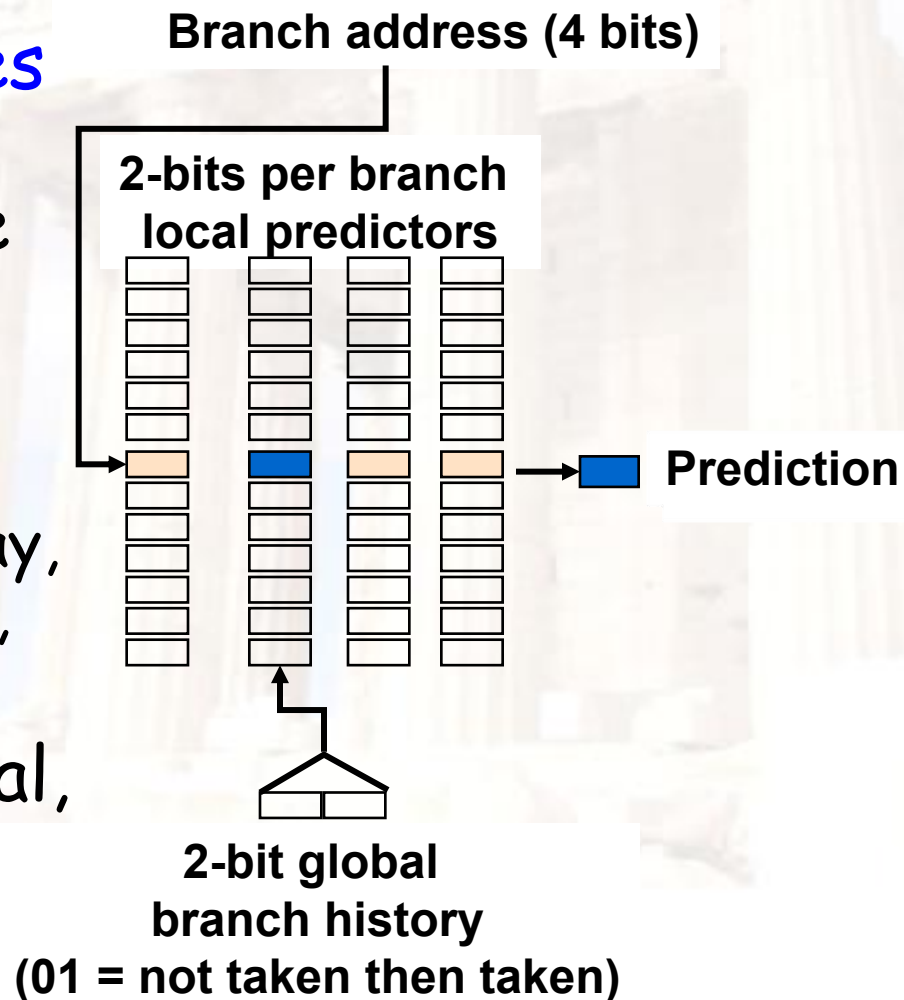


# Correlating Branches prediction buffer

❑ Idea: taken/not taken of recently executed branches is related to behavior of next branch (as well as the history of that branch behavior)

➤ Then behavior of recent branches selects between, say, 4 predictions of next branch, updating just that prediction

❑ (2,2) predictor: 2-bit global, 2-bit local







## 4. Correlating Branches prediction buffer

Assume branch history with branch i in blue box.

0 0 0 0 0 1 1 1 0 0 0 0 0 1 1 1 0 0 0 0 0 0 1 1 1

Predict:    **v**   **v**            **x**            **v** **v**            **x**            **v** **v**            **v**

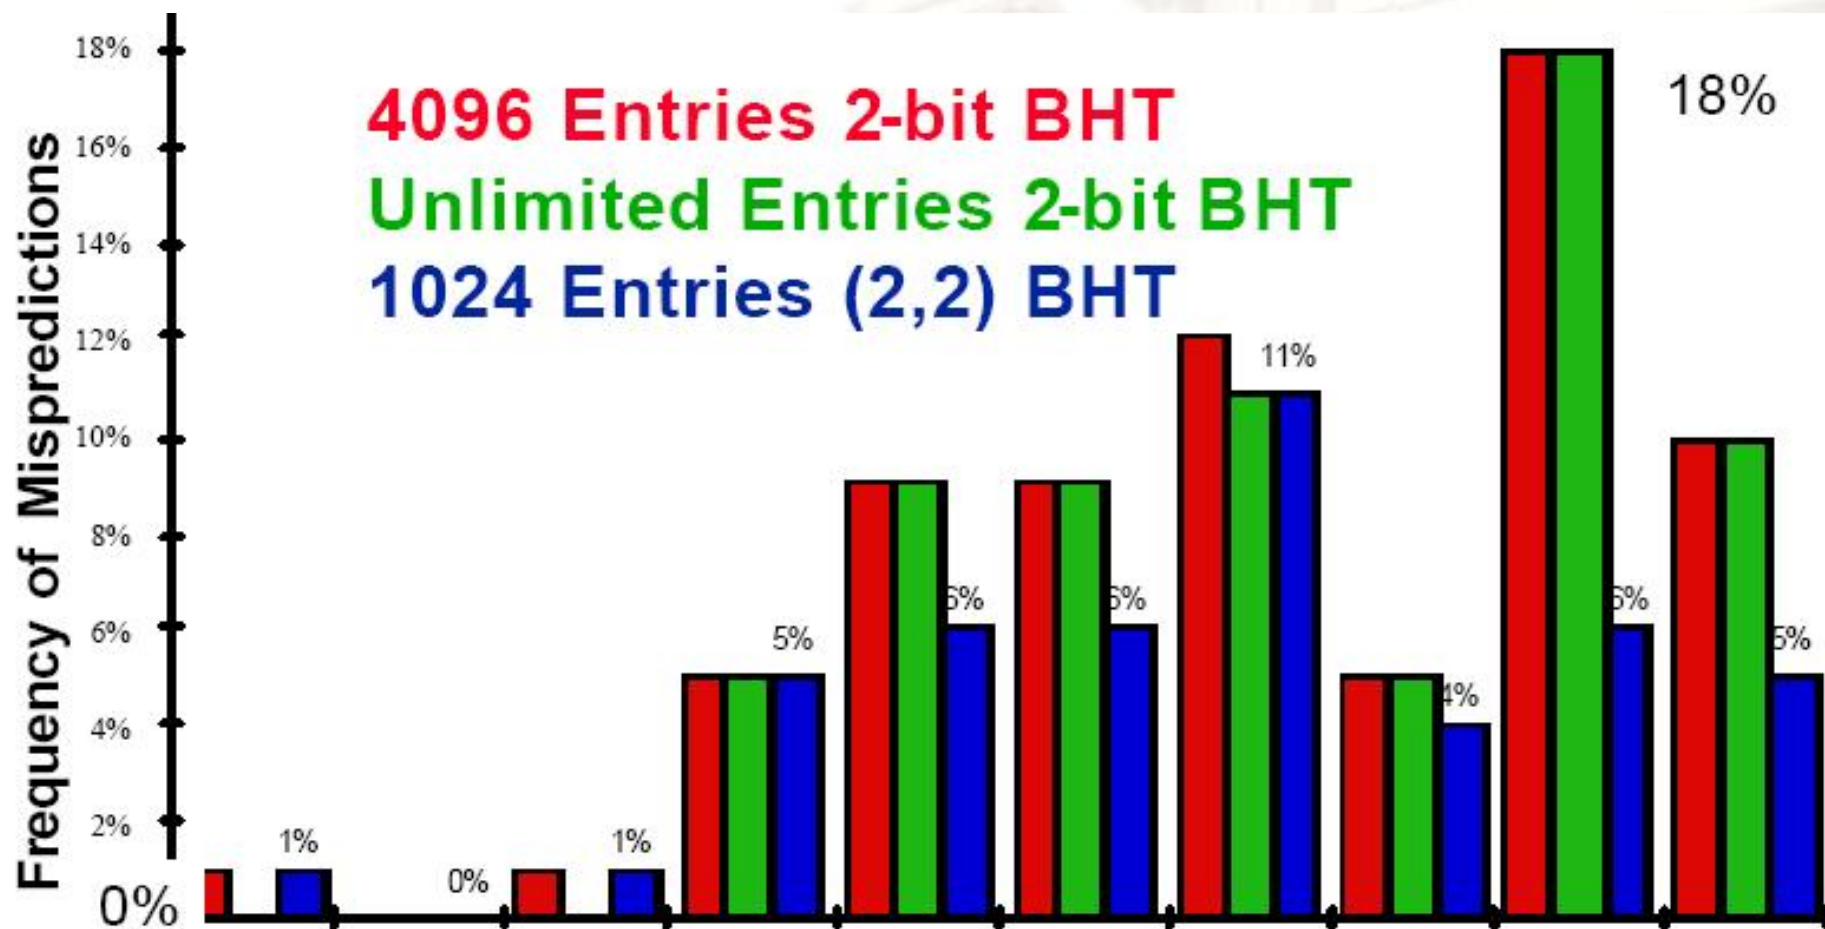




- ❑  $(m, n)$  predictor
- ❑  $m$ : previous  $m$  branch instruction
- ❑  $2^m$



# Comparision of 2-bit predictors





# Accuracy analysis

❑ Mispredict because either:

- Wrong guess for the branch
- Got branch history of wrong branch when index the table.




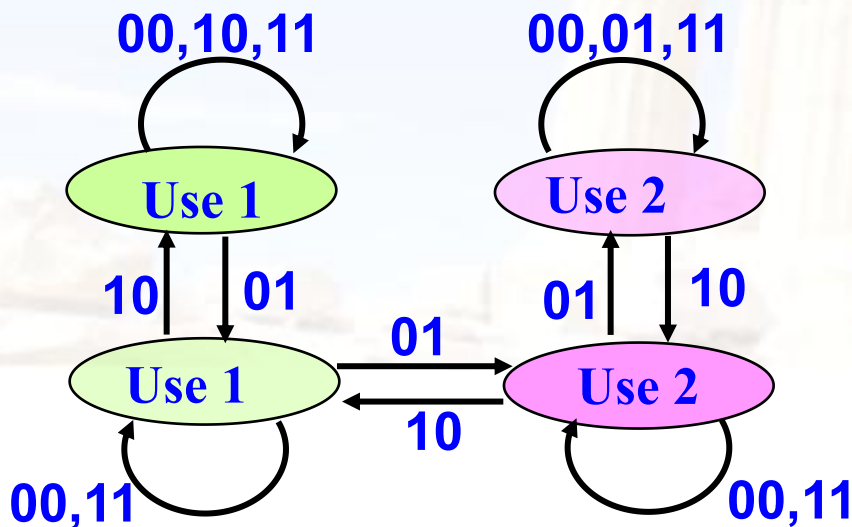
## 4、Tournament Predictors

- ❑ Motivation for correlating branch predictors is 2-bit predictor failed on important branches; by adding global information, performance improved
- ❑ **Tournament predictors**: use 2 predictors, 1 based on global information and 1 based on local information, and combine with a selector
- ❑ Hopes to select right predictor for right branch



# Tournament Predictor in Alpha 21264

- ❑ 4K 2-bit counters to choose from among a global predictor and a local predictor
- ❑ **Global predictor** also has 4K entries and is indexed by the history of the last 12 branches; each entry in the global predictor is a standard 2-bit predictor
  - 12-bit pattern:  $i$ th bit 0  $\Rightarrow$   $i$ th prior branch not taken;  $i$ th bit 1  $\Rightarrow$   $i$ th prior branch taken;







# T-Predictor in Alpha 21264 (cont.)

Local predictor consists of a 2-level predictor:

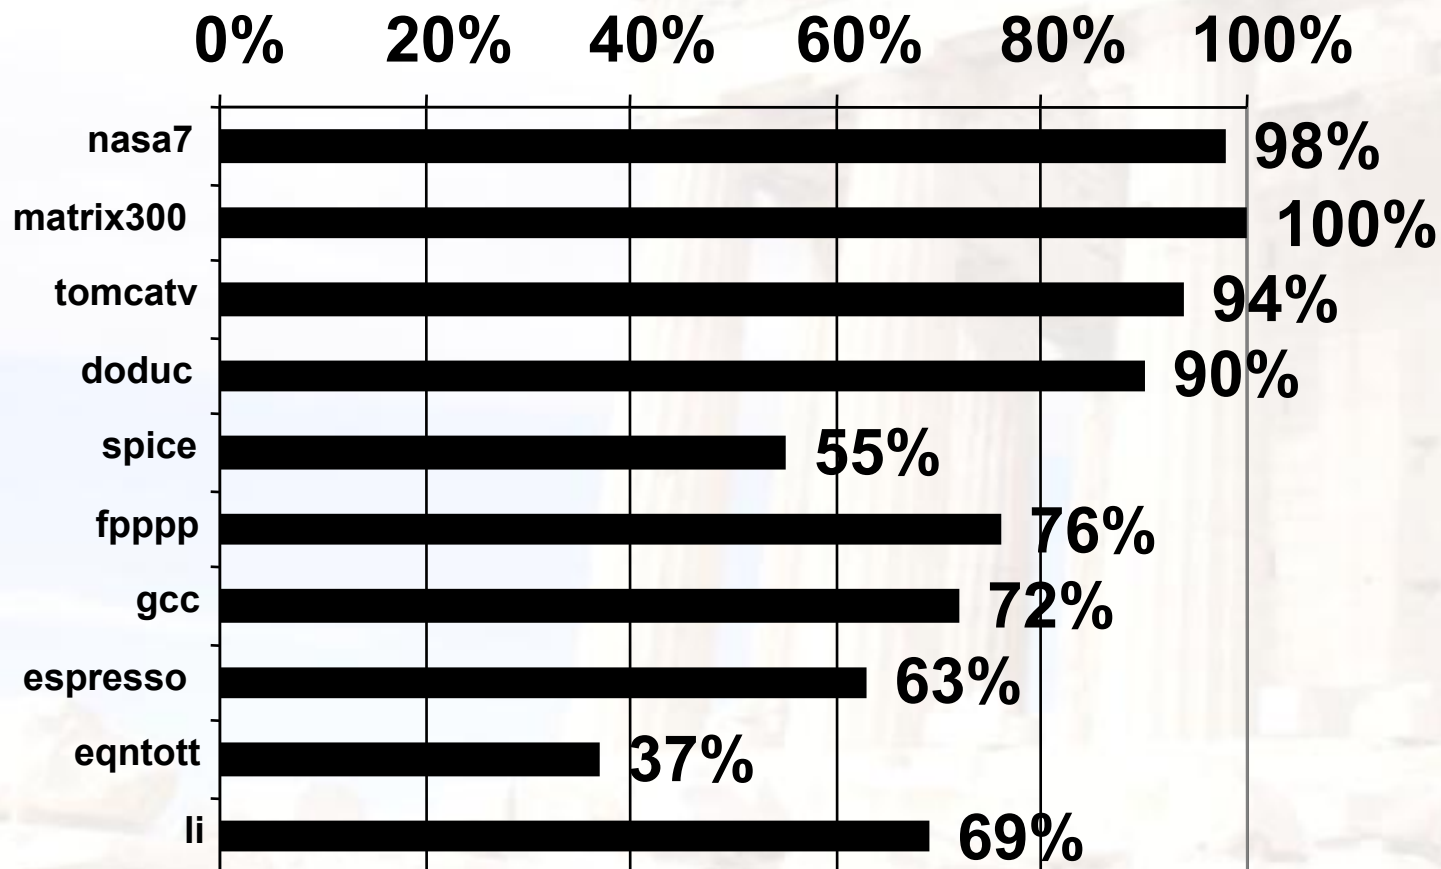
- **Top level** a local history table consisting of 1024 10-bit entries; each 10-bit entry corresponds to the most recent 10 branch outcomes for the entry. 10-bit history allows patterns 10 branches to be discovered and predicted.
- **Next level** Selected entry from the local history table is used to index a table of 1K entries consisting a 3-bit saturating counters, which provide the local prediction

Total size:  $4K \times 2 + 4K \times 2 + 1K \times 10 + 1K \times 3 = 29K \text{ bits!}$   
(~180,000 transistors)





# % of predictions from local predictor in Tournament Prediction Scheme





# Accuracy of Branch Prediction

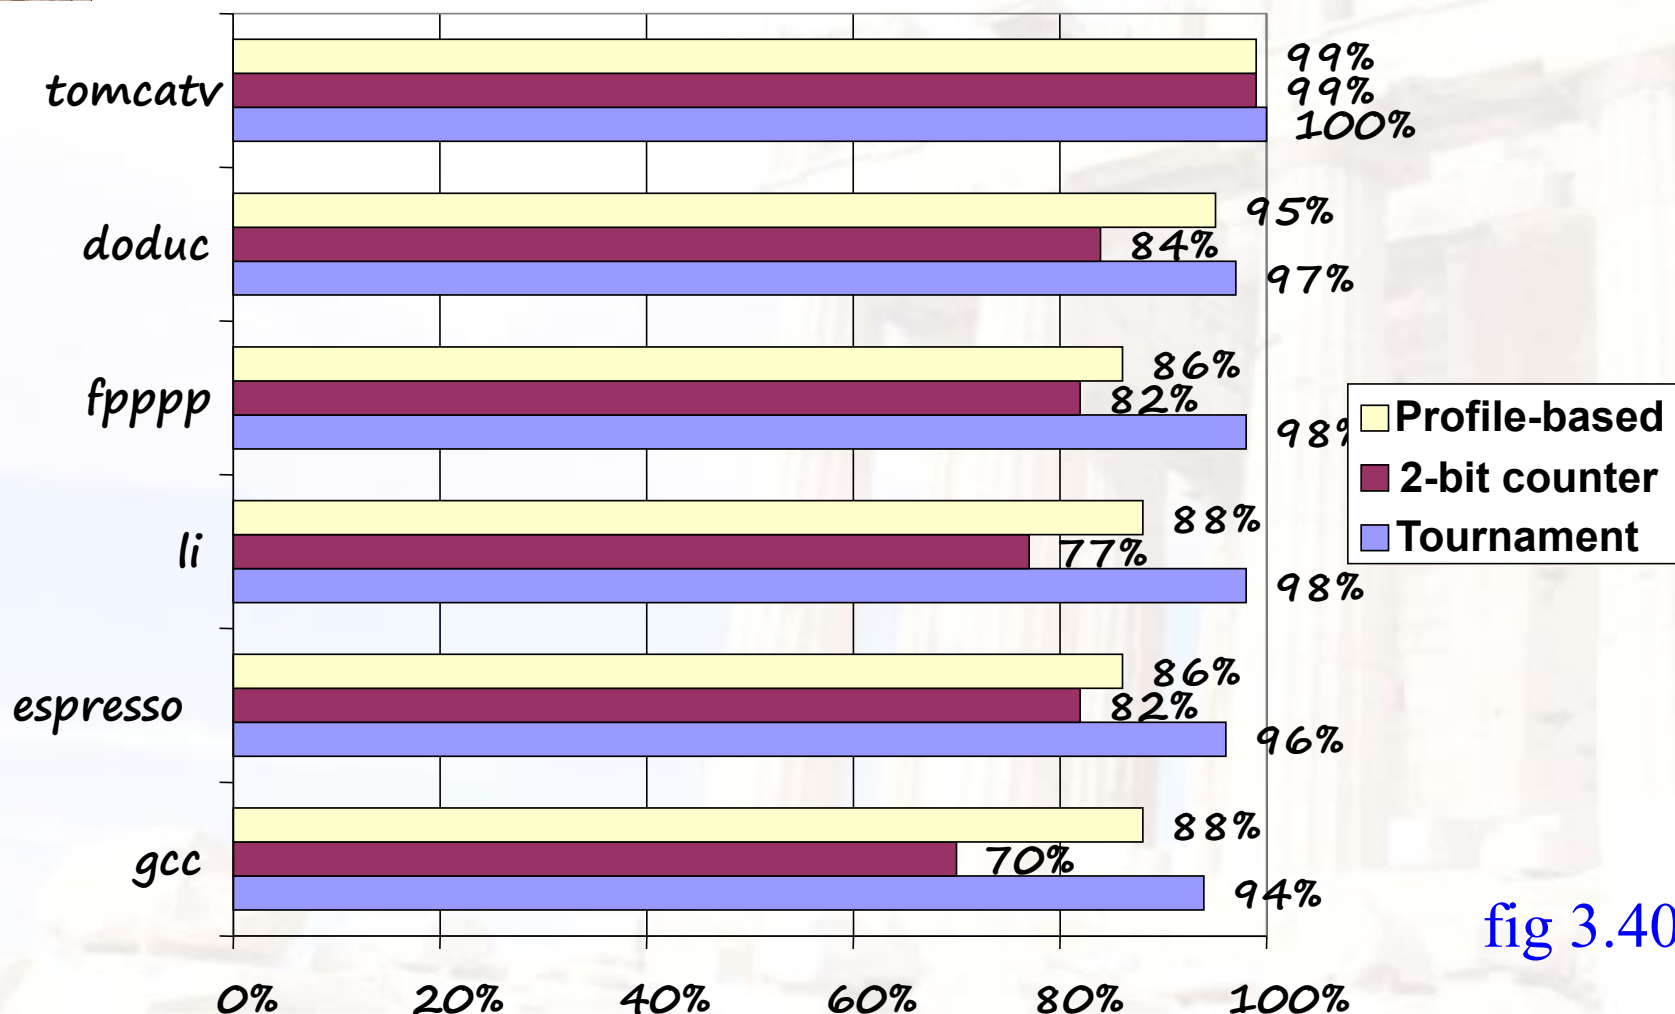
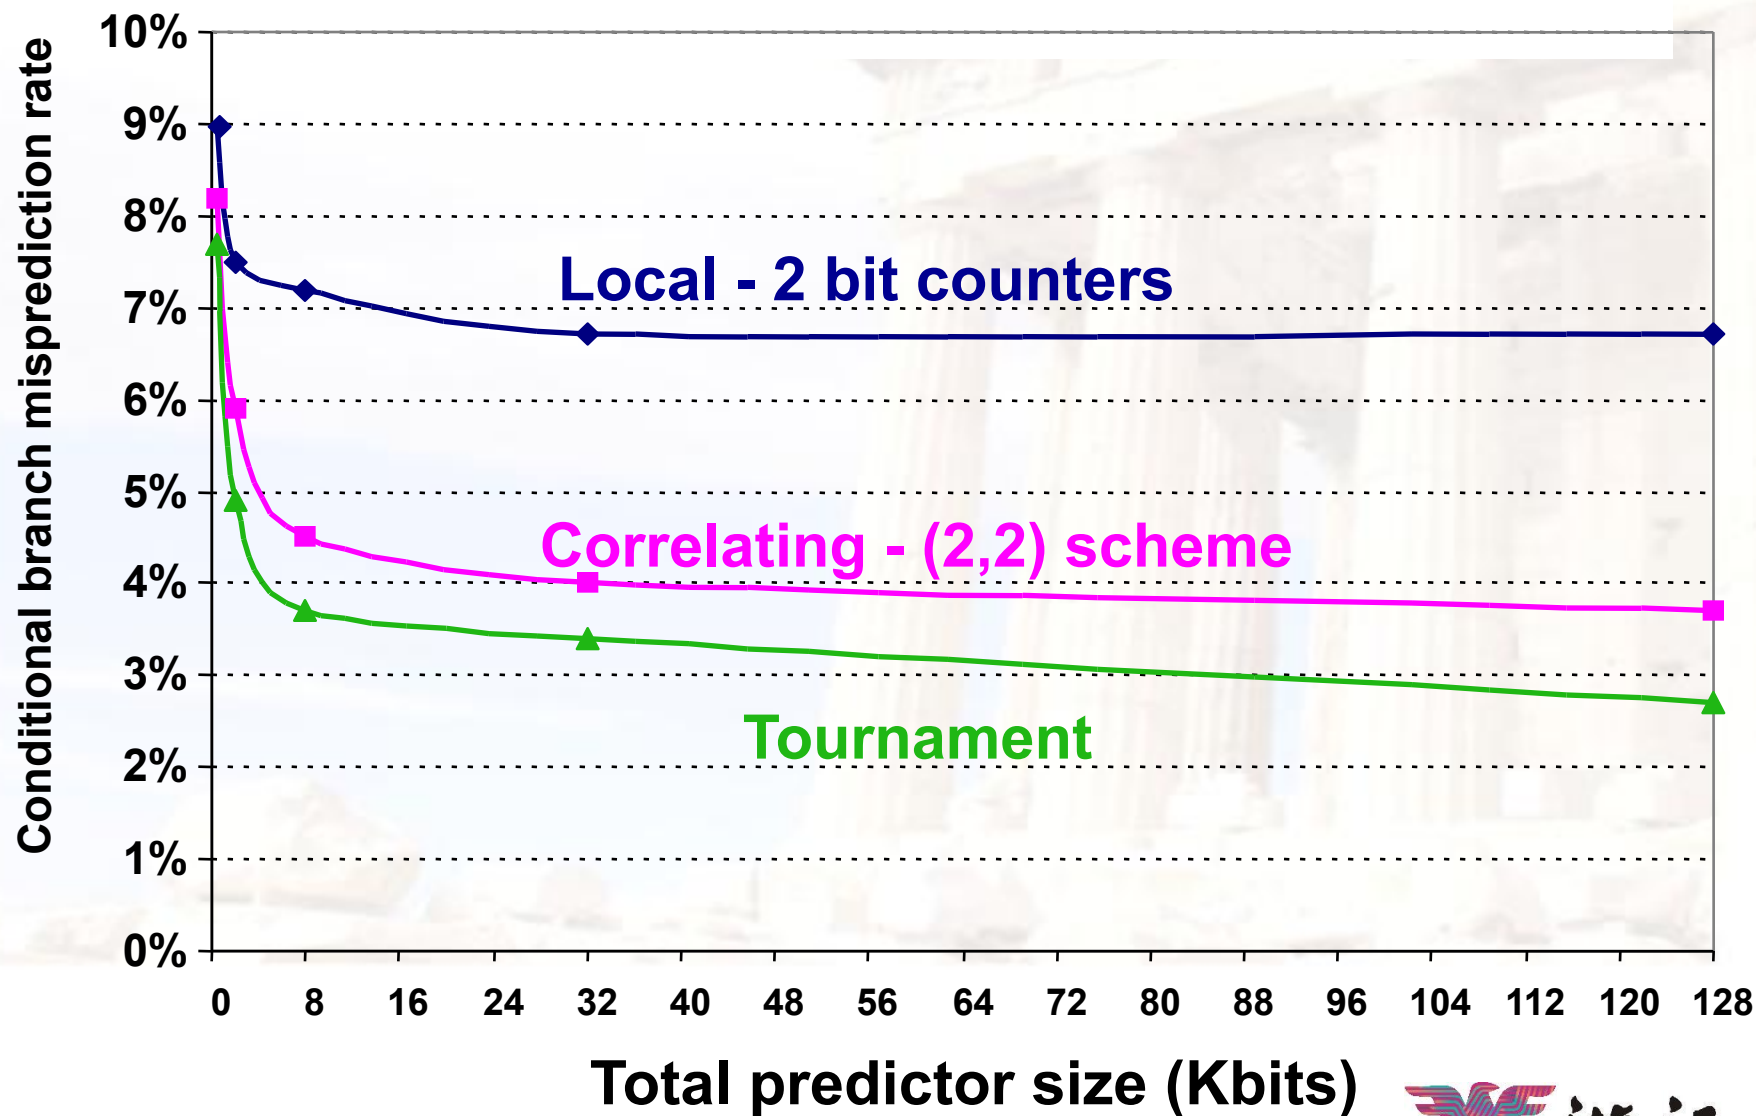


fig 3.40

□ Profile: branch profile from last execution  
(static in that is encoded in instruction, but profile)

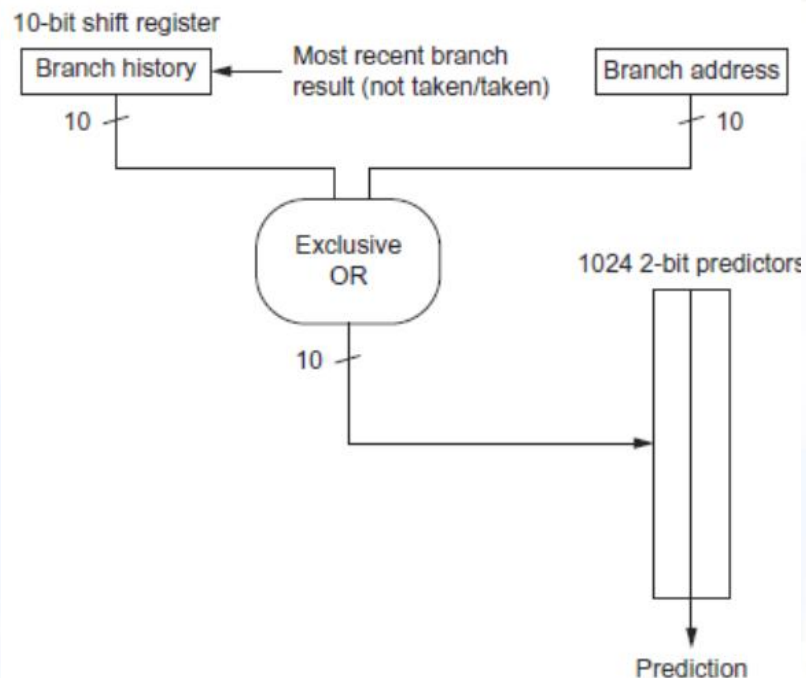


## Accuracy v. Size (SPEC89)

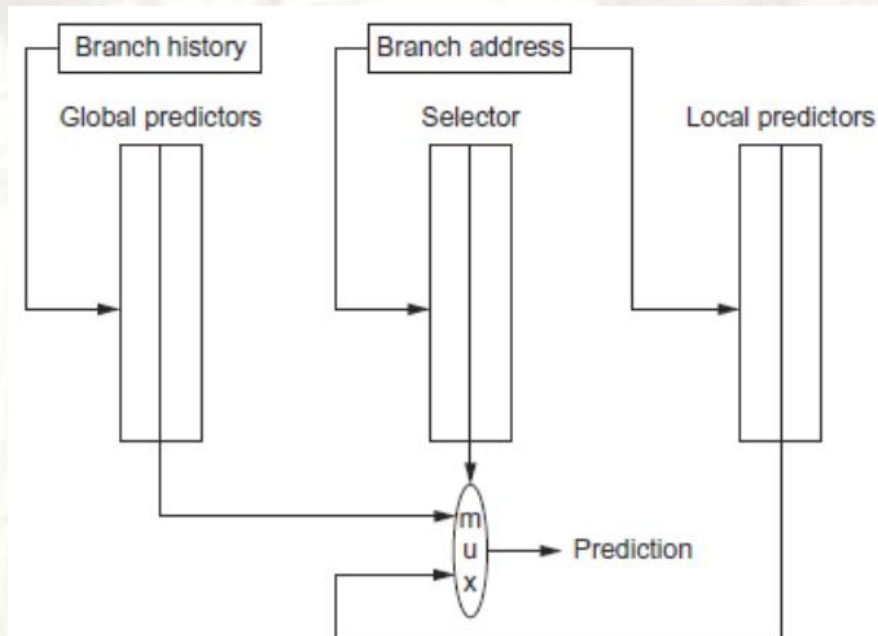




# Other Branch Prediction



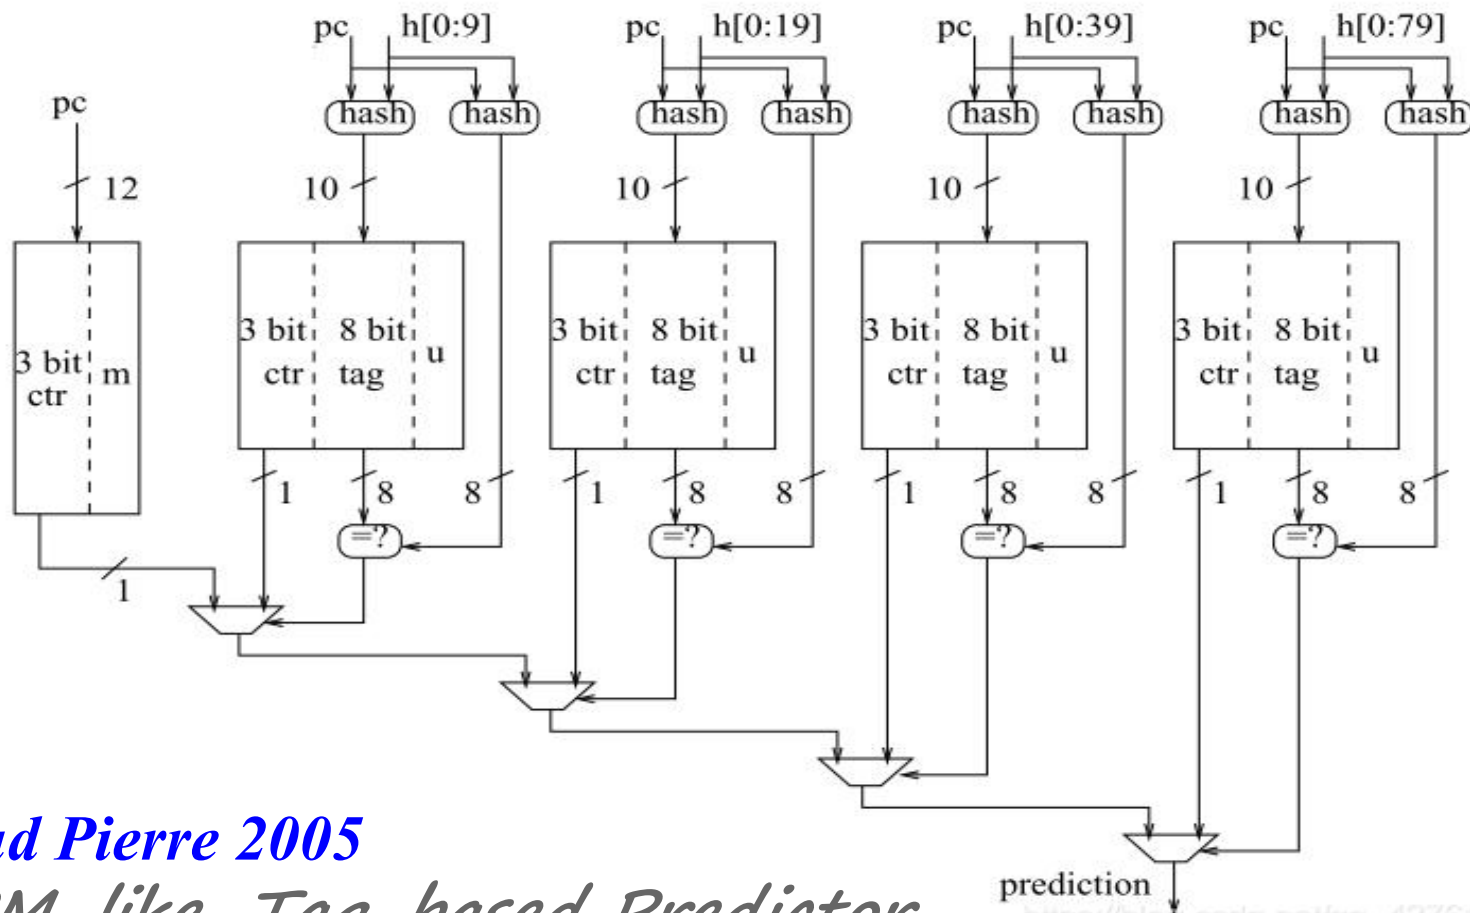
**gshare**



**tournament**



# PPM (Prediction by Partial Matching)



**Michaud Pierre 2005**

- A PPM-like, Tag-based Predictor.
- J. Instr. Level Parallelism 7 (2005)
- 第一届CBP取得了第五名的成绩



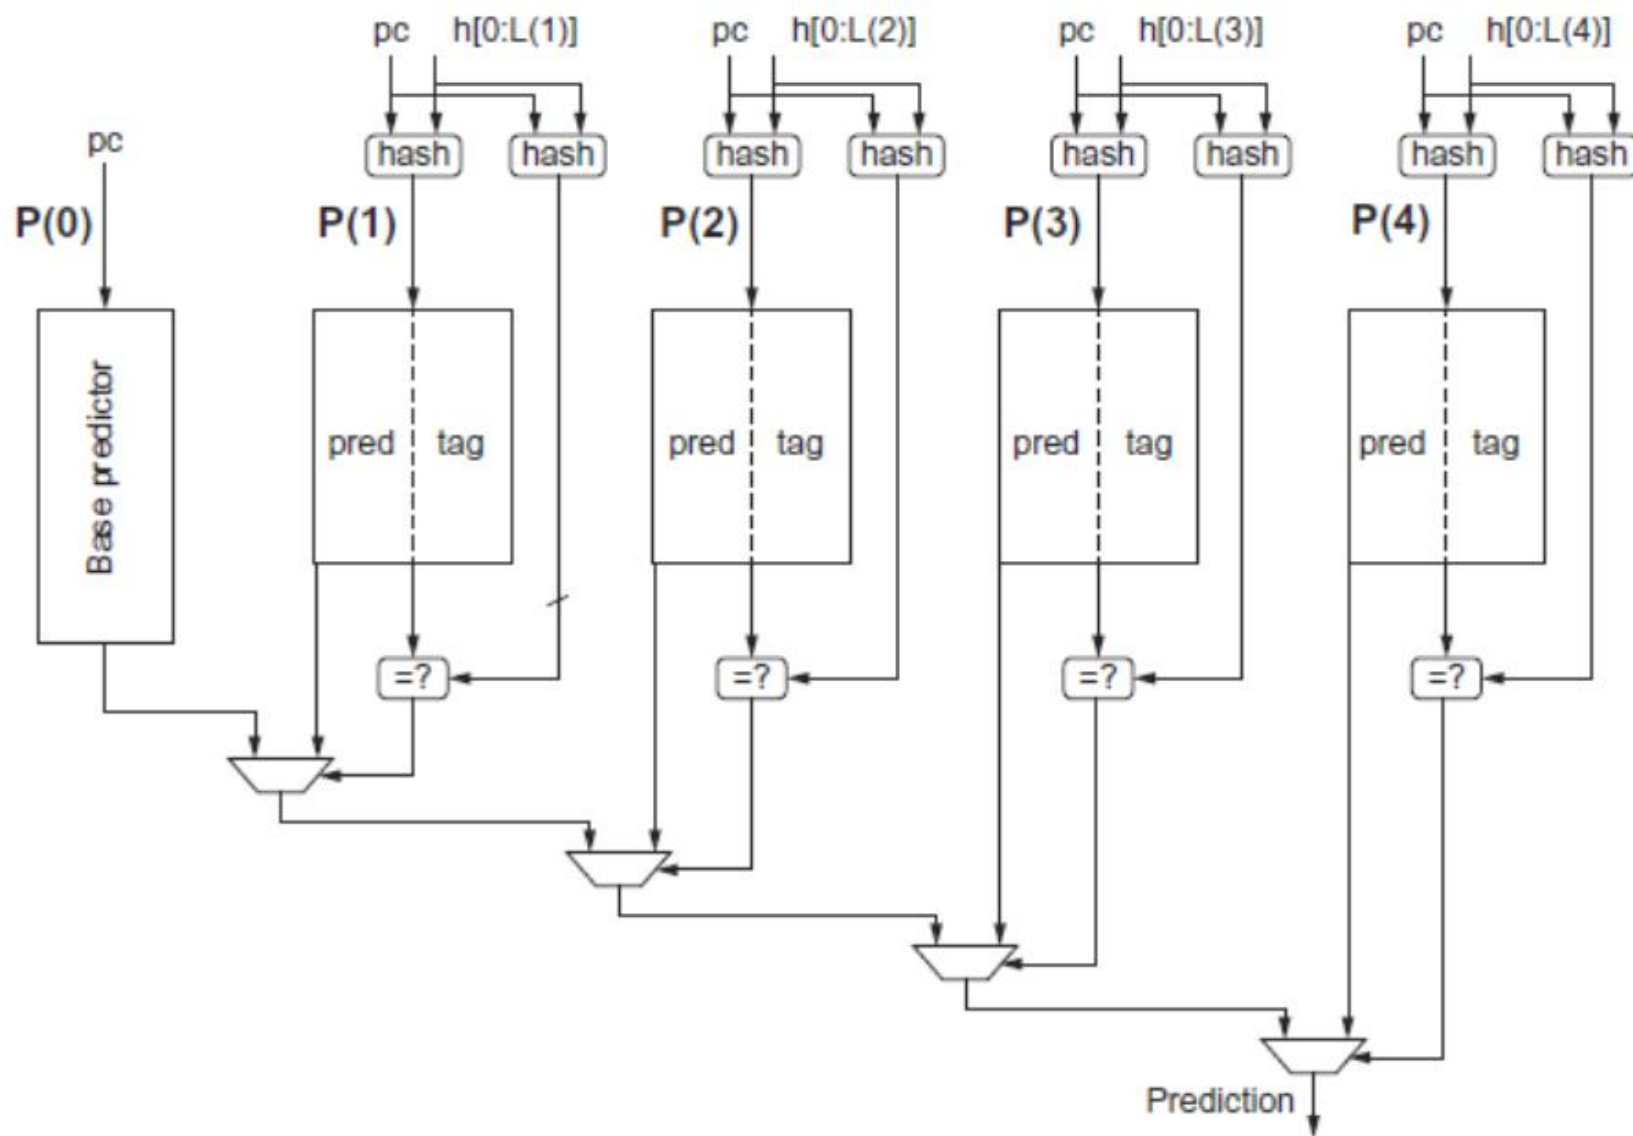


## Tagged Hybrid Predictors

- ❑ Need to have predictor for each branch and history
  - Problem: this implies huge tables
  - Solution:
    - Use hash tables, whose hash value is based on branch address and branch history
    - Longer histories may lead to increased chance of hash collision, so use multiple tables with increasingly shorter histories

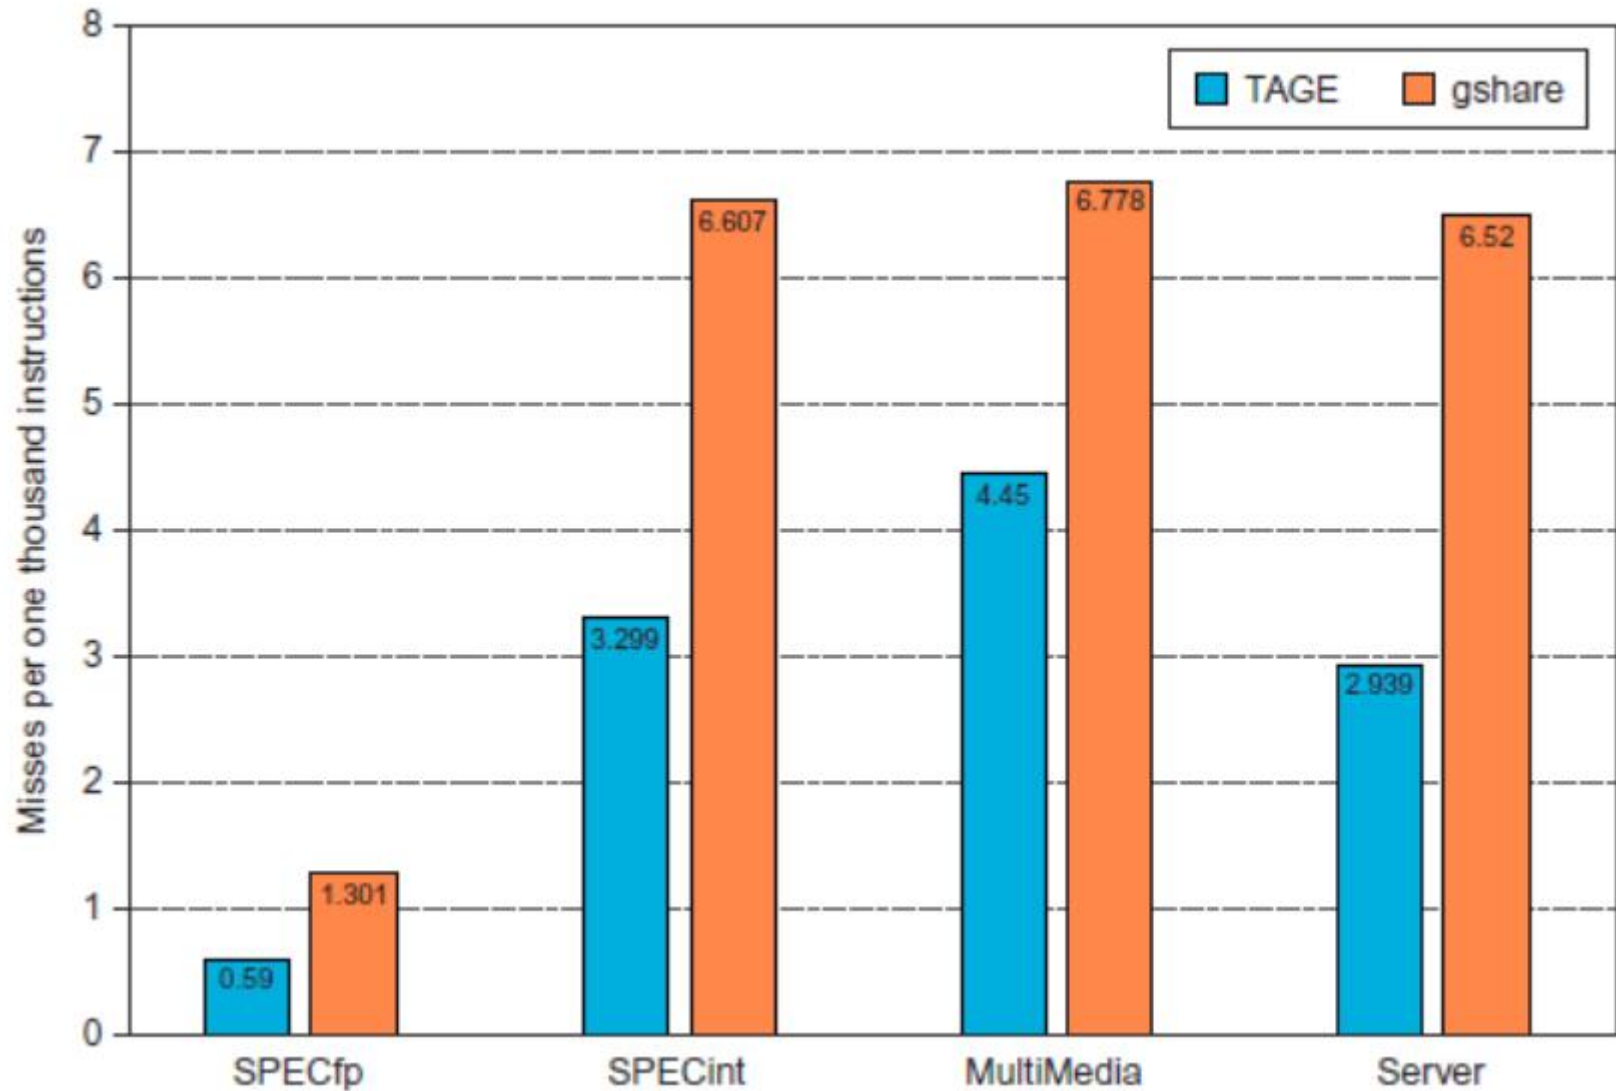


# TAGE(TAgged GEometric history length branch prediction:





# Tagged Hybrid Predictors

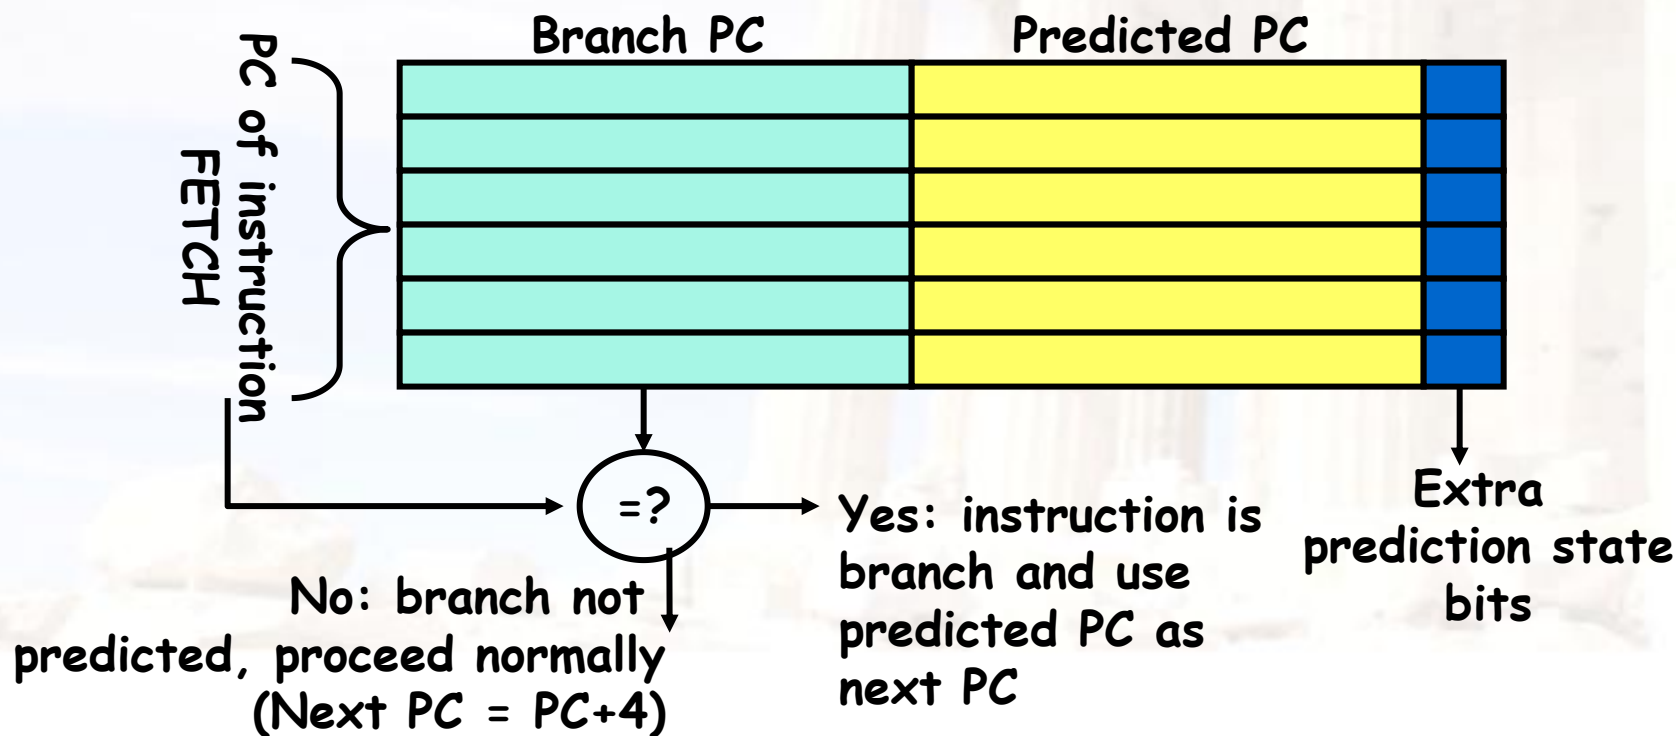




## 5. Branch Target Buffer

### ❑ Branch Target Buffer (Branch Target Cache):

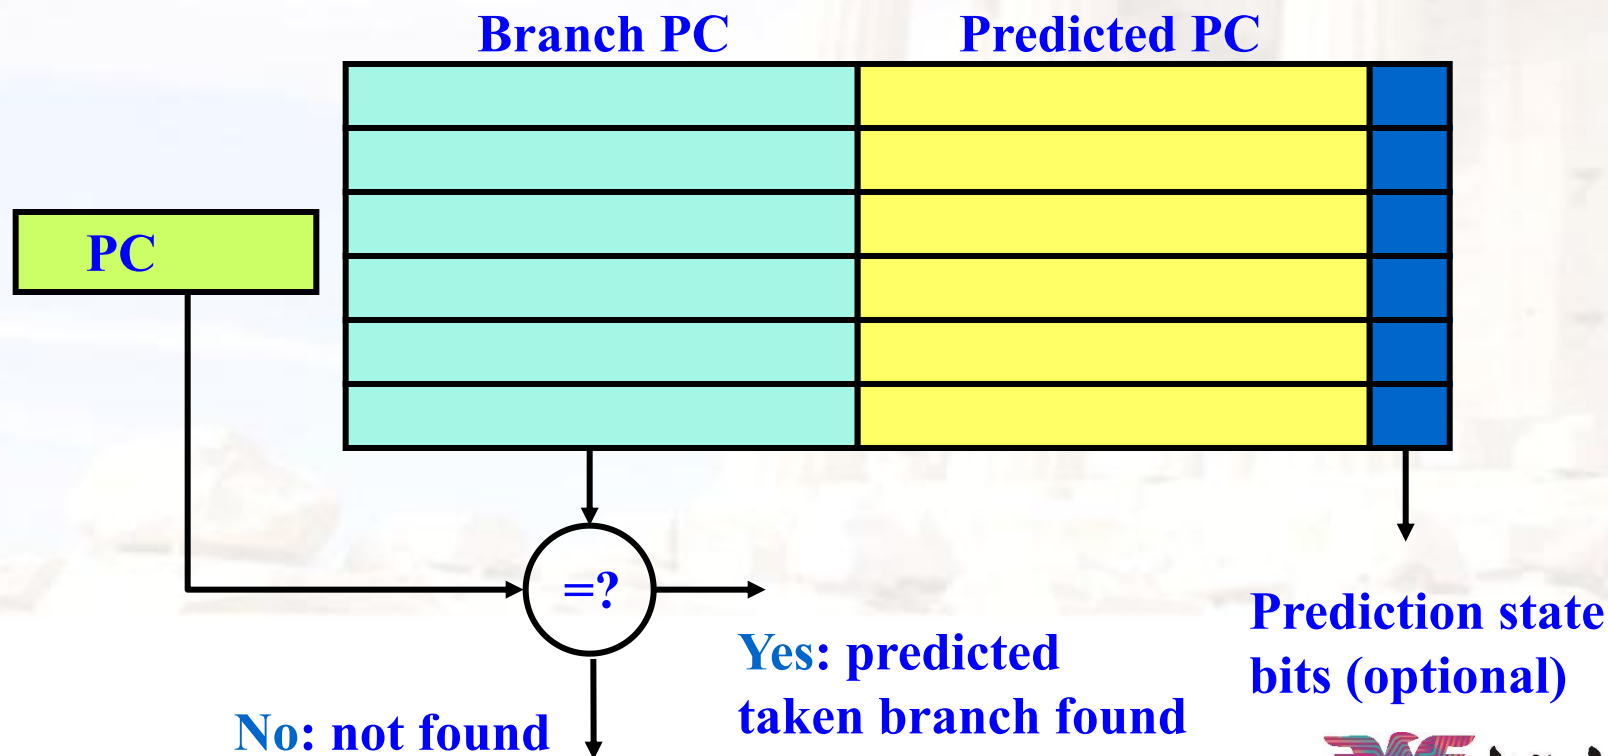
- Address of branch index to get **prediction AND branch address** (if taken)
- Note: must check for branch match now, since can't use wrong branch address





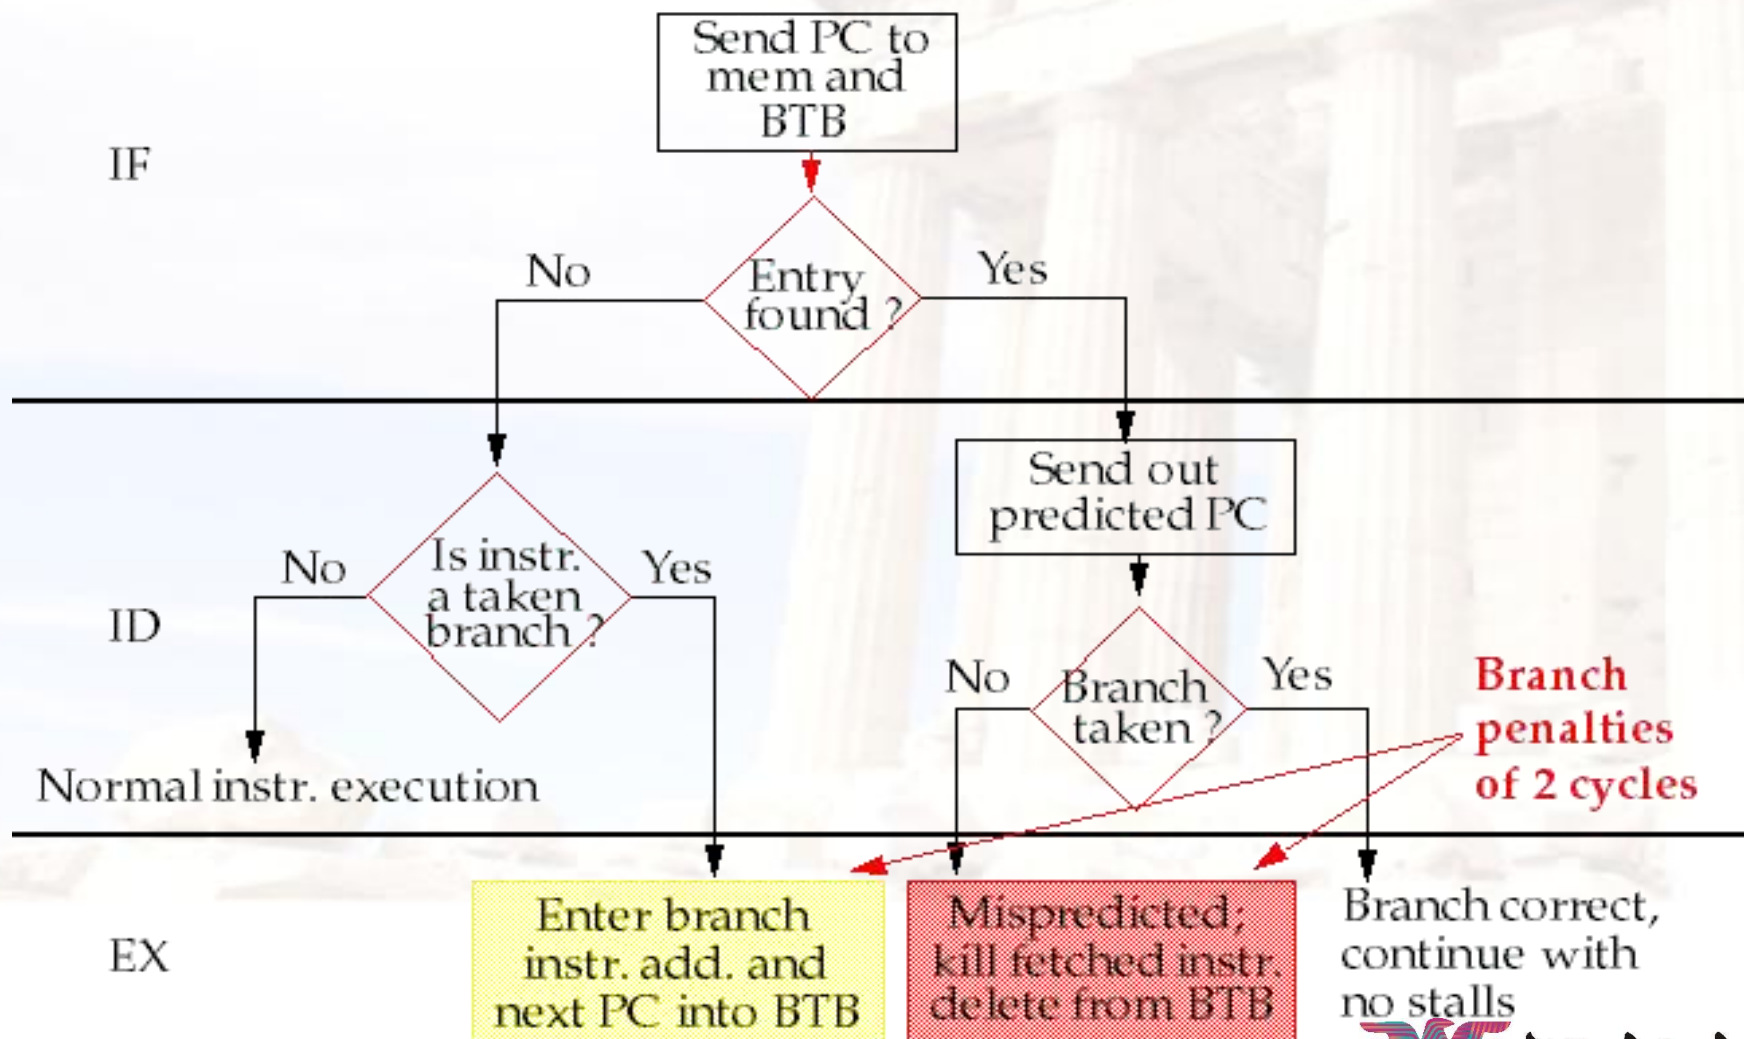
# Branch Target “Cache”

- ❑ Branch Target cache - Only predicted taken branches
- ❑ “Cache” - Content Addressable Memory (CAM) or Associative Memory (see figure)
- ❑ Use a big Branch History Table & a small Branch Target Cache





# Steps in handling an instruction with a Branch-Target Buffer







# Variations of Branch Target Buffer

- ❑ Instead of storing just the branch address, the BTB can store the actual instruction as well.
- ❑ Storing both the branch address and the actual instruction in the buffer
- ❑ Very big branch Target buffer to cache addresses or instructions from multiple paths, ex, the predicted and unpredicted direction





## 6. Integrated Instruction Fetch Units

- ❑ Increasing the number of instructions executed per clock, instruction fetch will become a significant bottleneck.
- ❑ Integrated instruction fetch unit:
  - Integrated branch prediction
  - Instruction prefetch  
(see P438 hardware prefetch & P483 ex. Instruction prefetcher )
  - Instruction memory access and buffering



## 7、Return Address Predictors

- ❑ Technique for predicting indirect jumps, whose destination address varies at run time.
- ❑ Register Indirect branch hard to predict address
- ❑ SPEC89: **procedure return** account for 85% of the indirect jumps.
- ❑ Branch-target buffer doesn't work well for procedure return



# Return Addresses Buffer

- ❑ Since stack discipline for procedures, **save return address in small buffer that acts like a stack**:
  - Caches the most recent return addresses.
  - Push return address into stack at a call;
  - pop an address off at a return
  
- ❑ SPEC benchmark shows a stack with 8-16 entries works quite well. (fig-3.22)



# Dynamic Branch Prediction Summary

- ❑ Prediction becoming important part of scalar execution
- ❑ Branch History Table: 2 bits for loop accuracy
- ❑ **Correlation**: Recently executed branches correlated with next branch.
  - Either different branches
  - Or different executions of same branches
- ❑ **Tournament Predictor**: more resources to competitive solutions and pick between them
- ❑ **Tagged Hybrid Predictors**: have predictor for each branch and history
- ❑ **Branch Target Buffer**: include branch address & prediction
- ❑ Predicated Execution can reduce number of branches, number of mispredicted branches
- ❑ **Return address stack** for prediction of indirect branches



## Chapter3 ILP

# Hardware Based Speculation



# Tomasulo Drawbacks

## ❑ Complexity

- delays of 360/91, MIPS 10000, Alpha 21264, IBM PPC 620 in CA:AQA 2/e, but not in silicon!

## ❑ Many associative stores (CDB) at high speed, Performance limited by Common Data Bus

## ❑ Non-precise interrupts!

- We will address this later

## ❑ Partially overlapped basic blocks

- integer units get ahead, beyond FP operation
- the successor basic block can not start execution until the branch is resolved, in spite of that it can be issued.



# Hardware-based Speculation

- ❑ Speculating on the outcome of branches and executing the program as if the predictions were correct.
  - Need to handle the situation where the speculation is incorrect.
- ❑ Key ideas:
  - Dynamic branch prediction to choose which instruction to execute
  - Speculation to allow the execution of instructions before the control dependences are resolved.
  - Dynamic scheduling to deal with the scheduling of different combinations of basic blocks.
- ❑ Data flow execution:
  - Operations execute as soon as their operands are available.





# Hardware-Based Speculation

- ❑ Execute instructions along predicted execution paths but only commit the results if prediction was correct
- ❑ Instruction commit: allowing an instruction to update the register file when instruction is no longer speculative
- ❑ Need an additional piece of hardware to prevent any irrevocable action until an instruction commits
  - I.e. updating state or taking an executionb

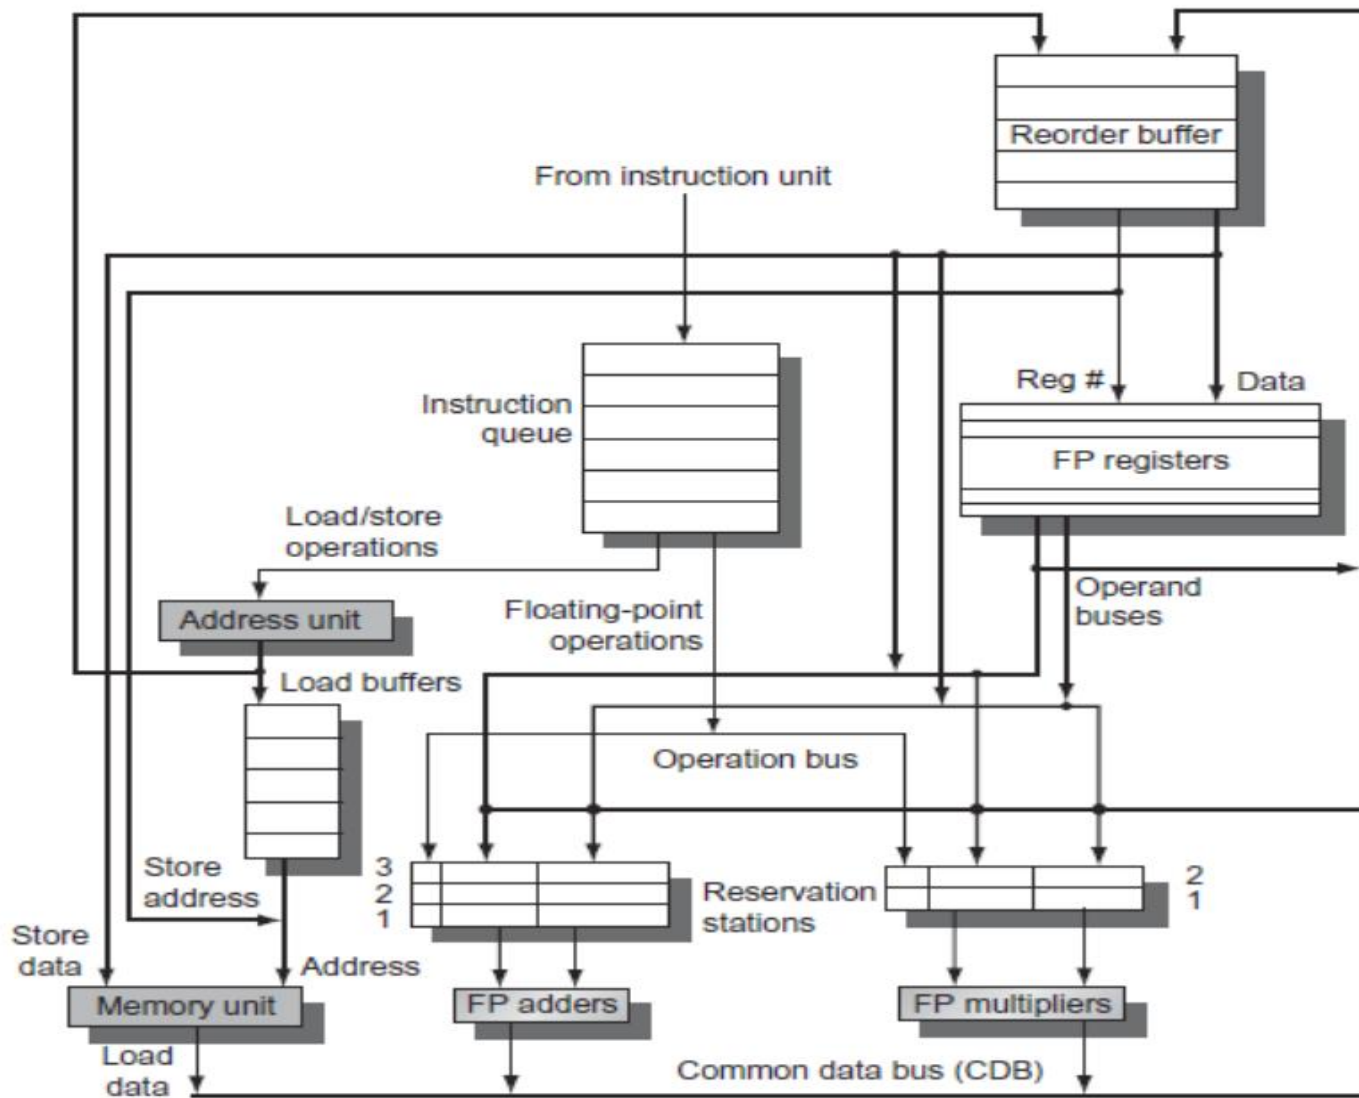


# Speculative execution based on Tomasulo's algorithm

- ❑ Separate the process of completing execution and the bypassing of results among instructions from instruction commit (reg file or memory update).
- ❑ Allows other (speculative) instructions to execute, but no results are committed until we know the instruction is no longer speculative.
- ❑ Allow instructions to **execute out of order** but force them to **commit in order**, which helps with handling exceptions properly.

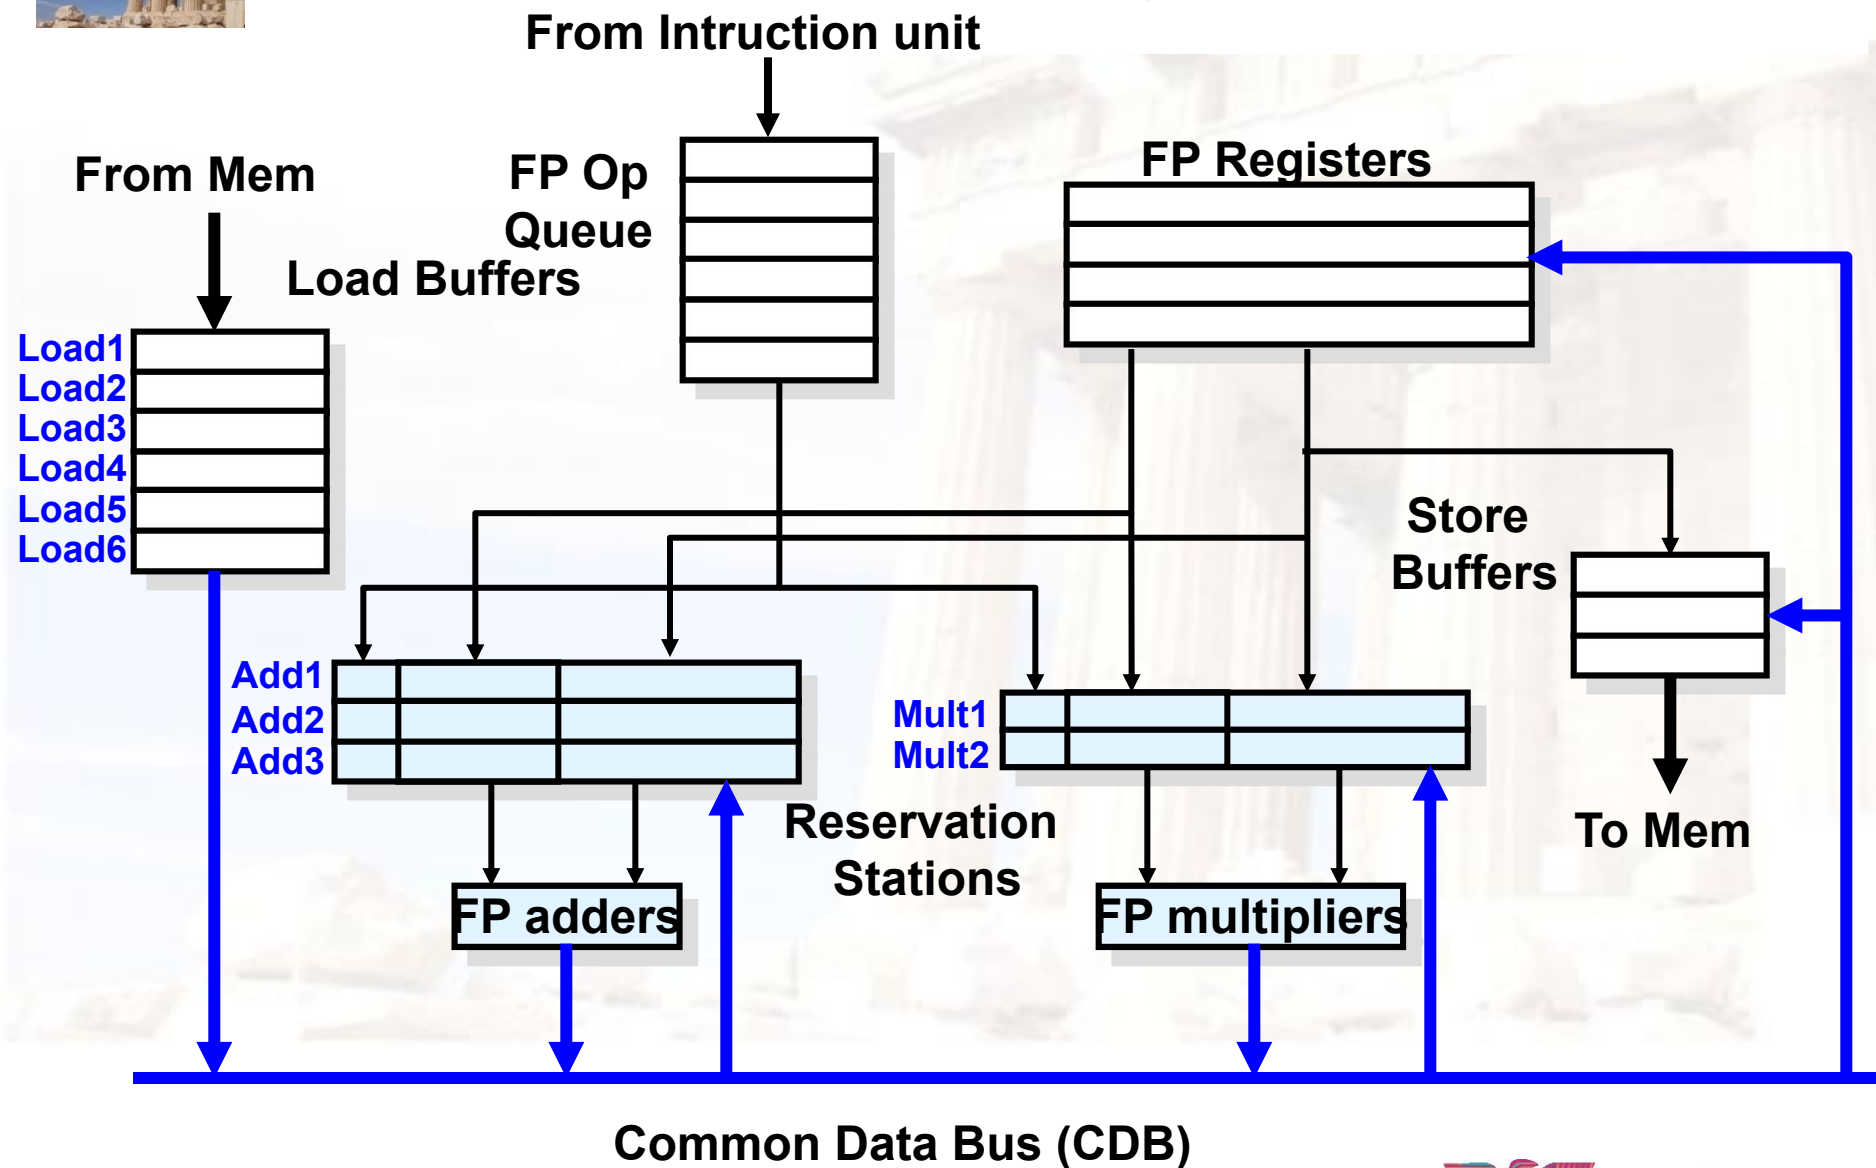


# Extend Tomasulo's Algorithm to handle speculation





# Review: Tomasulo Algorithm





# Difference from Tomasulo Algorithm

- ❑ **Reorder Buffer** is added
- ❑ Eliminate store buffer (Mem write is ordered in Reorder buffer)
- ❑ **Register renaming is done via reorder buffer** instead of reservation station
- ❑ Reservation only for storing opcodes and operands during the instruction flying between issue and execution.
- ❑ ROB hold the instruction results and supply operands in the interval between completion of instruction execution and instruction commit.



# Entry of ROB

- ☐ Instruction type
- ☐ Destination field
- ☐ Value field
- ☐ Ready field
- ☐ Exception vector



# 4 Steps of Speculative Tomasulo Algorithm

## 1. Issue—get instruction from FP Op Queue

If reservation station **and reorder buffer slot** free, issue instr & send operands **& reorder buffer no. for destination** (this stage sometimes called "**dispatch**") , update control entries "in use".

## 2. Execution—operate on operands (EX)

When both operands ready then execute; if not ready, watch CDB for result; when both in reservation station, execute; **checks RAW**

## 3. Write result—finish execution (WB)

Write on Common Data Bus to all awaiting FUs **& reorder buffer**; mark the reservation station available.

## 4. Commit—**update register with reorder result**

When instruction at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer. **Mispredicted branch flushes reorder buffer (sometimes called "graduation")**





# What to do when commit ?

when an instruction reaches the head of ROB and its result is present in the buffer, then

## ❑ For non-Branch instruction

- Update the register with the result
- Remove the instruction from ROB
- When it's a STORE, update the memory instead of Reg.

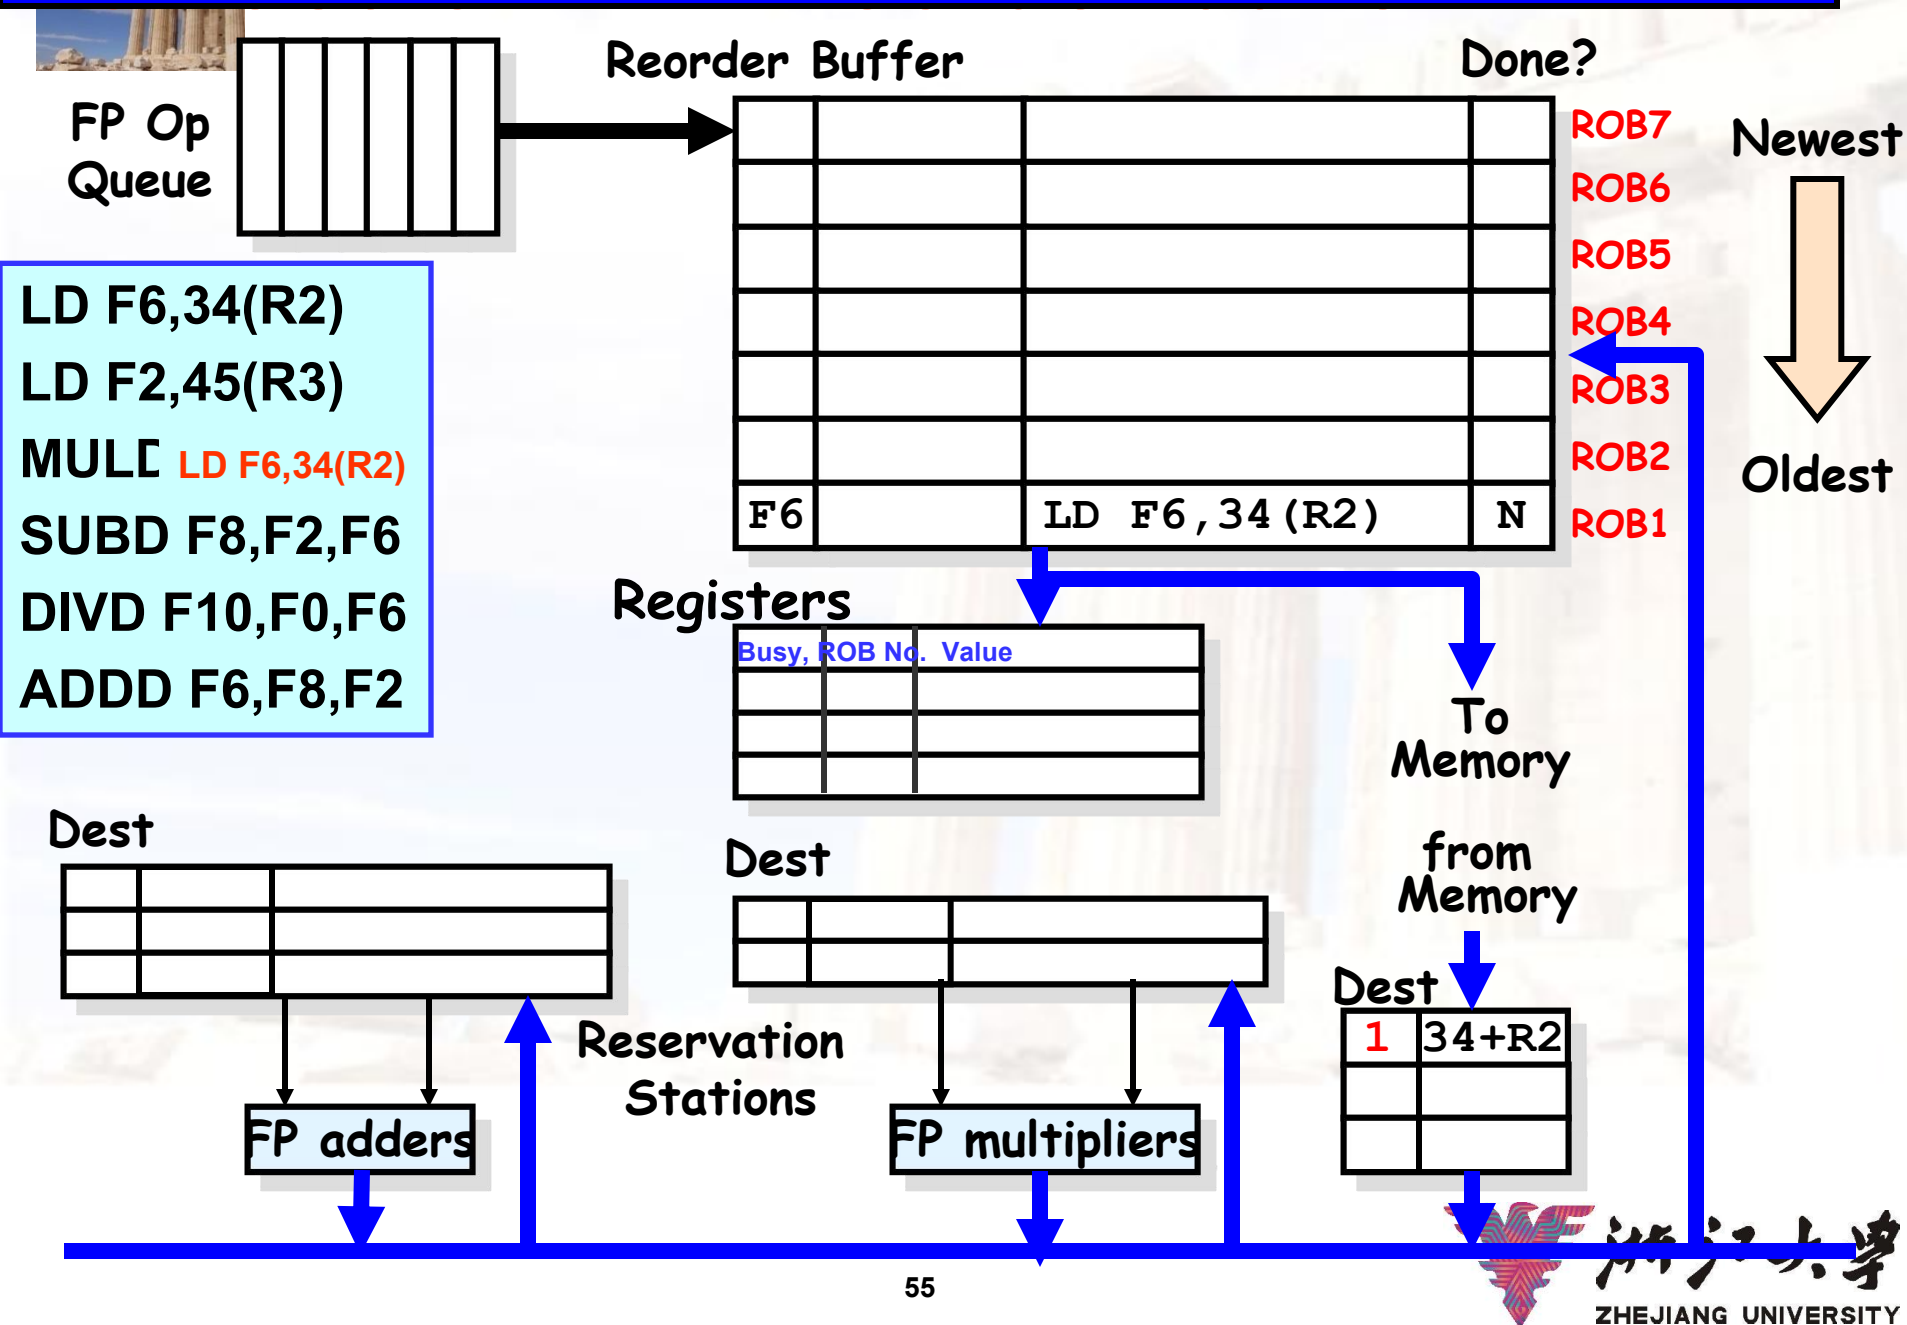
## ❑ For a mispredicted branch

- ROB is flushed and execution is restarted at the correct successor of the branch

## ❑ For a right-predicted branch

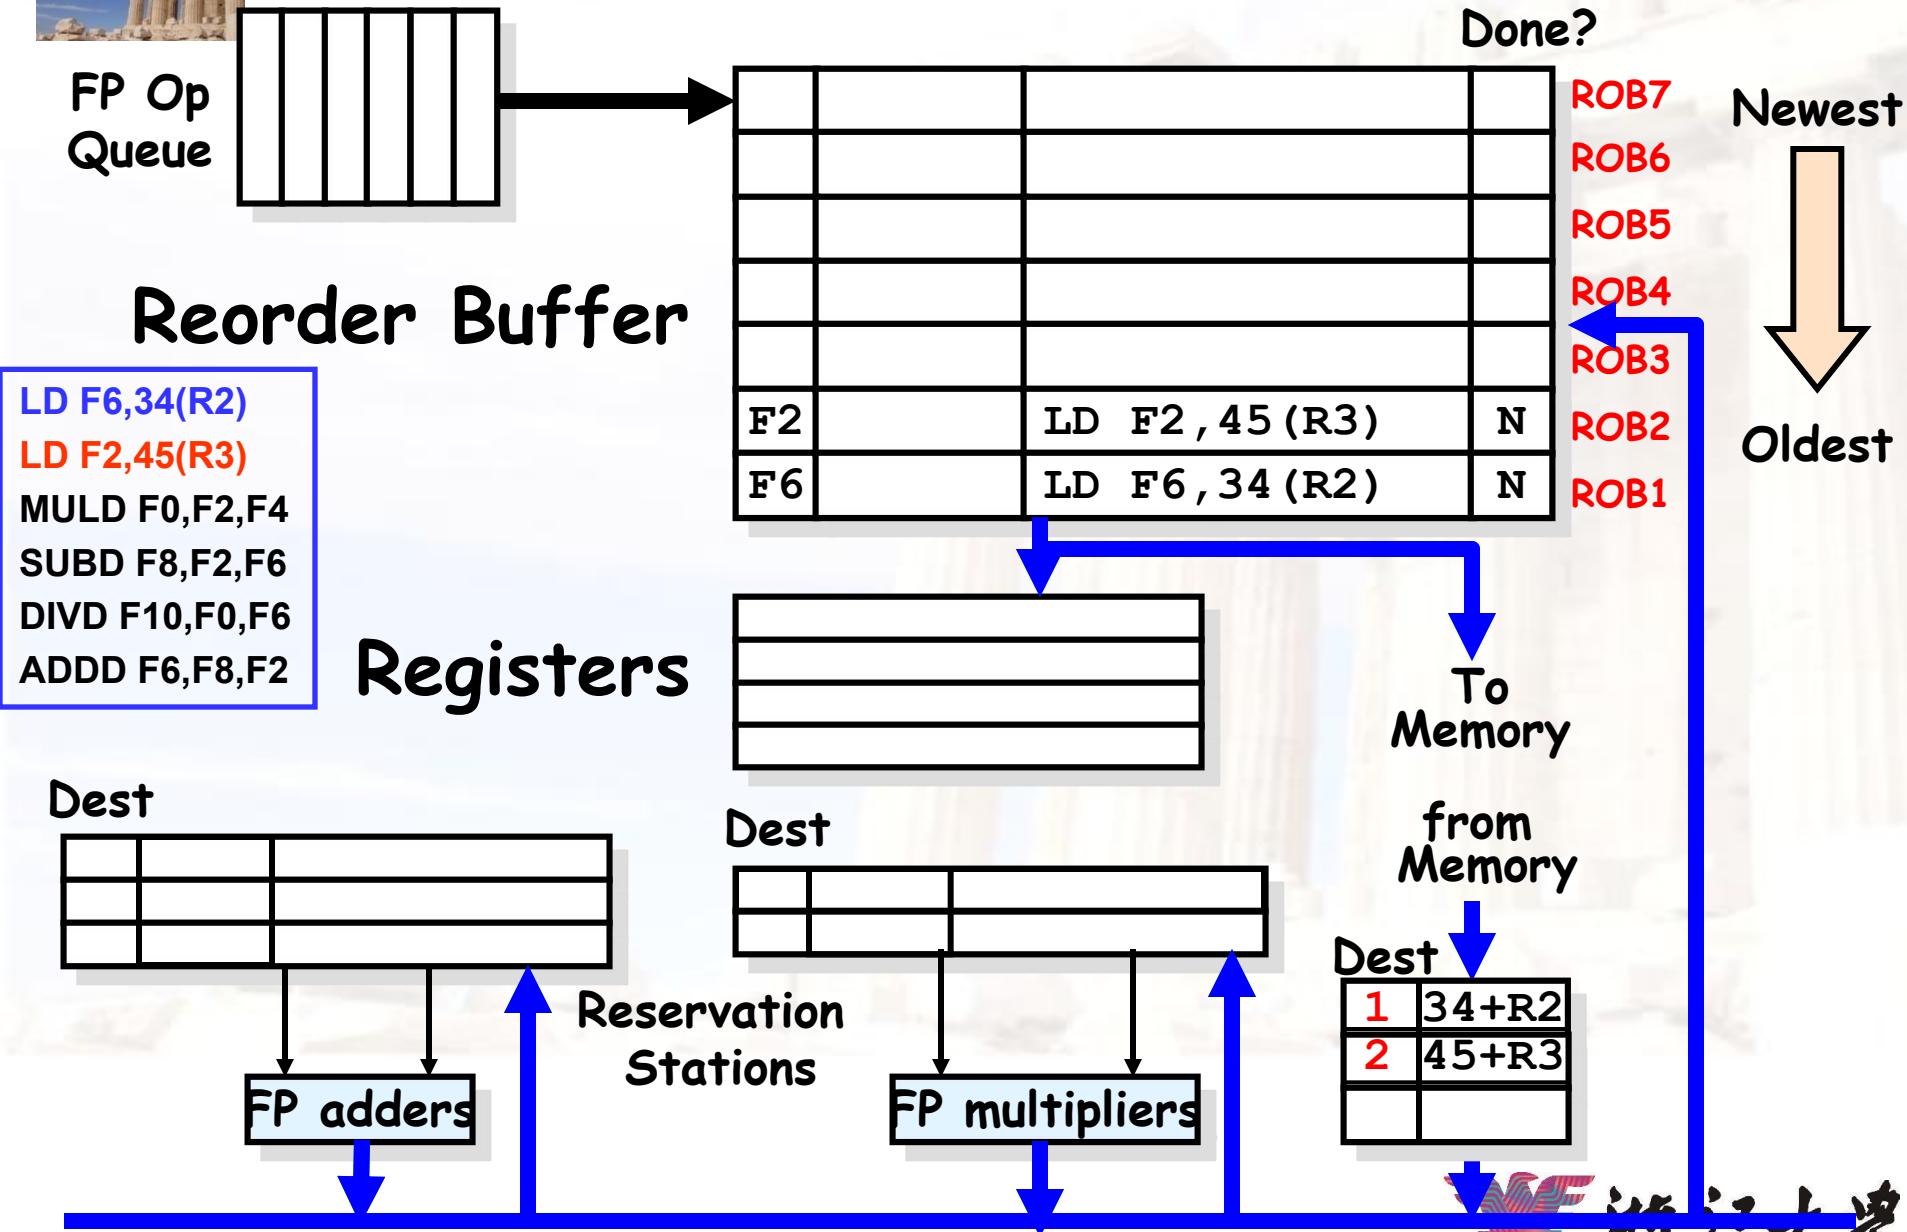
- End of branch instruction.

# 2 clock cycles for pipelined memory access






# Tomasulo With Reorder buffer:





# Tomasulo With Reorder buffer:

FP Op Queue



Reorder Buffer


			Done?
F0		MULD F0, F2, F4	N
F2		LD F2, 45 (R3)	N
F6		LD F6, 34 (R2)	N

ROB7  
ROB6  
ROB5  
ROB4  
ROB3  
ROB2  
ROB1

Newest  
↓  
Oldest

LD F6,34(R2)  
LD F2,45(R3)  
MULD F0,F2,F4  
SUBD F8,F2,F6  
DIVD F10,F0,F6  
ADD F6,F8,F2

Registers



To  
Memory

Dest


FP adders

Reservation  
Stations

Dest

3	MULD	ROB2, R(F4)

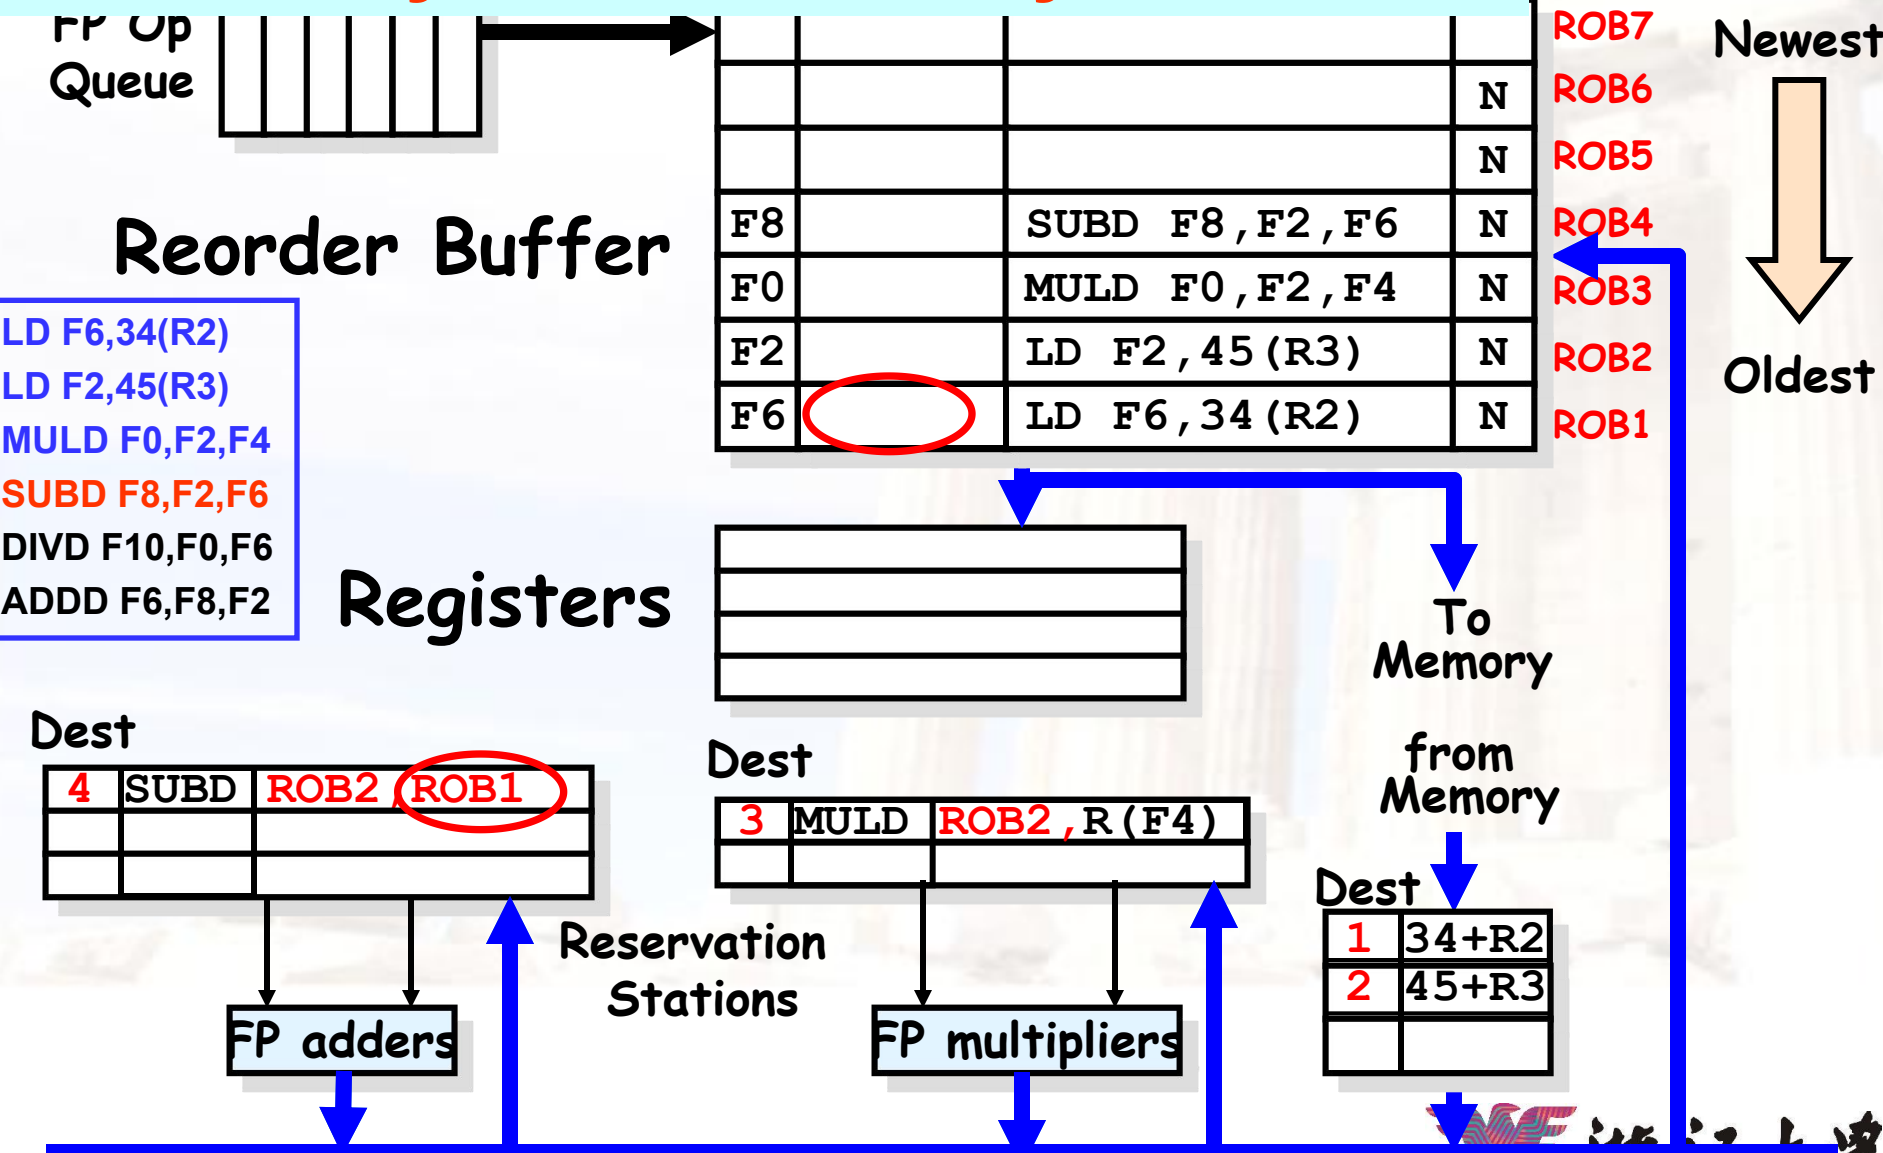
FP multipliers

from  
Memory

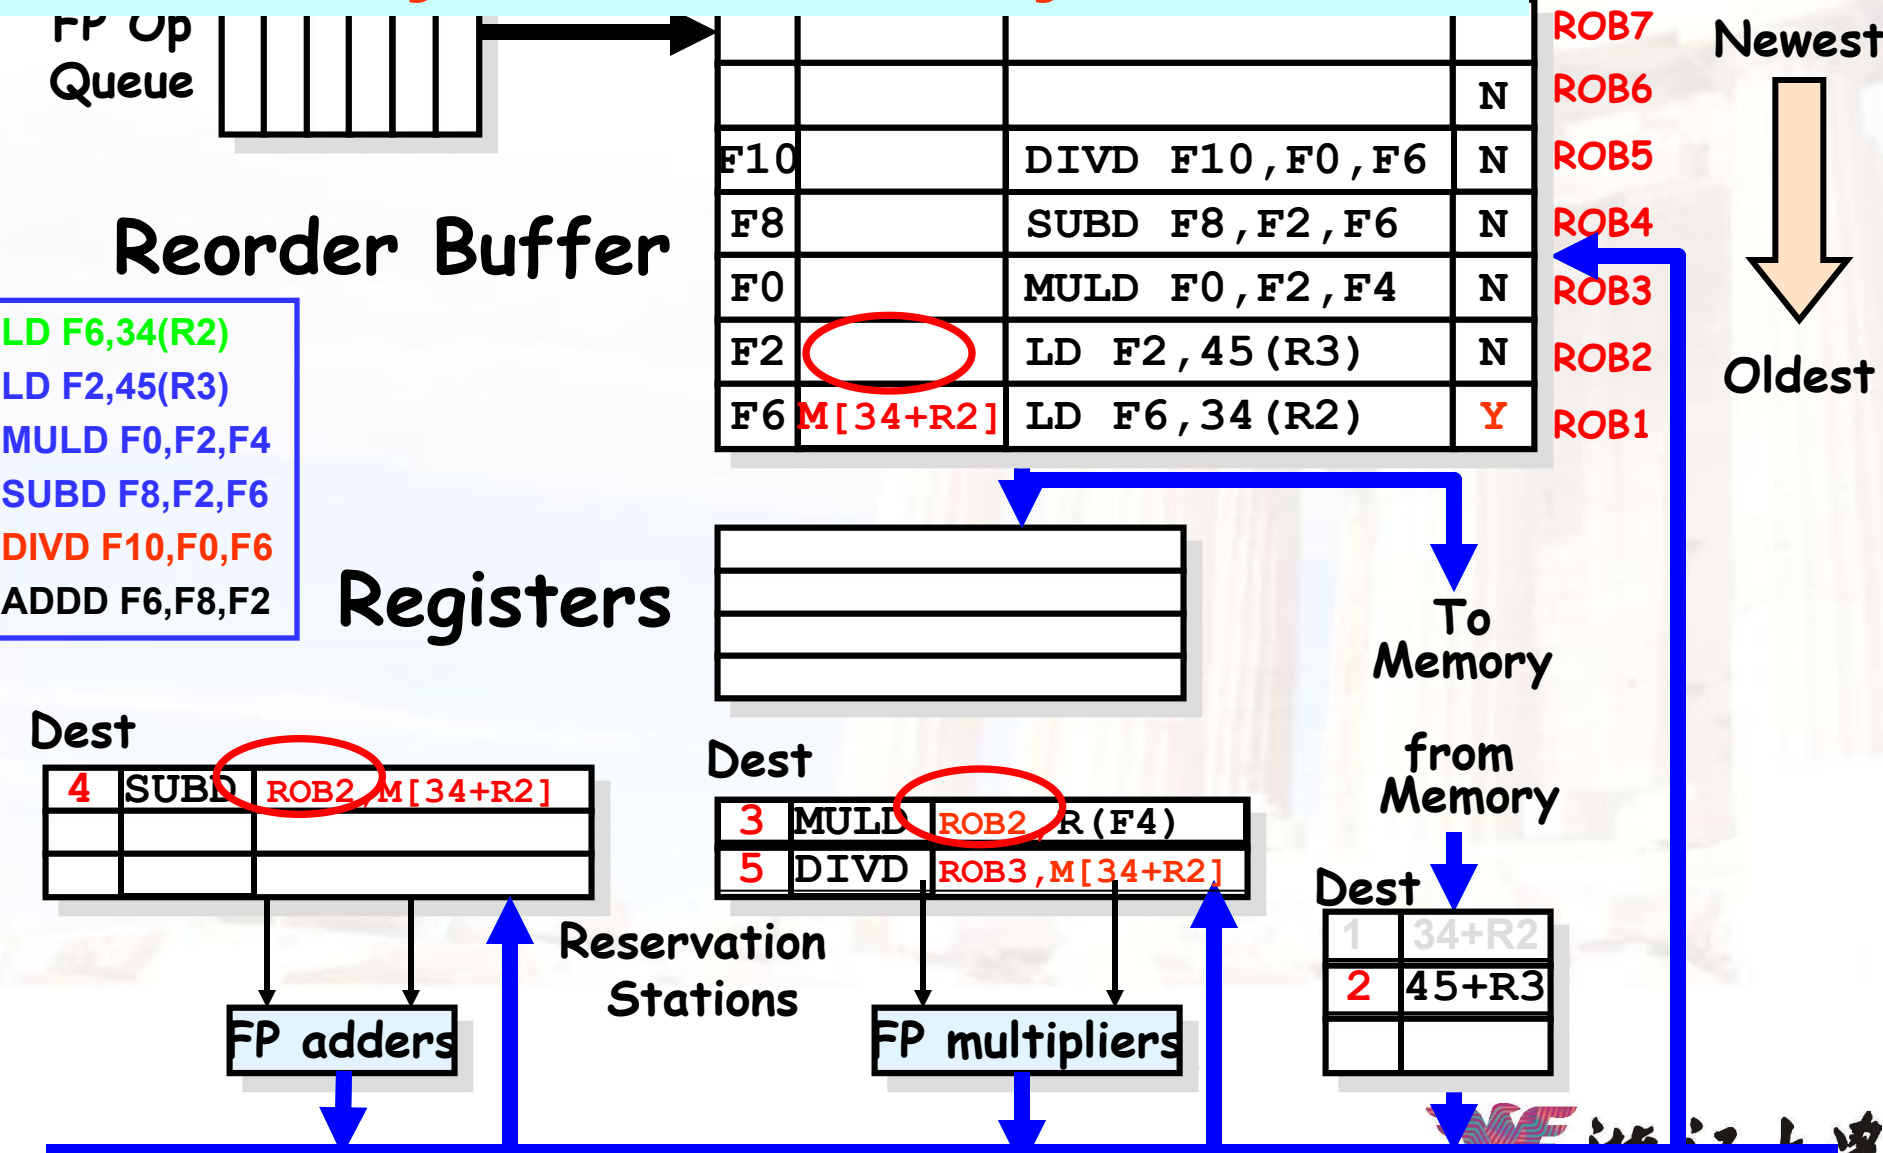
Dest

1	34+R2
2	45+R3

**Note: The first LD will finish  
memory access next cycle**

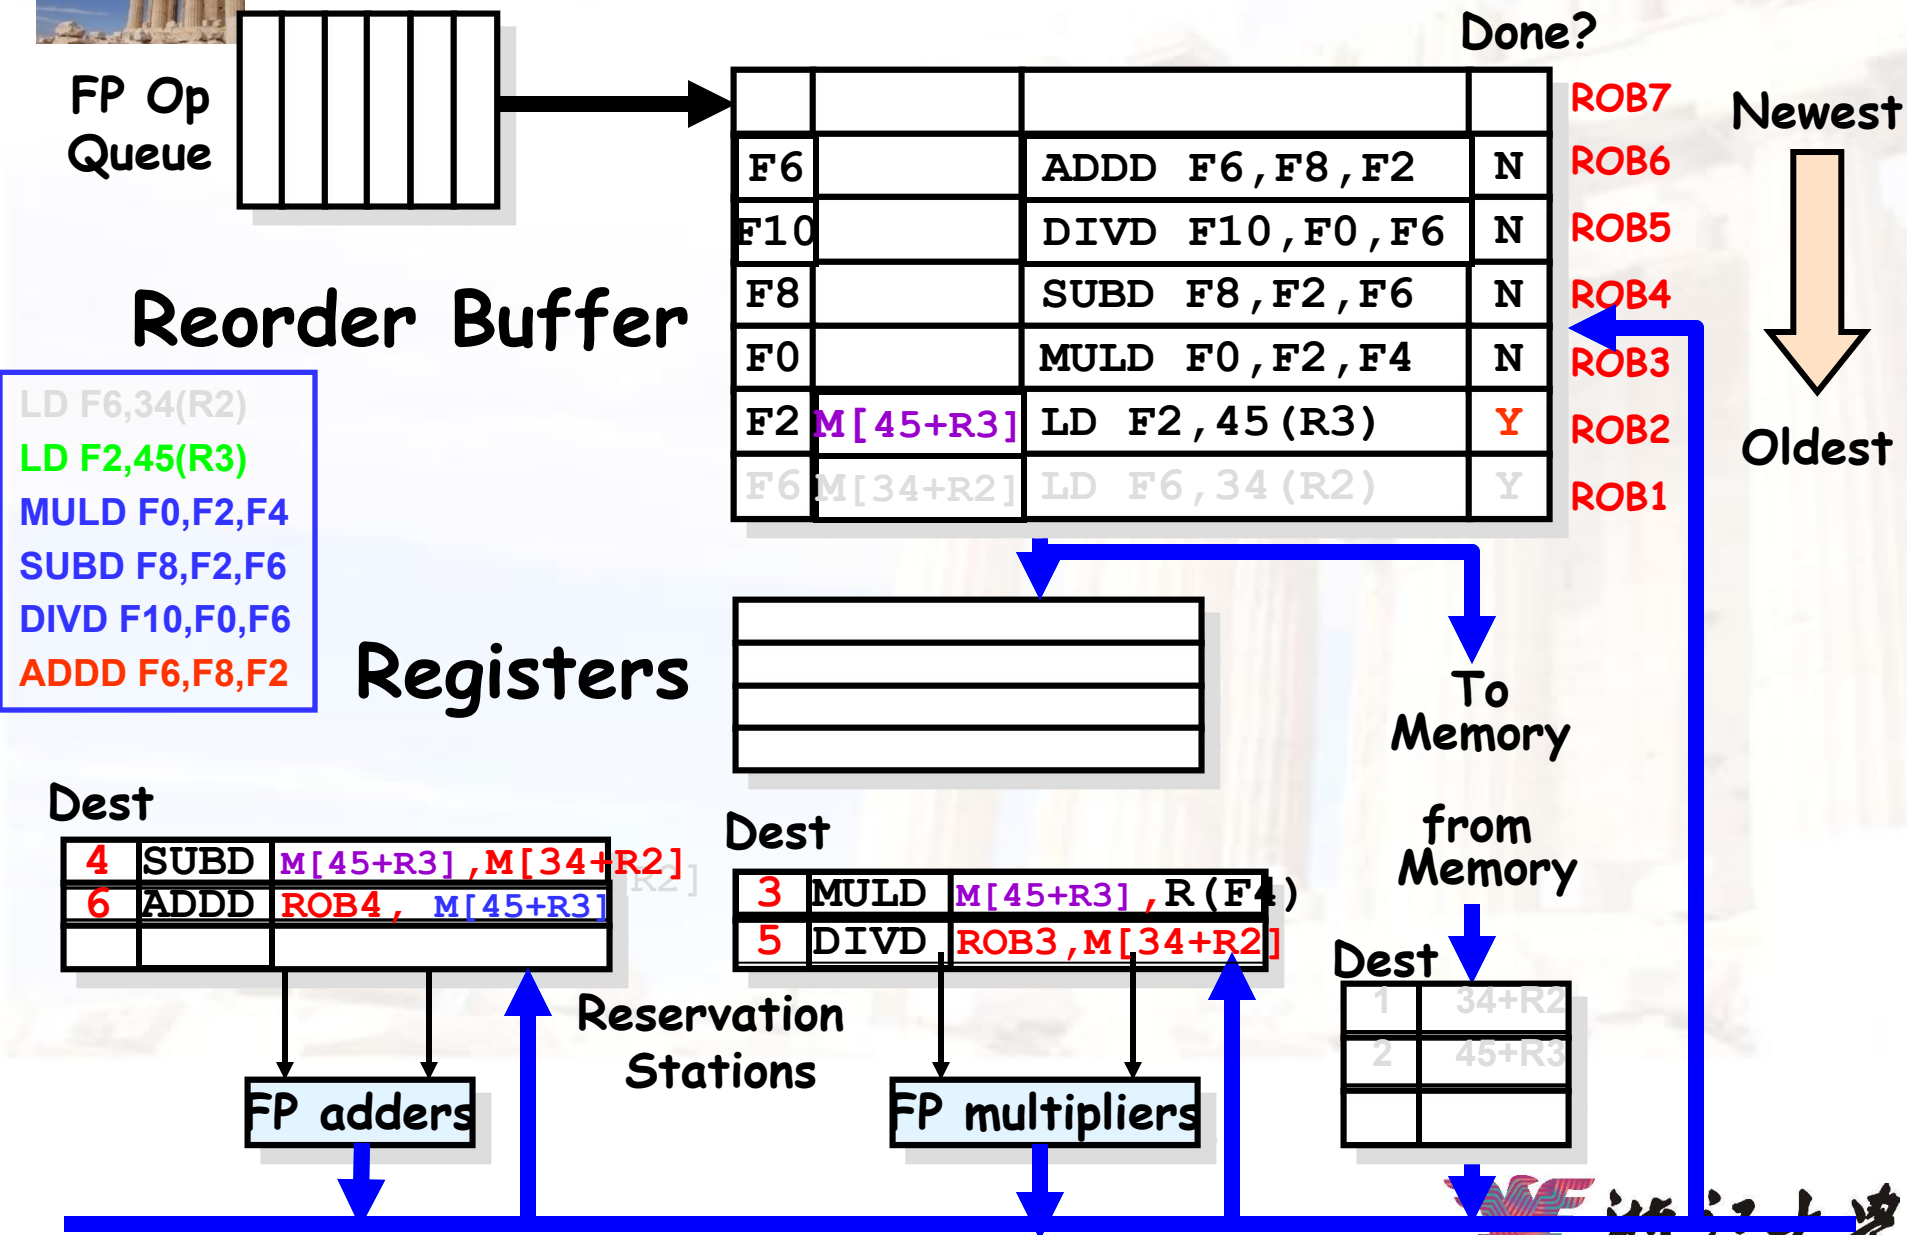


# Note: The 2nd LD will finish memory access next cycle





# Tomasulo With Reorder buffer:








# Tomasulo With Reorder buffer:

FP Op Queue




Reorder Buffer

			Done?
F6		ADDD F6, F8, F2	N
F10		DIVD F10, F0, F6	N
F8	V[1]	SUBD F8, F2, F6	Y
F0		MULD F0, F2, F4	N
F2	M[45+R3]	LD F2, 45 (R3)	Y
F6	M[34+R2]	LD F6, 34 (R2)	Y

ROB7  
ROB6  
ROB5  
ROB4  
ROB3  
ROB2  
ROB1


Newest



Oldest

LD F6,34(R2)  
LD F2,45(R3)  
MULD F0,F2,F4  
SUBD F8,F2,F6  
DIVD F10,F0,F6  
ADDD F6,F8,F2

Registers



To  
Memory

Dest

6	ADDD	V[1], M[45+R3]

FP adders

Reservation  
Stations

Dest

3	MULD	M[45+R3], R(F4)
5	DIVD	ROB3, M[34+R2]

FP multipliers

from  
Memory

Dest



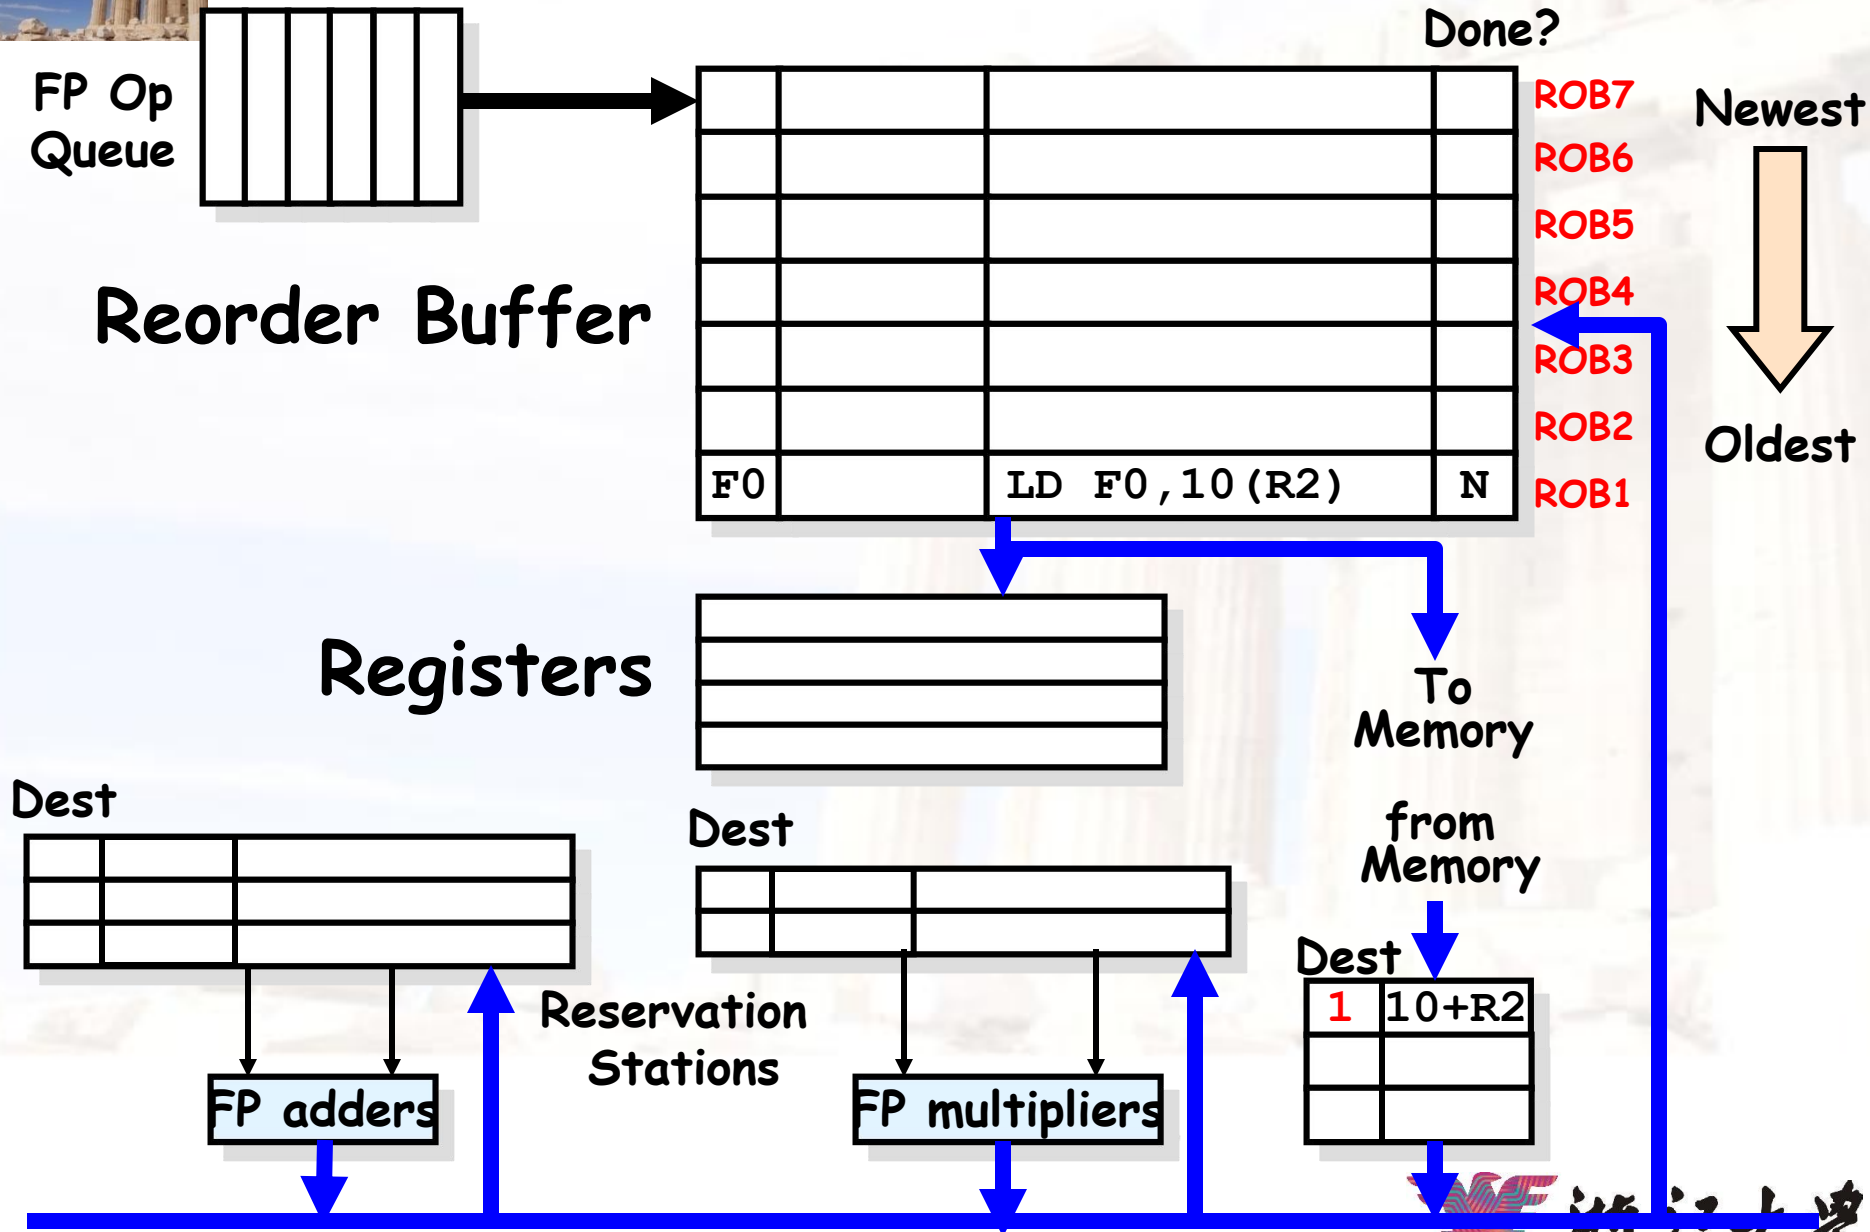

# Speculation performance

❑ Overlapped basic block

❑ Is there any new problems ?

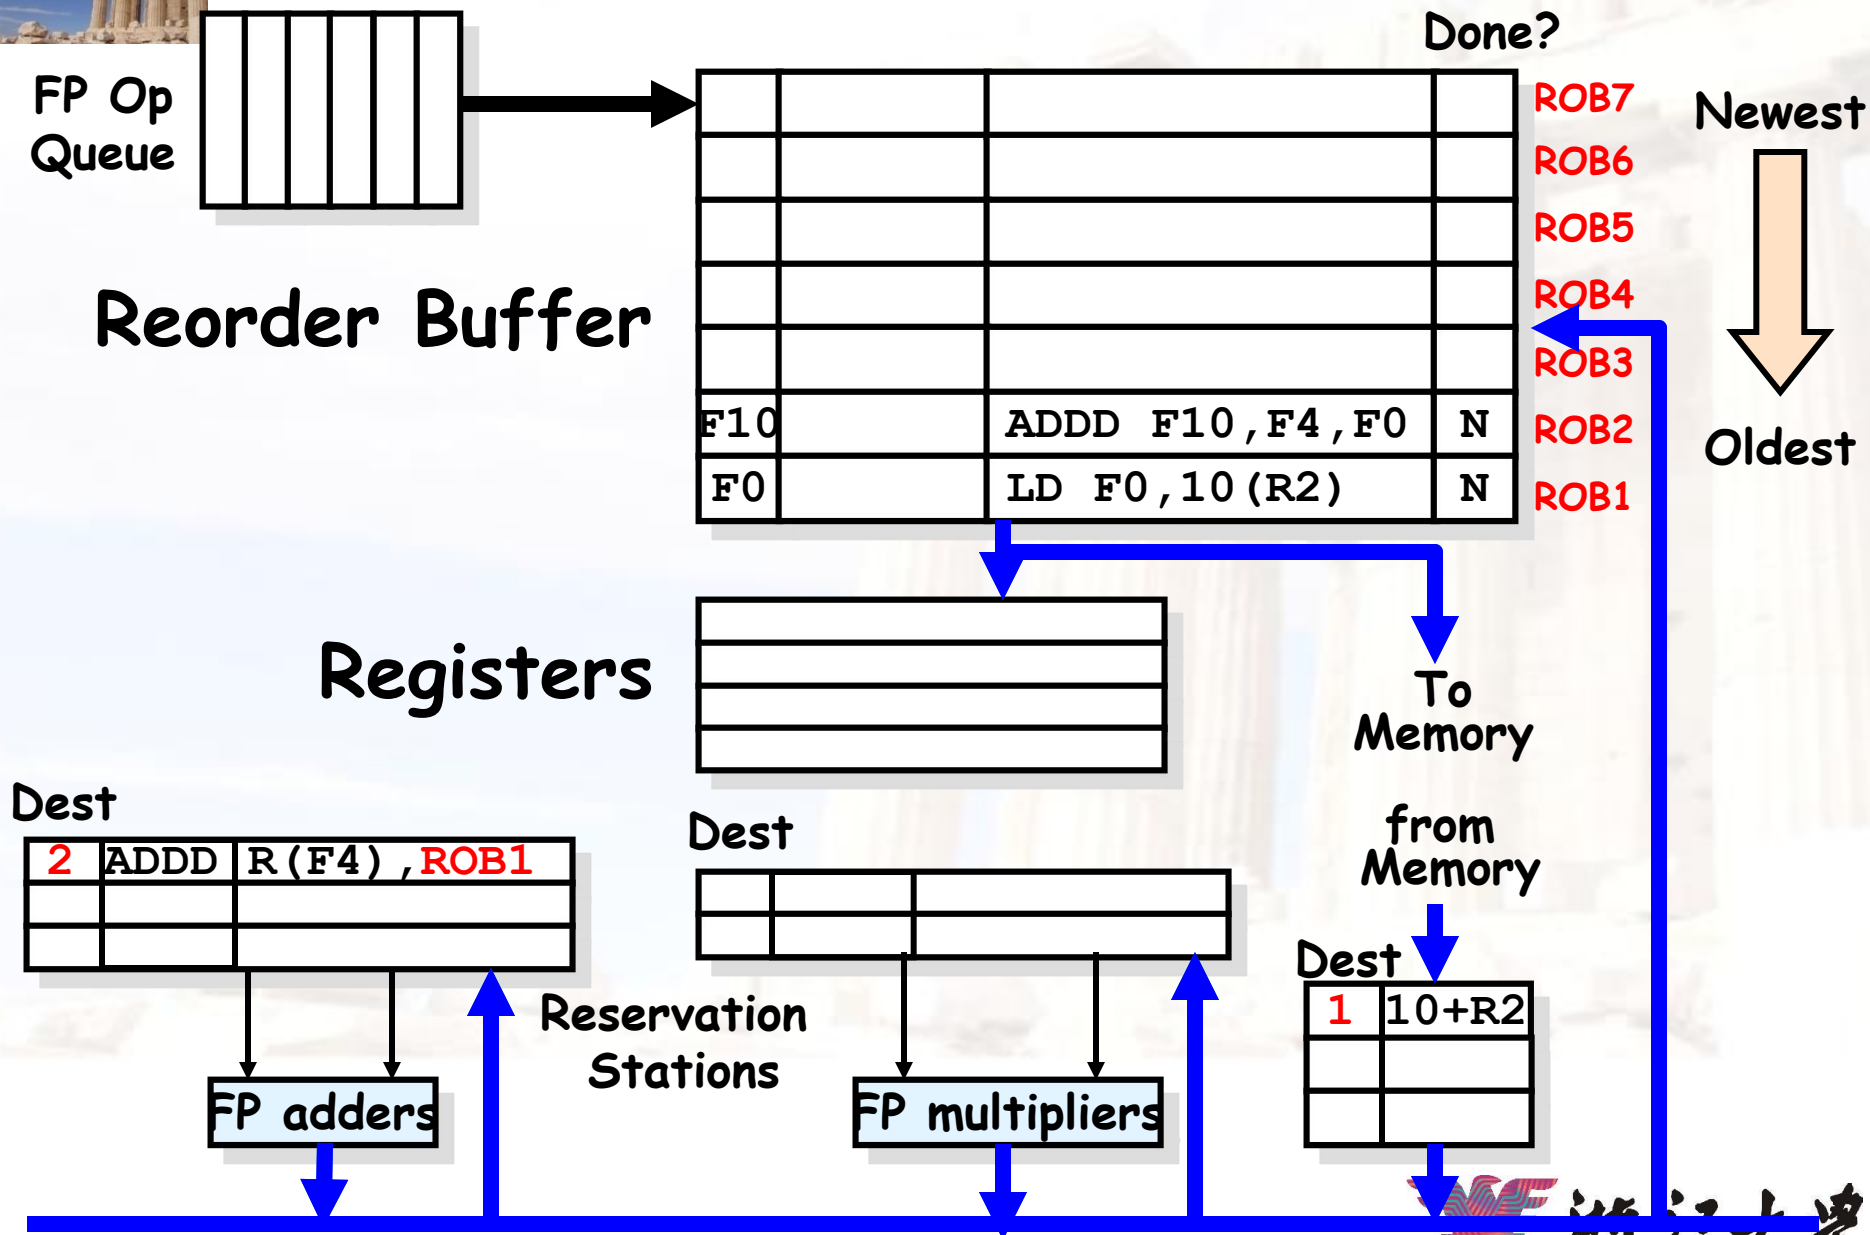
➤ Pay attention to the memory accessing

# Tomasulo With Reorder buffer:



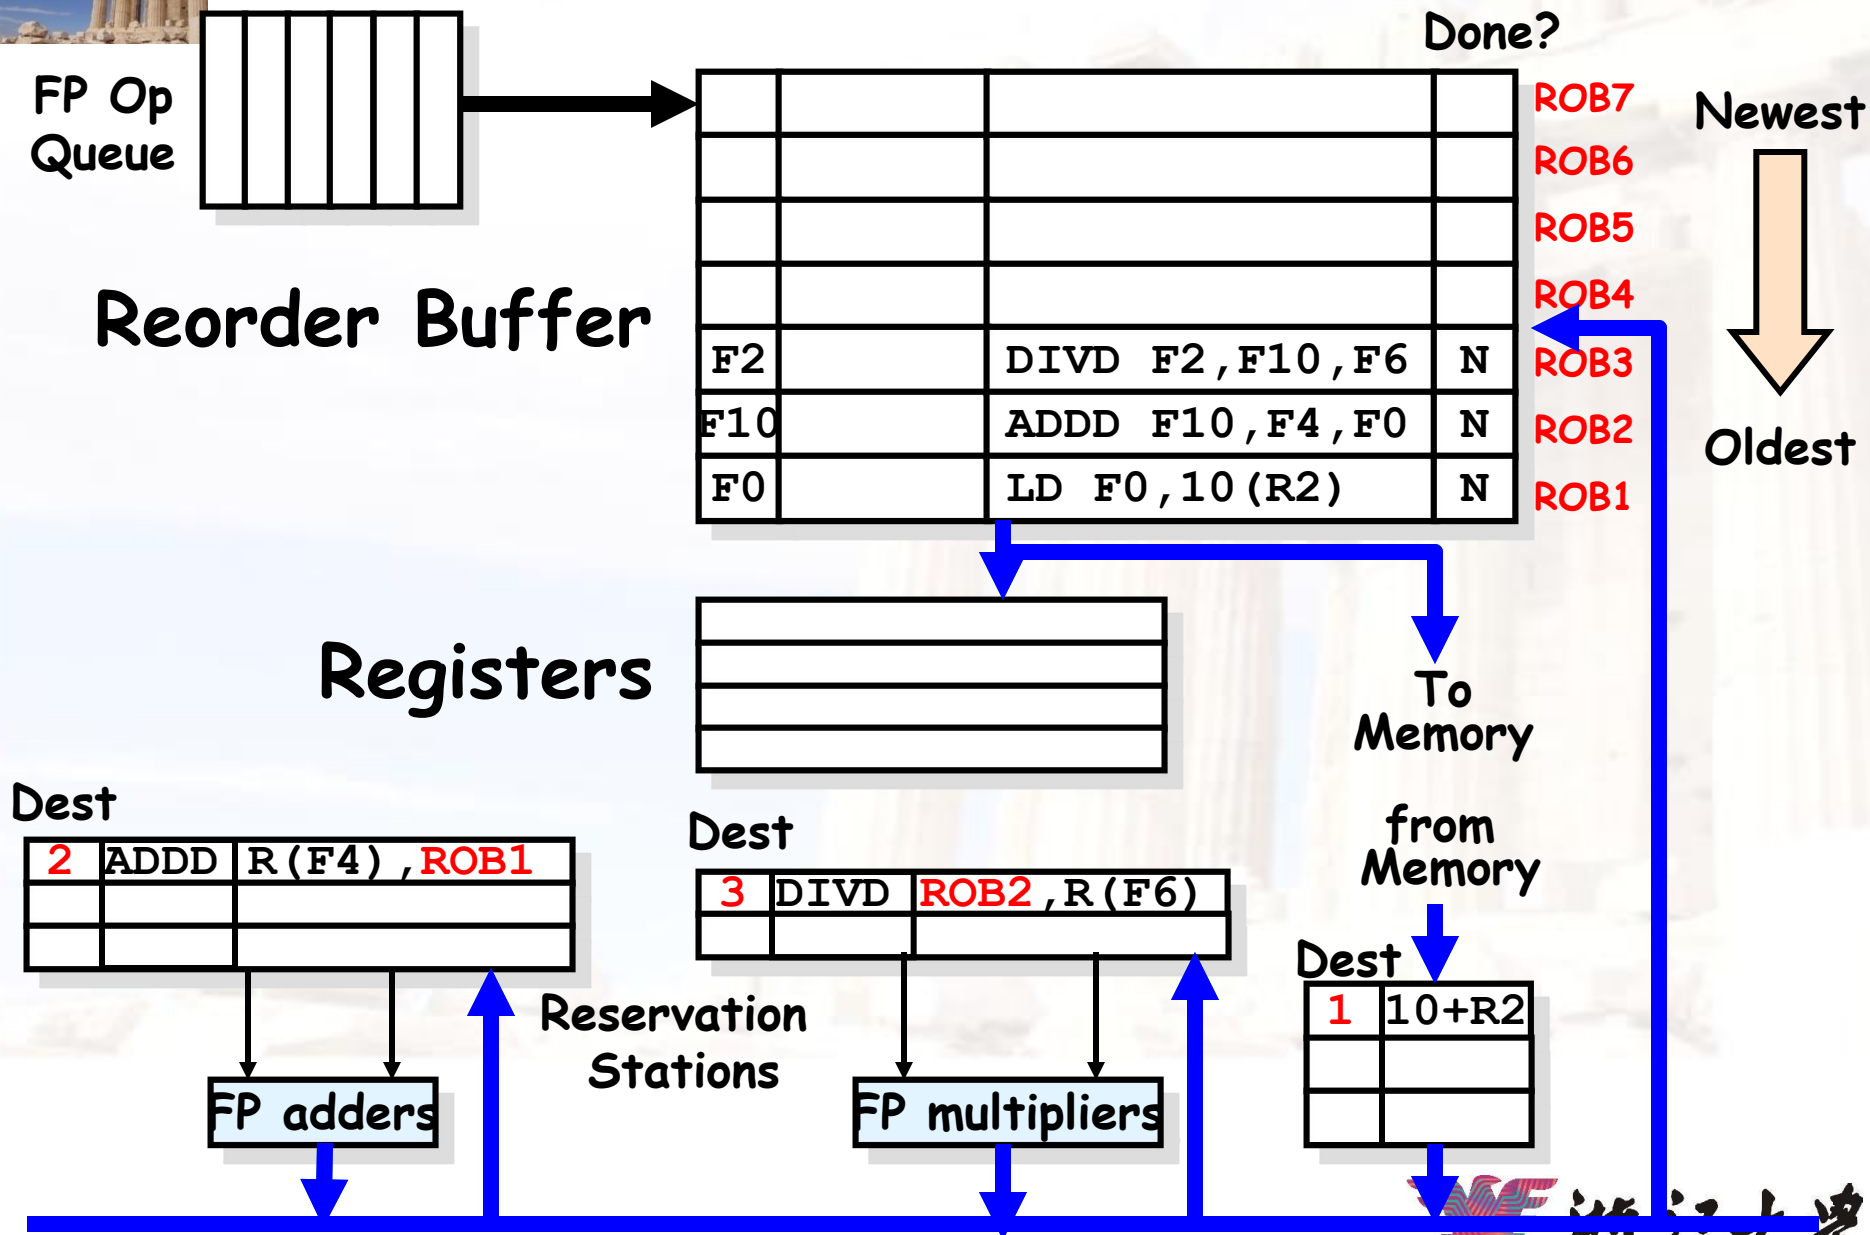


# Tomasulo With Reorder buffer:



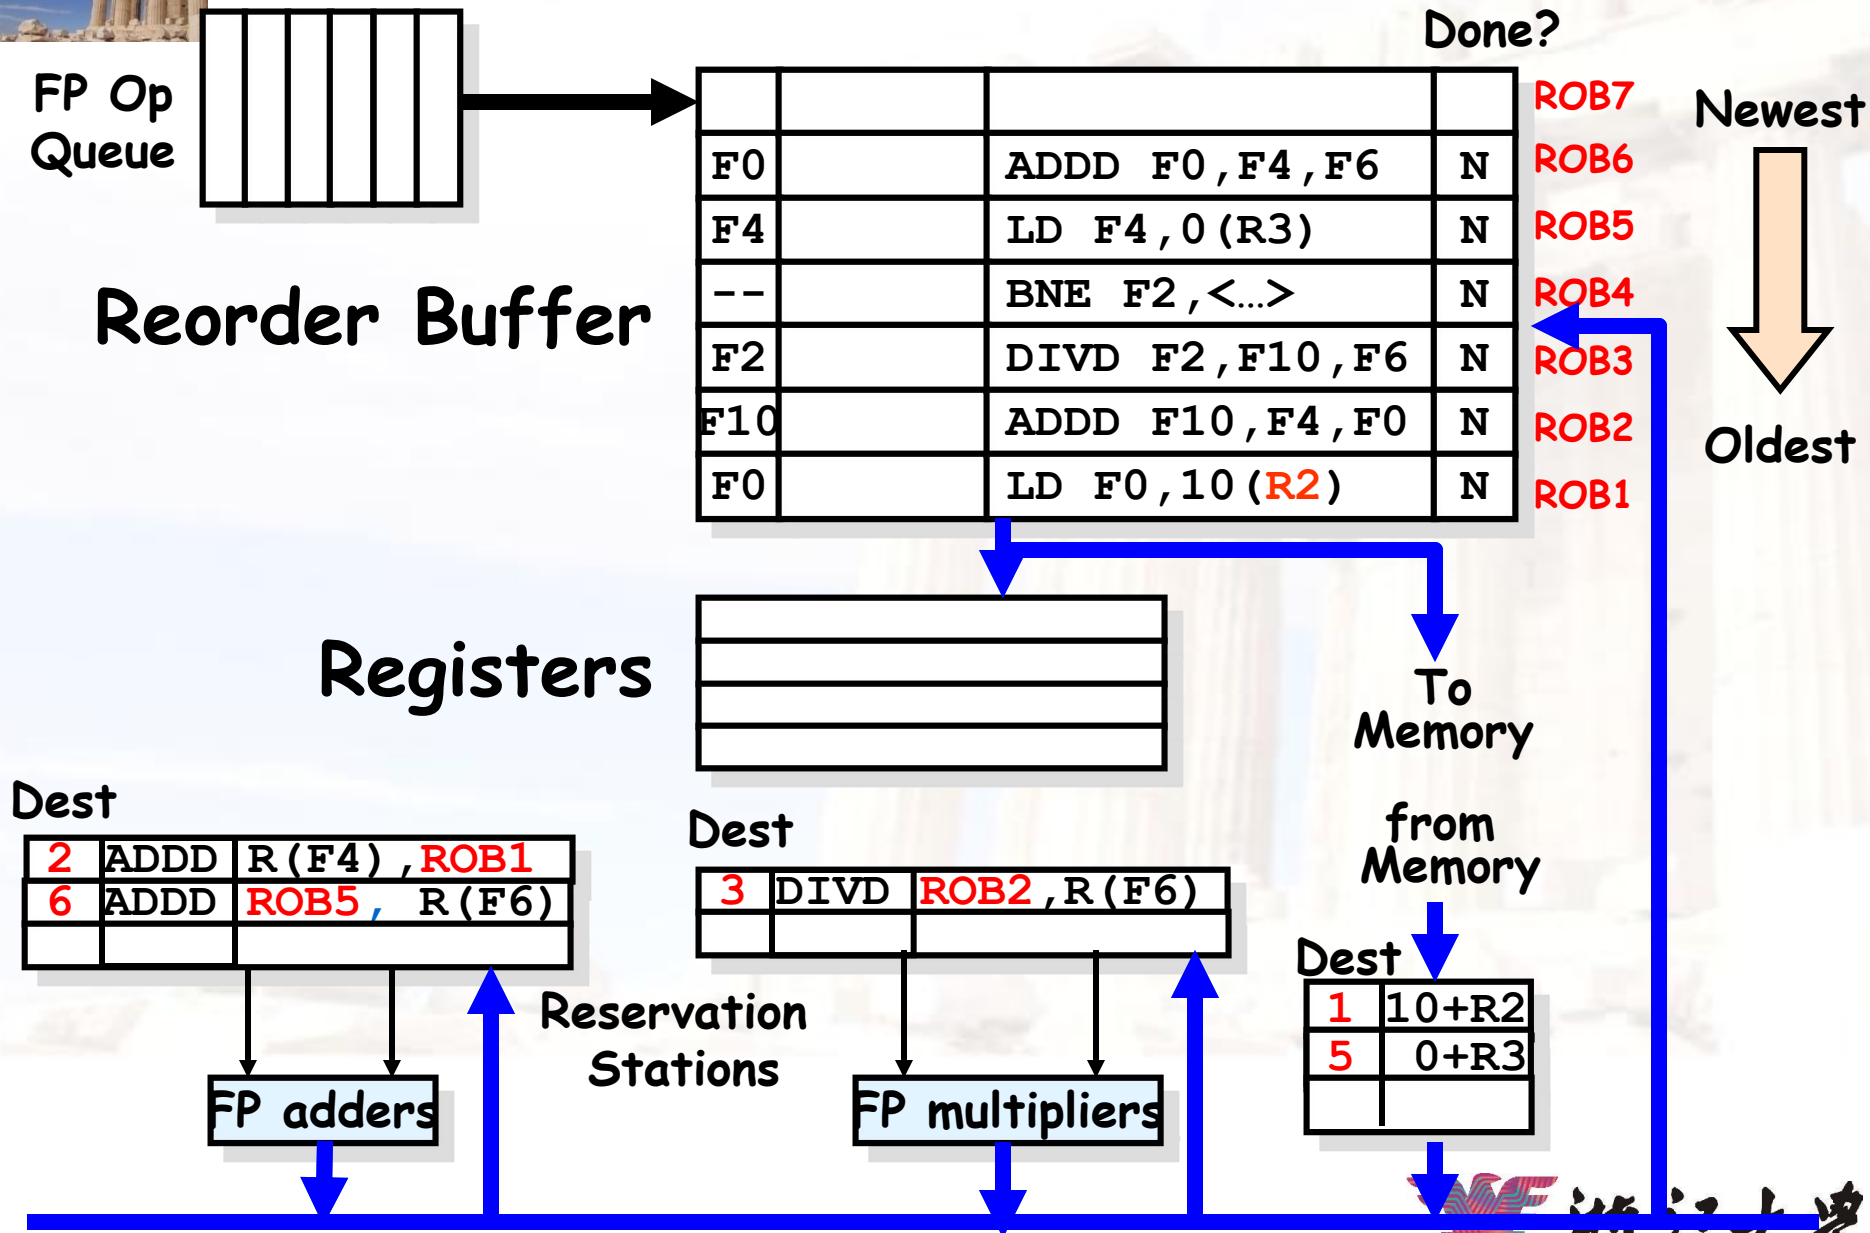


# Tomasulo With Reorder buffer:



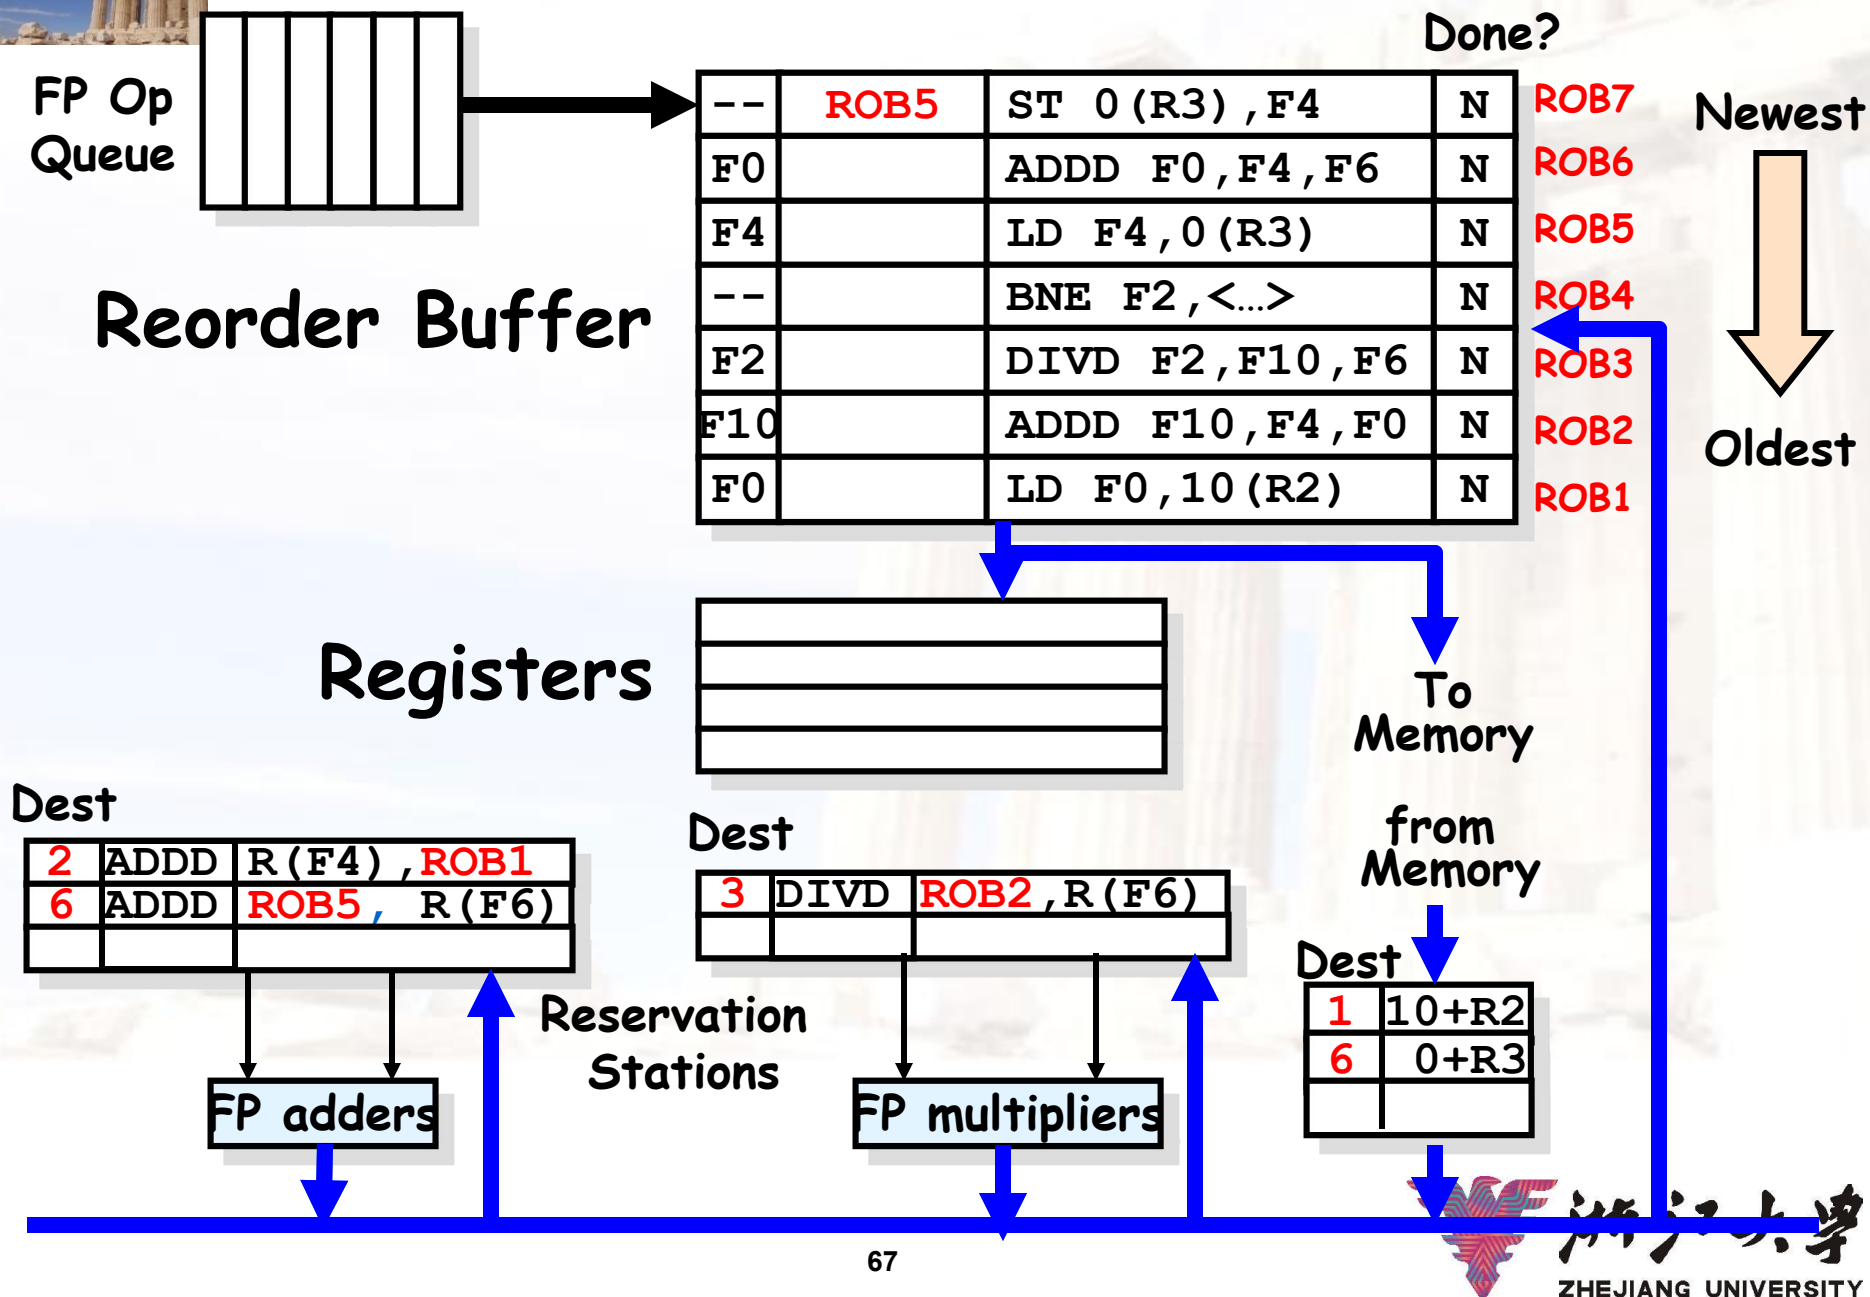


# Tomasulo With Reorder buffer:





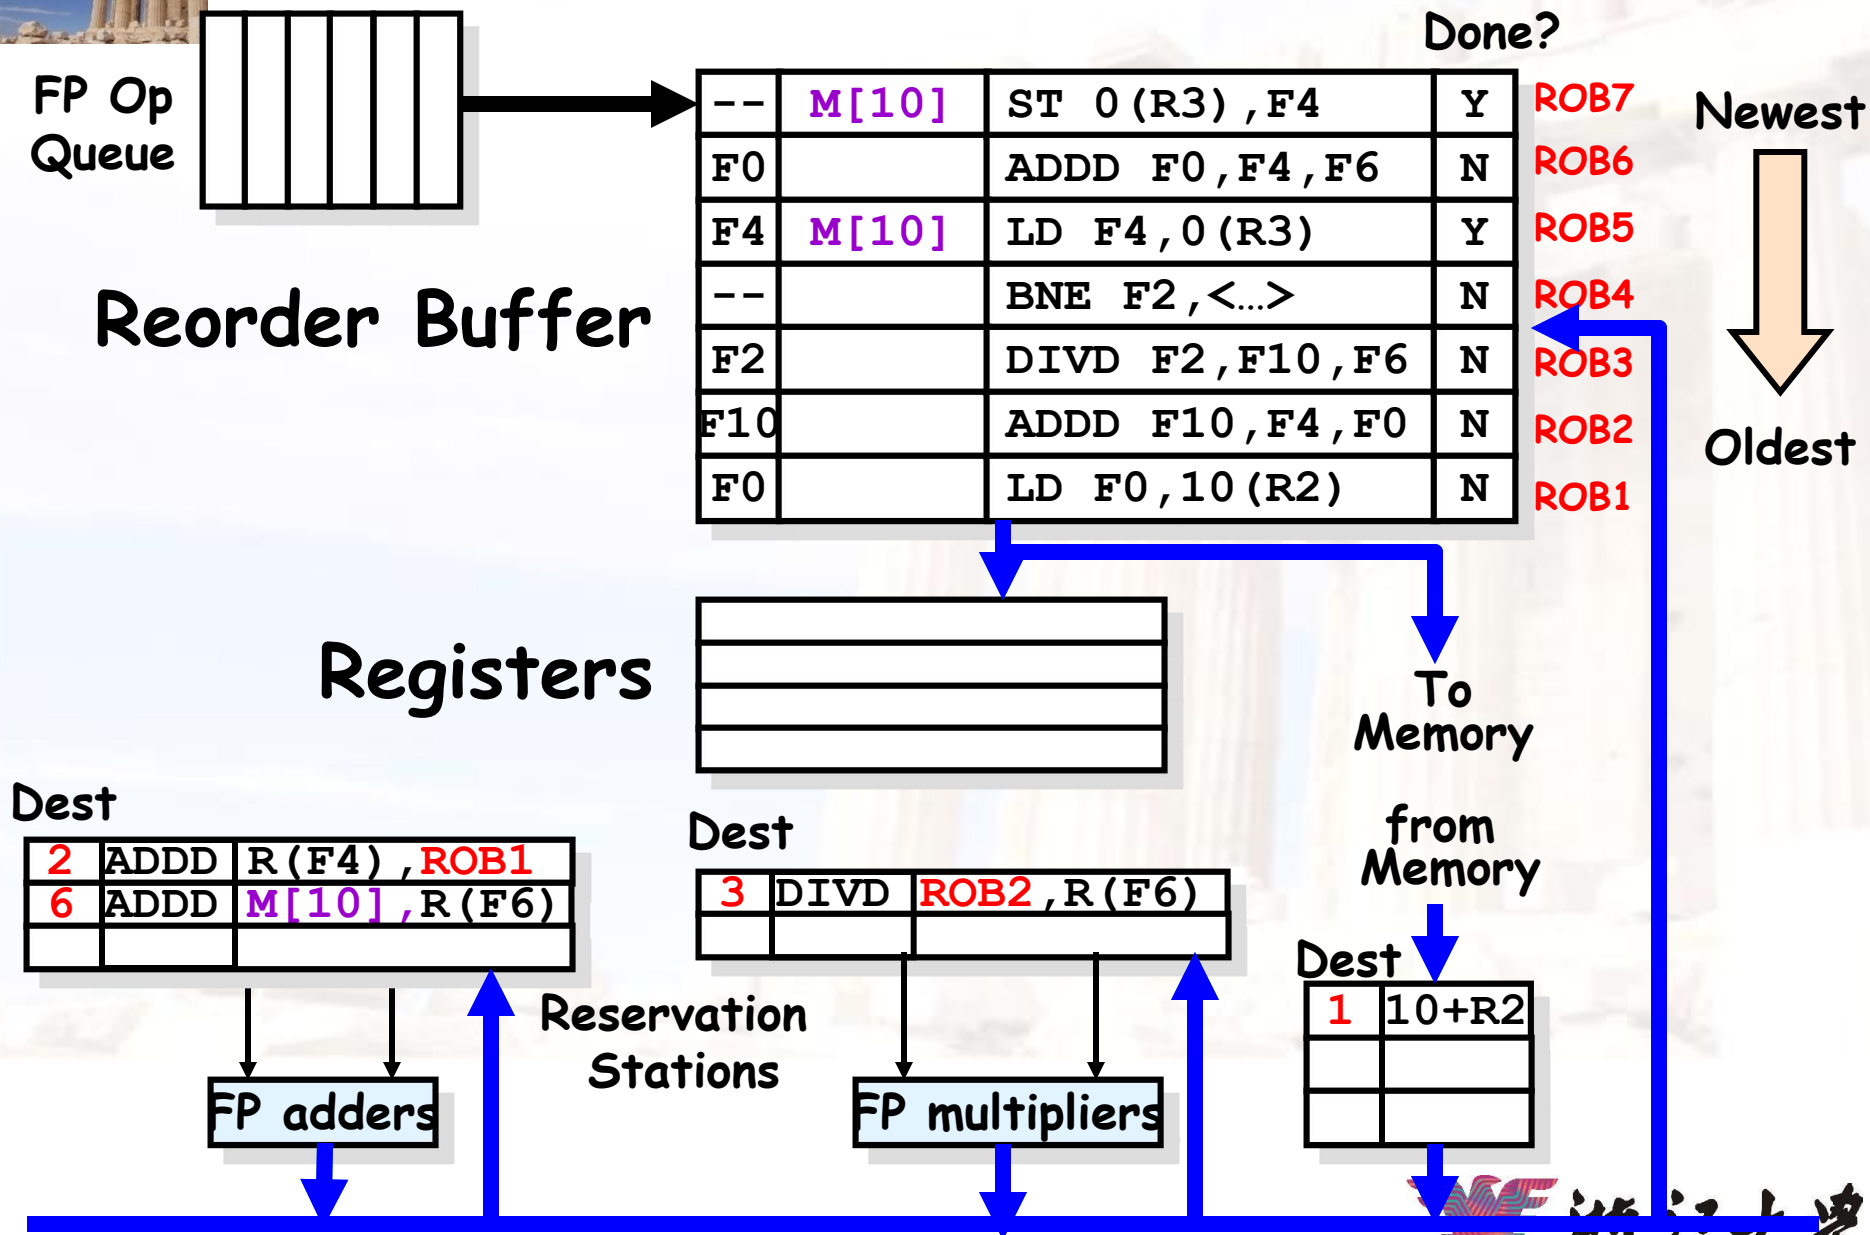
# Tomasulo With Reorder buffer:





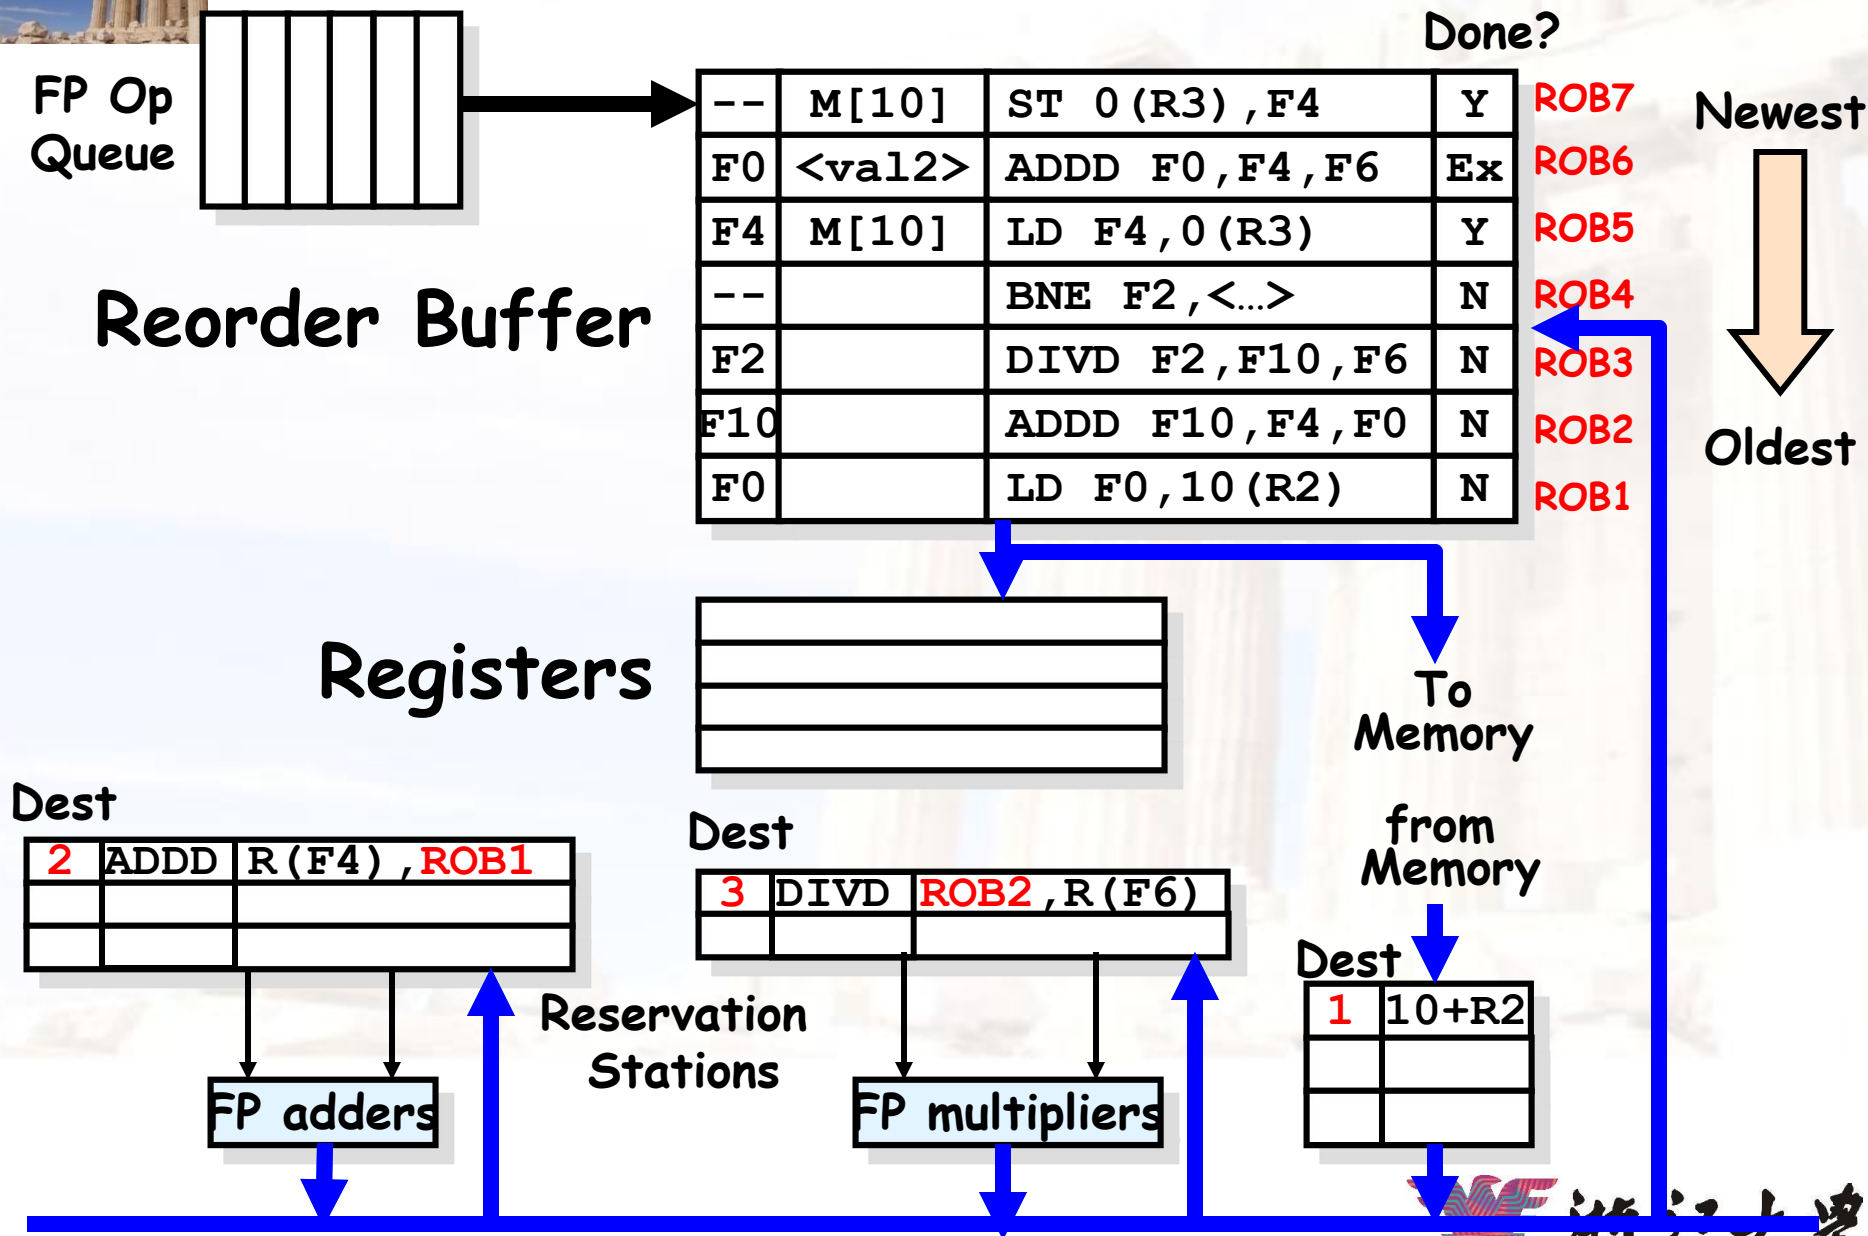


# Tomasulo With Reorder buffer:





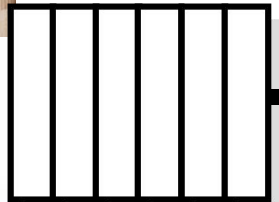
# Tomasulo With Reorder buffer:





# Tomasulo With Reorder buffer:

FP Op Queue



Reorder Buffer

What about memory hazards???


Registers

--	M[10]	ST 0 (R3) , F4	Y
F0	<val2>	ADDD F0 , F4 , F6	Ex
F4	M[10]	LD F4 , 0 (R3)	Y
--		BNE F2 , <...>	N
F2		DIVD F2 , F10 , F6	N
F10		ADDD F10 , F4 , F0	N
F0		LD F0 , 10 (R2)	N

Done?

ROB7  
ROB6  
ROB5  
ROB4  
ROB3  
ROB2  
ROB1

Newest  
↓  
Oldest



Dest

2	ADDD	R (F4) , ROB1

FP adders

Dest

3	DIVD	ROB2 , R (F6)

FP multipliers

Reservation Stations

To Memory  
from Memory

Dest

1	10+P2



# new problem--memory disambiguation

❑ Question: Given a load that follows a store in program order, are the two related?

➤ (Alternatively: is there a RAW hazard between the store and the load)?

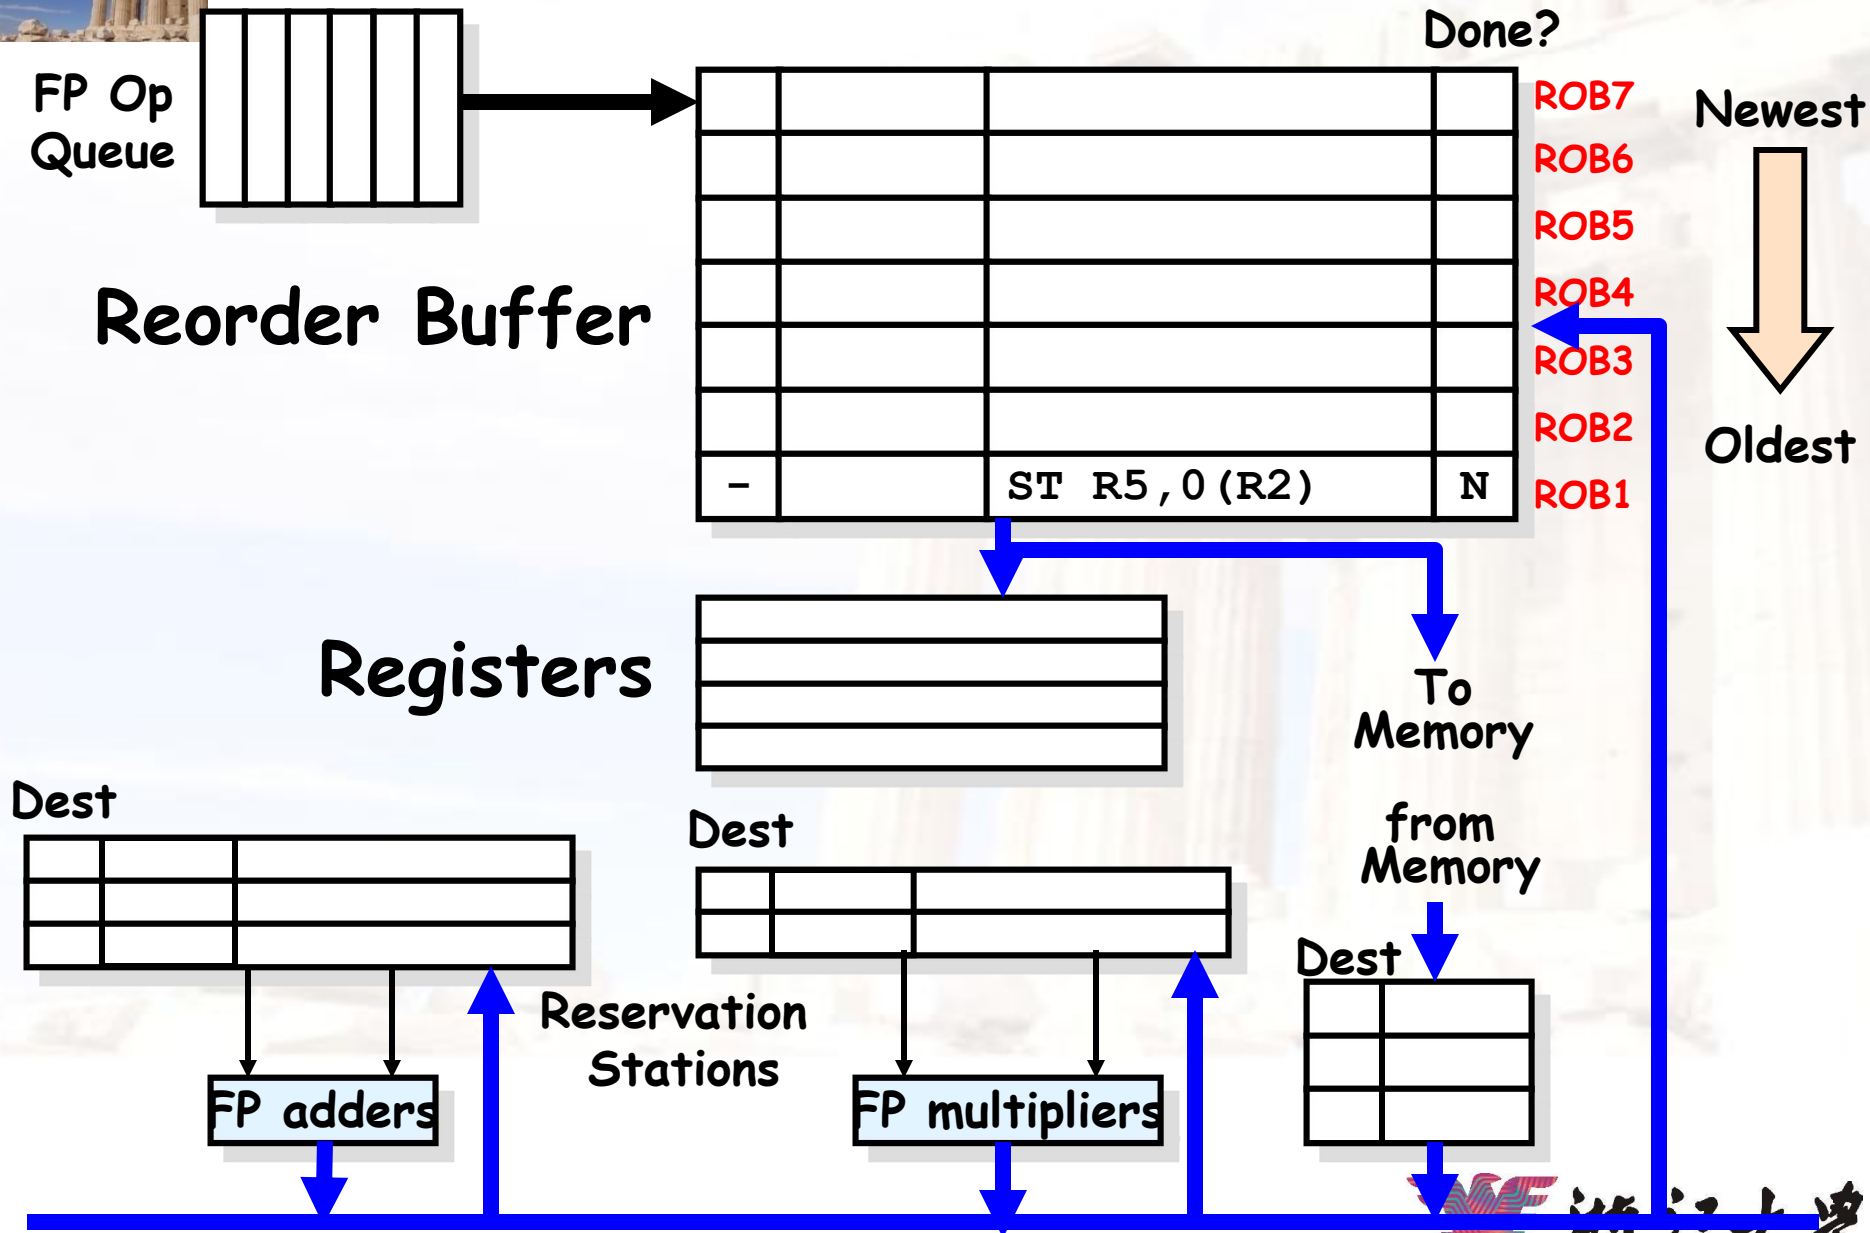
Eg:           ST     O(R2), R5  
              LD     R4,    O(R3)

❑ Can we go ahead and start the load early?

➤ Store address could be delayed for a long time by some calculation that leads to R2 (divide?).

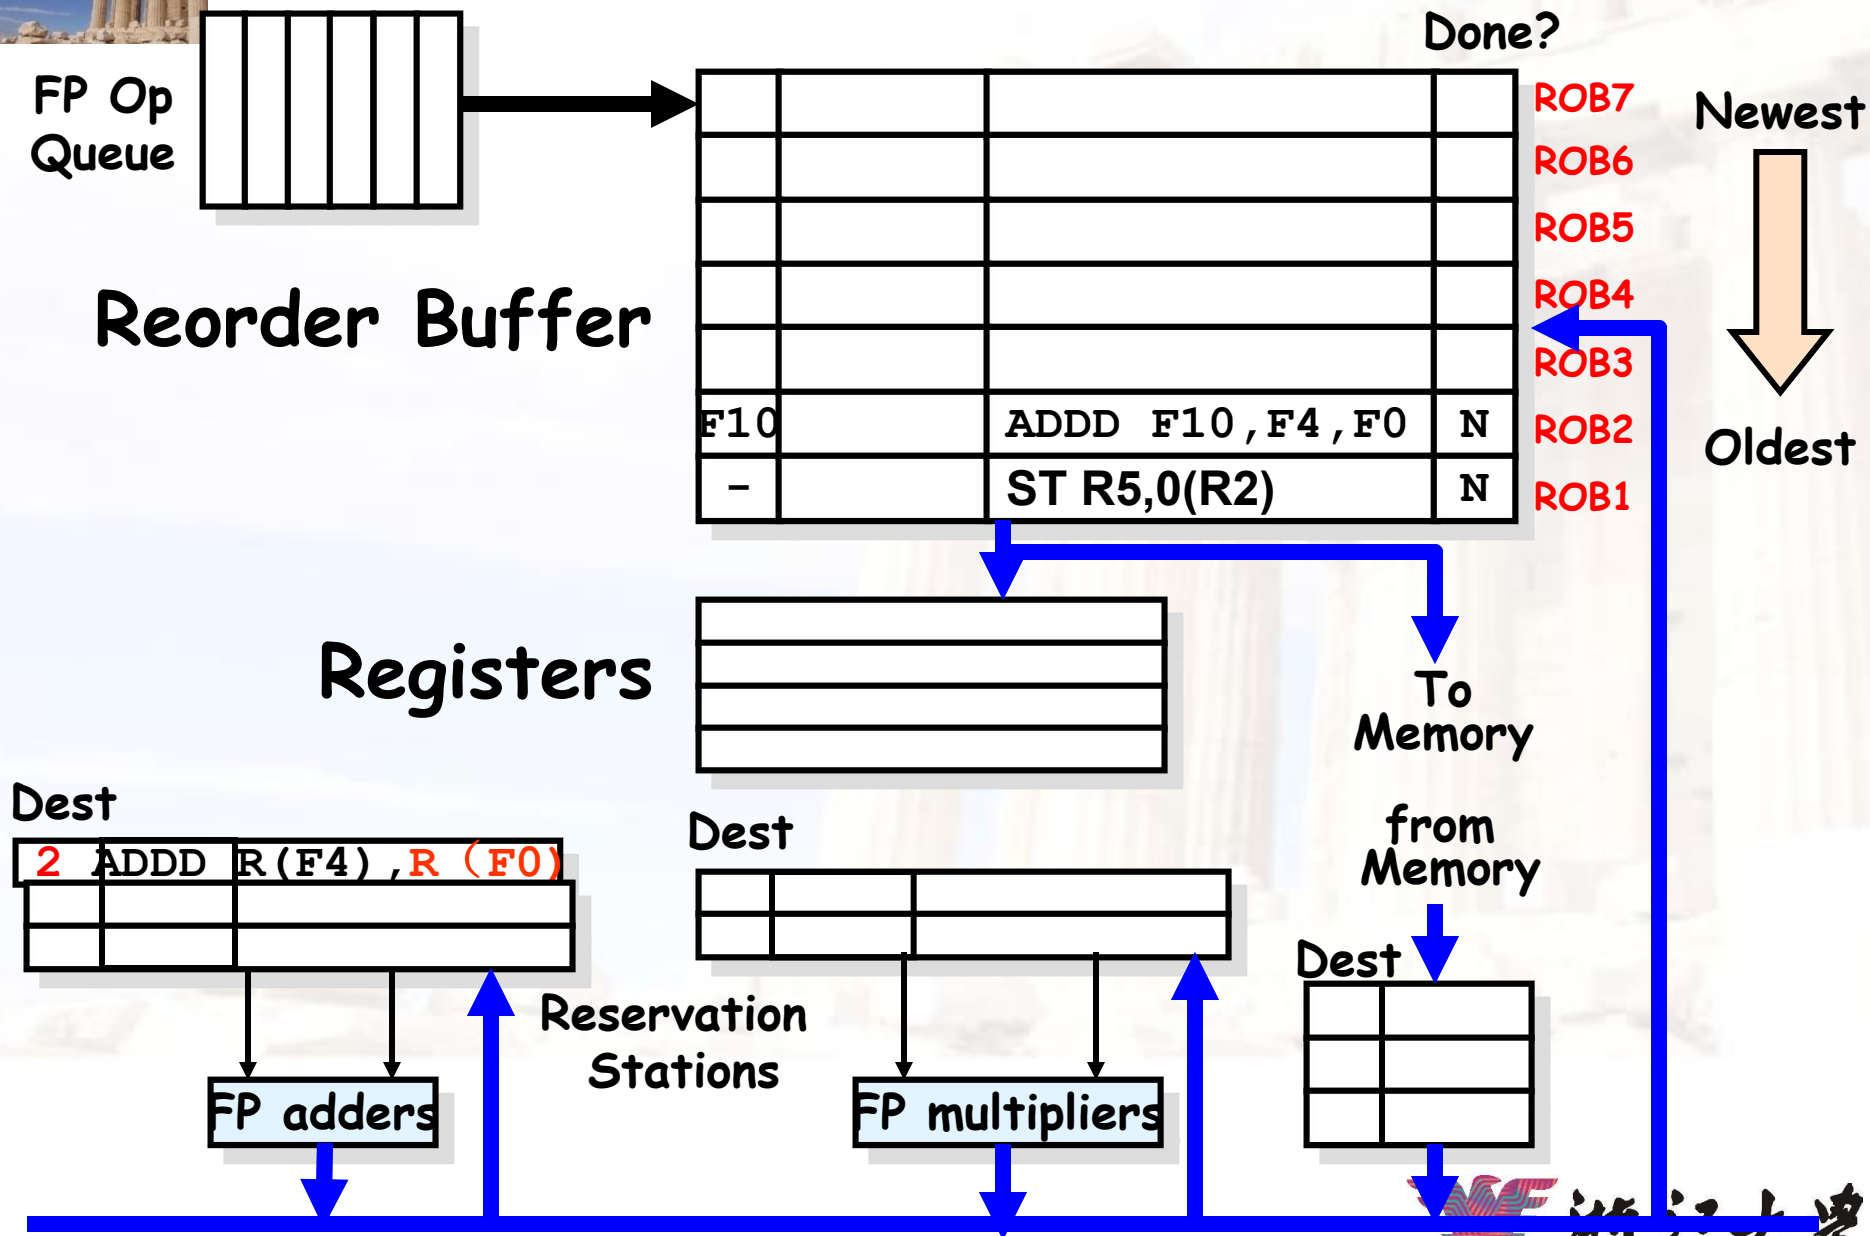
➤ We might want to issue/begin execution of both operations in same cycle.

# Tomasulo With Reorder buffer:



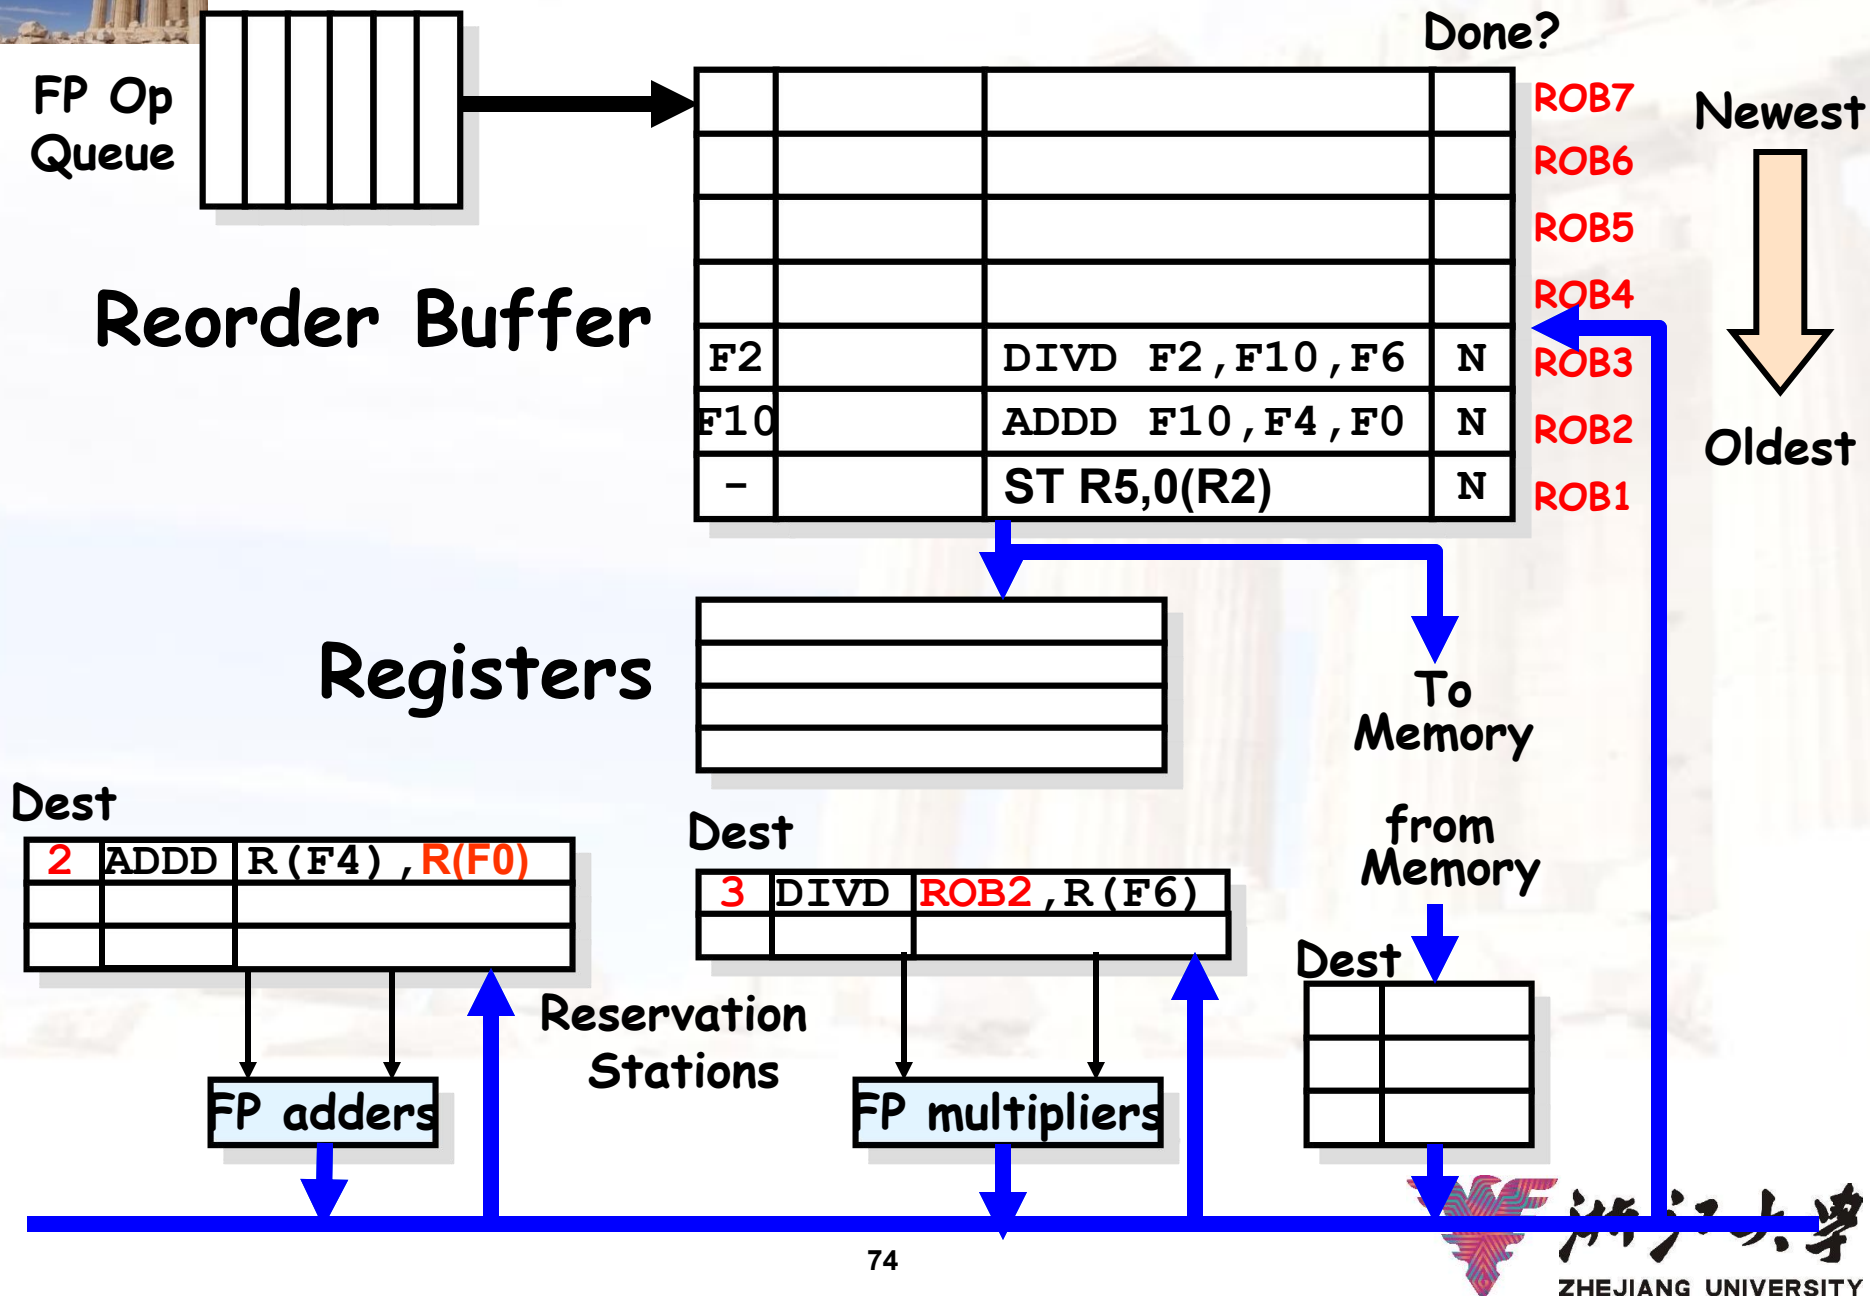


# Tomasulo With Reorder buffer:





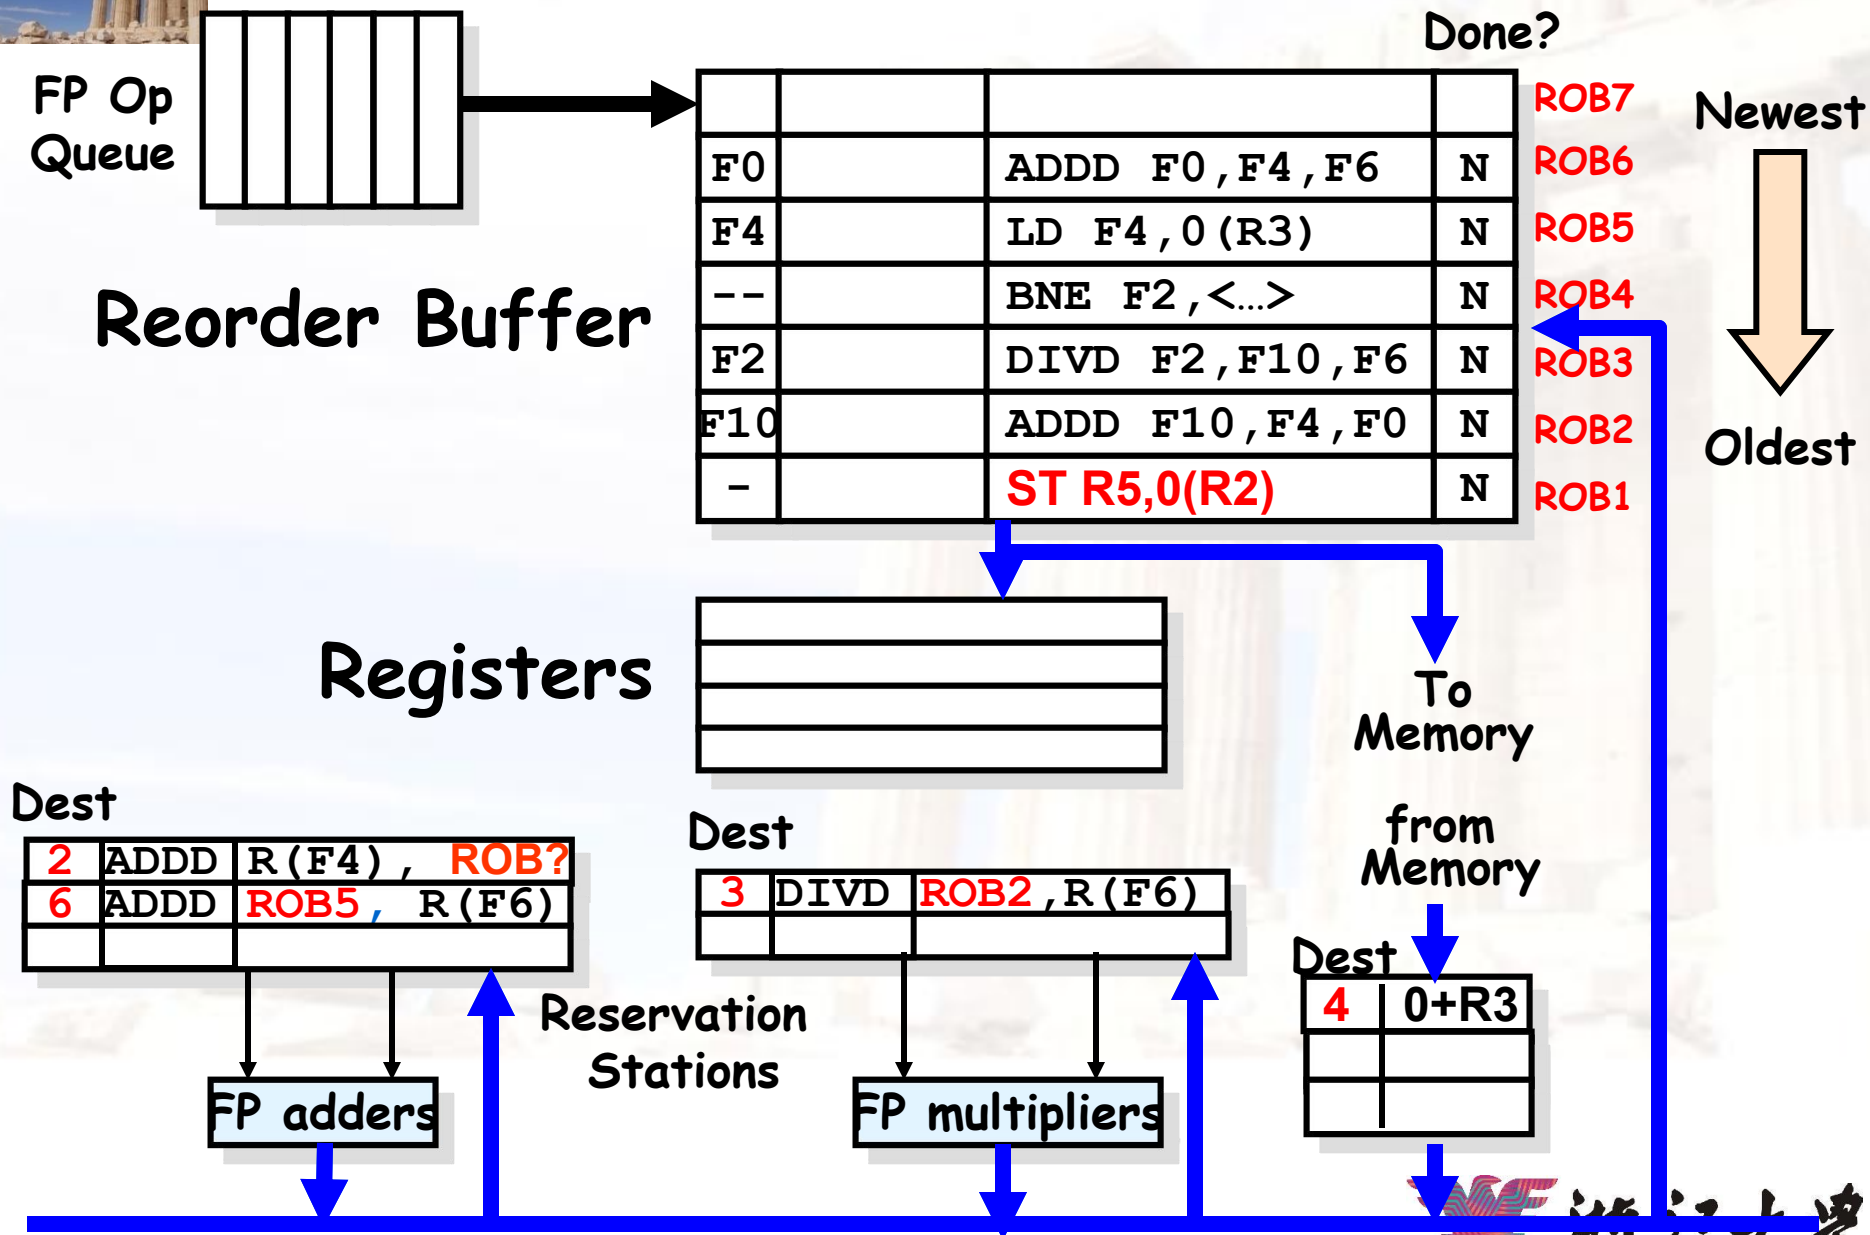
# Tomasulo With Reorder buffer:





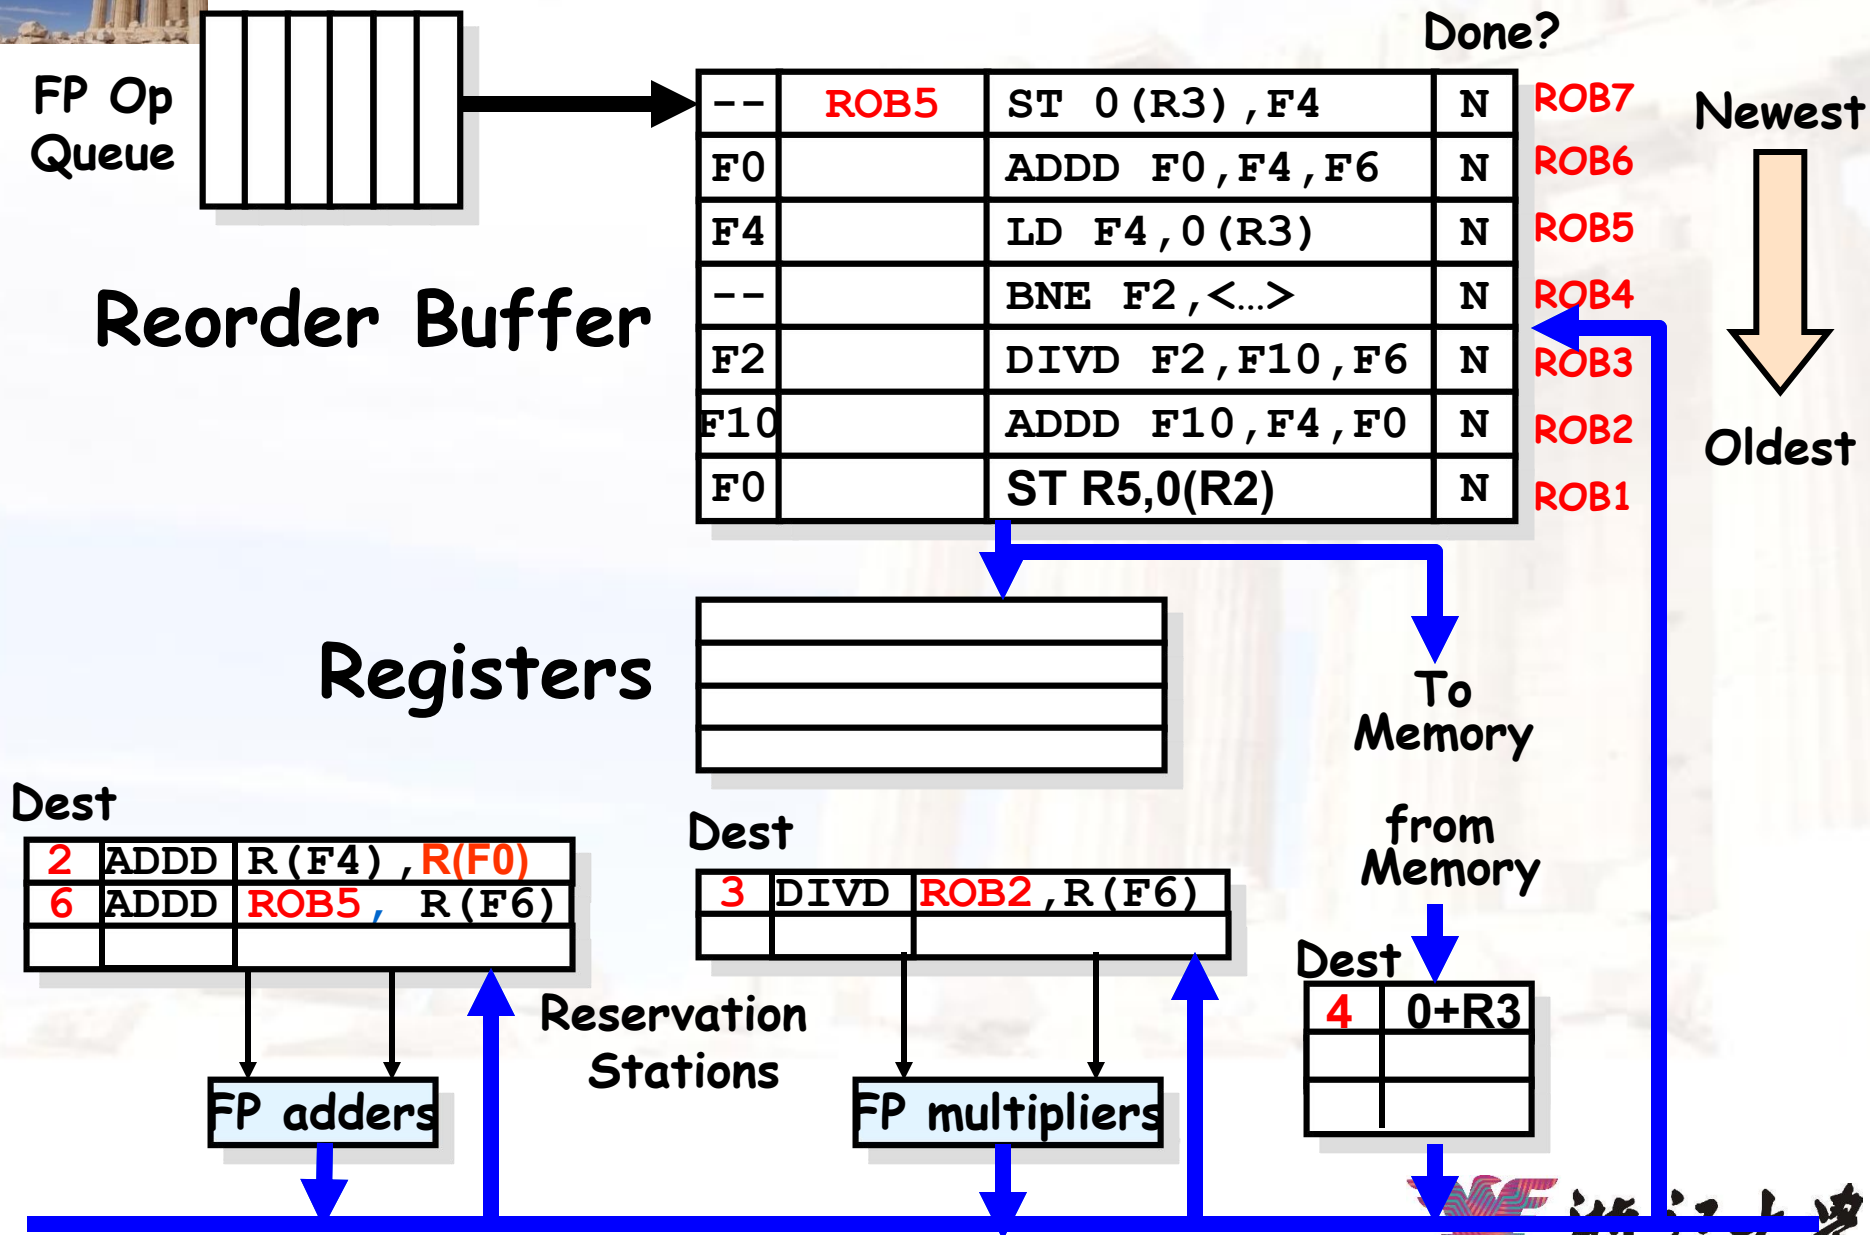


# Tomasulo With Reorder buffer:





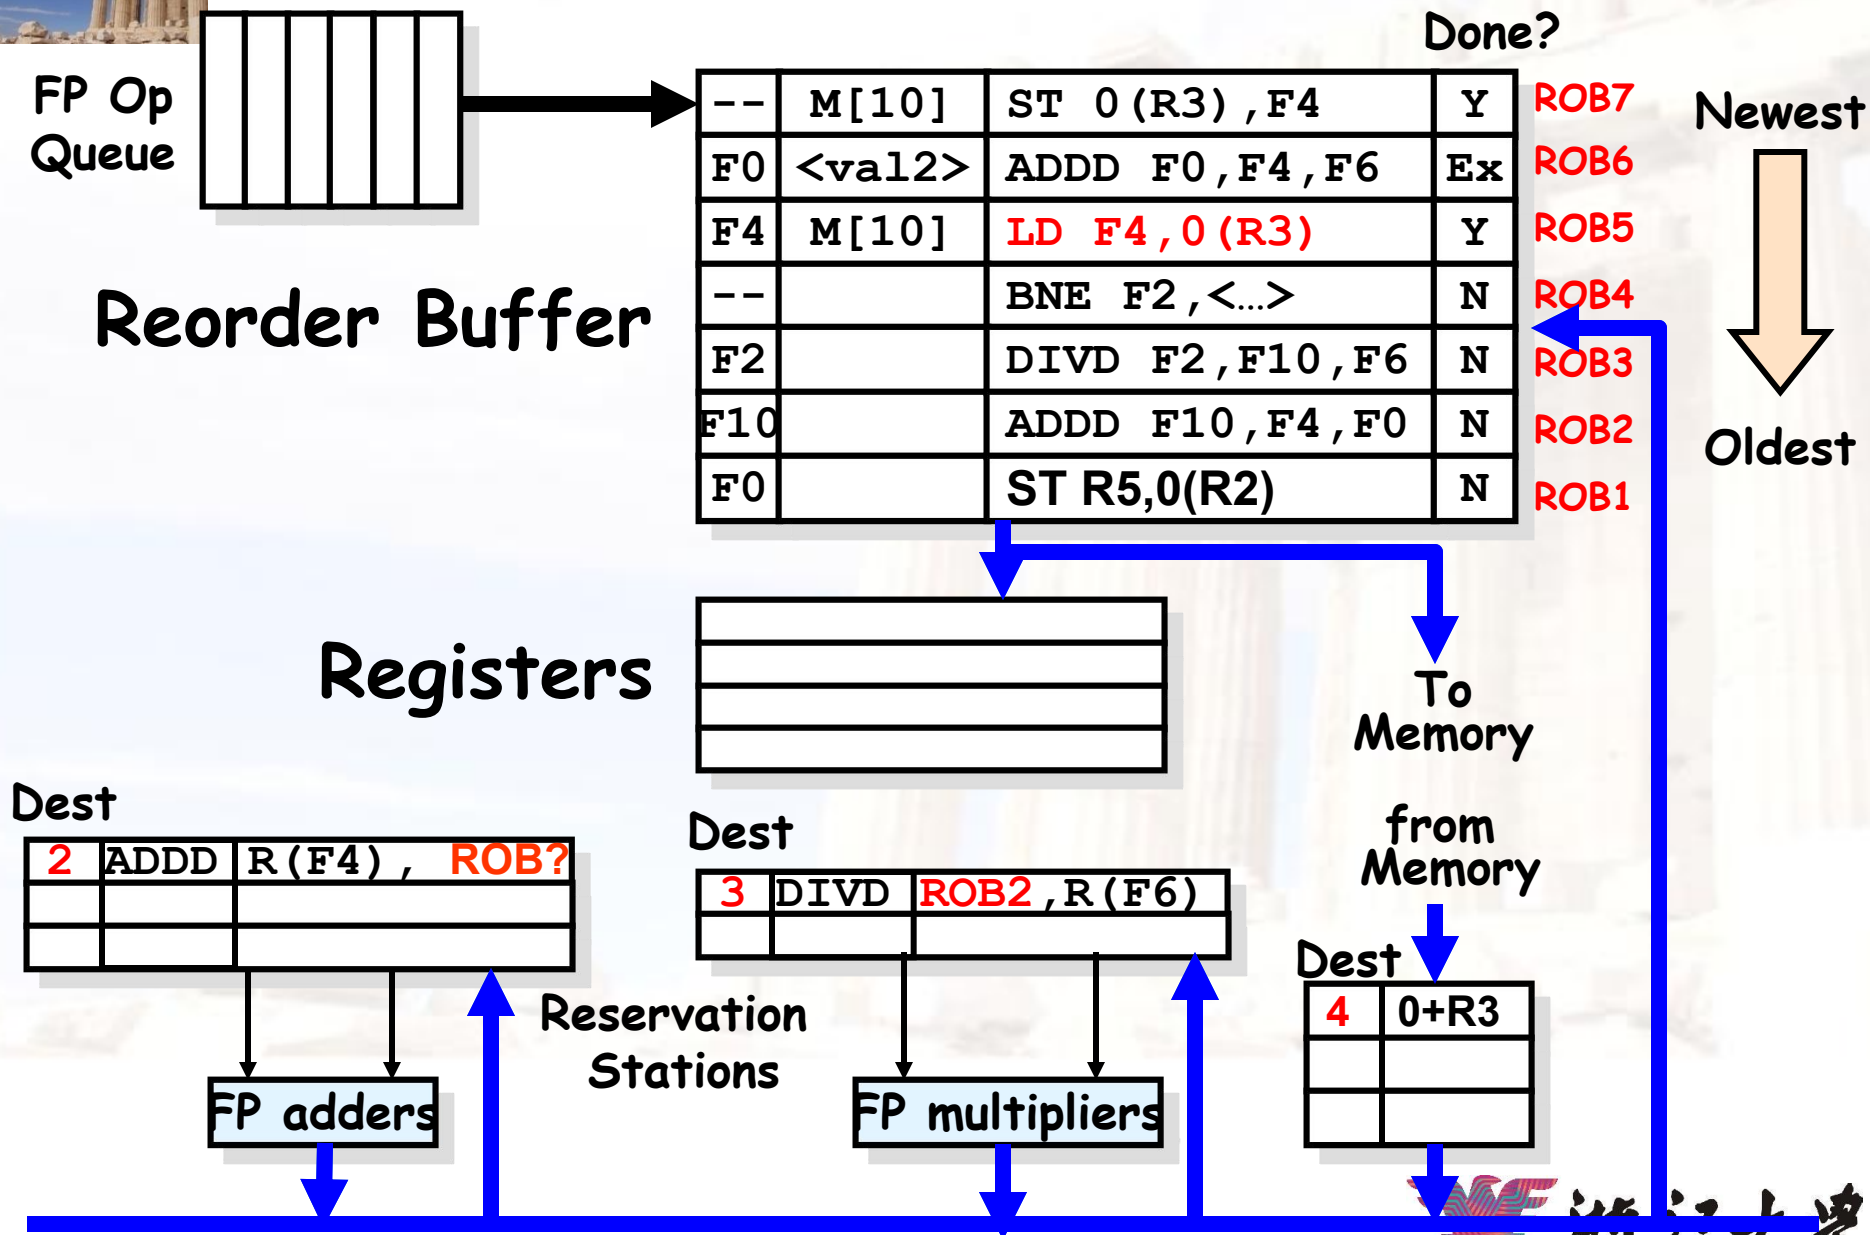
# Tomasulo With Reorder buffer:







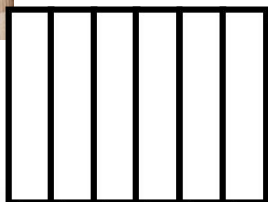
# Tomasulo With Reorder buffer:





# Tomasulo With Reorder buffer:

FP Op Queue




Reorder Buffer

--	M[10]	ST 0(R3), F4	Y
F0	<val2>	ADDD F0, F4, F6	Ex
F4	M[10]	LD F4, 0(R3)	Y
--		BNE F2, <...>	N
F2		DIVD F2, F10, F6	N
F10		ADDD F10, F4, F0	N
F0		ST R5, 0(R2)	N

Done?

ROB7  
ROB6  
ROB5  
ROB4  
ROB3  
ROB2  
ROB1


Newest



Oldest

What if  $0(R3) = 0(R2)$

Registers



To Memory

Dest

2	ADDD	R(F4), ROB?

FP adders

Reservation Stations

Dest

3	DIVD	ROB2, R(F6)

FP multipliers

from Memory

Dest

4	0+R3



# Memory Disambiguation: Sorting out RAW Hazards in memory

• Question: Given a load that follows a store in program order, are the two related?

-(Alternatively: is there a RAW hazard between the store and the load)?

Eg:      ST      O(R2),    R5  
         LD      R6,      O(R3)

• Can we go ahead and start the load early?

-Store address could be delayed for a long time by some calculation that leads to R2 (divide?).

-We might want to issue/begin execution of both operations in same cycle.

-Today: Answer is that we are not allowed to start load until we know that address  $O(R2) \neq O(R3)$





# Hardware Support for Memory Disambiguation

- ❑ Need buffer to keep track of all outstanding stores to memory, in program order.
  - Keep track of address (when becomes available) and value (when becomes available)
  - FIFO ordering: will retire stores from this buffer in program order
- ❑ When issuing a load, record current head of store queue (know which stores are ahead of you).



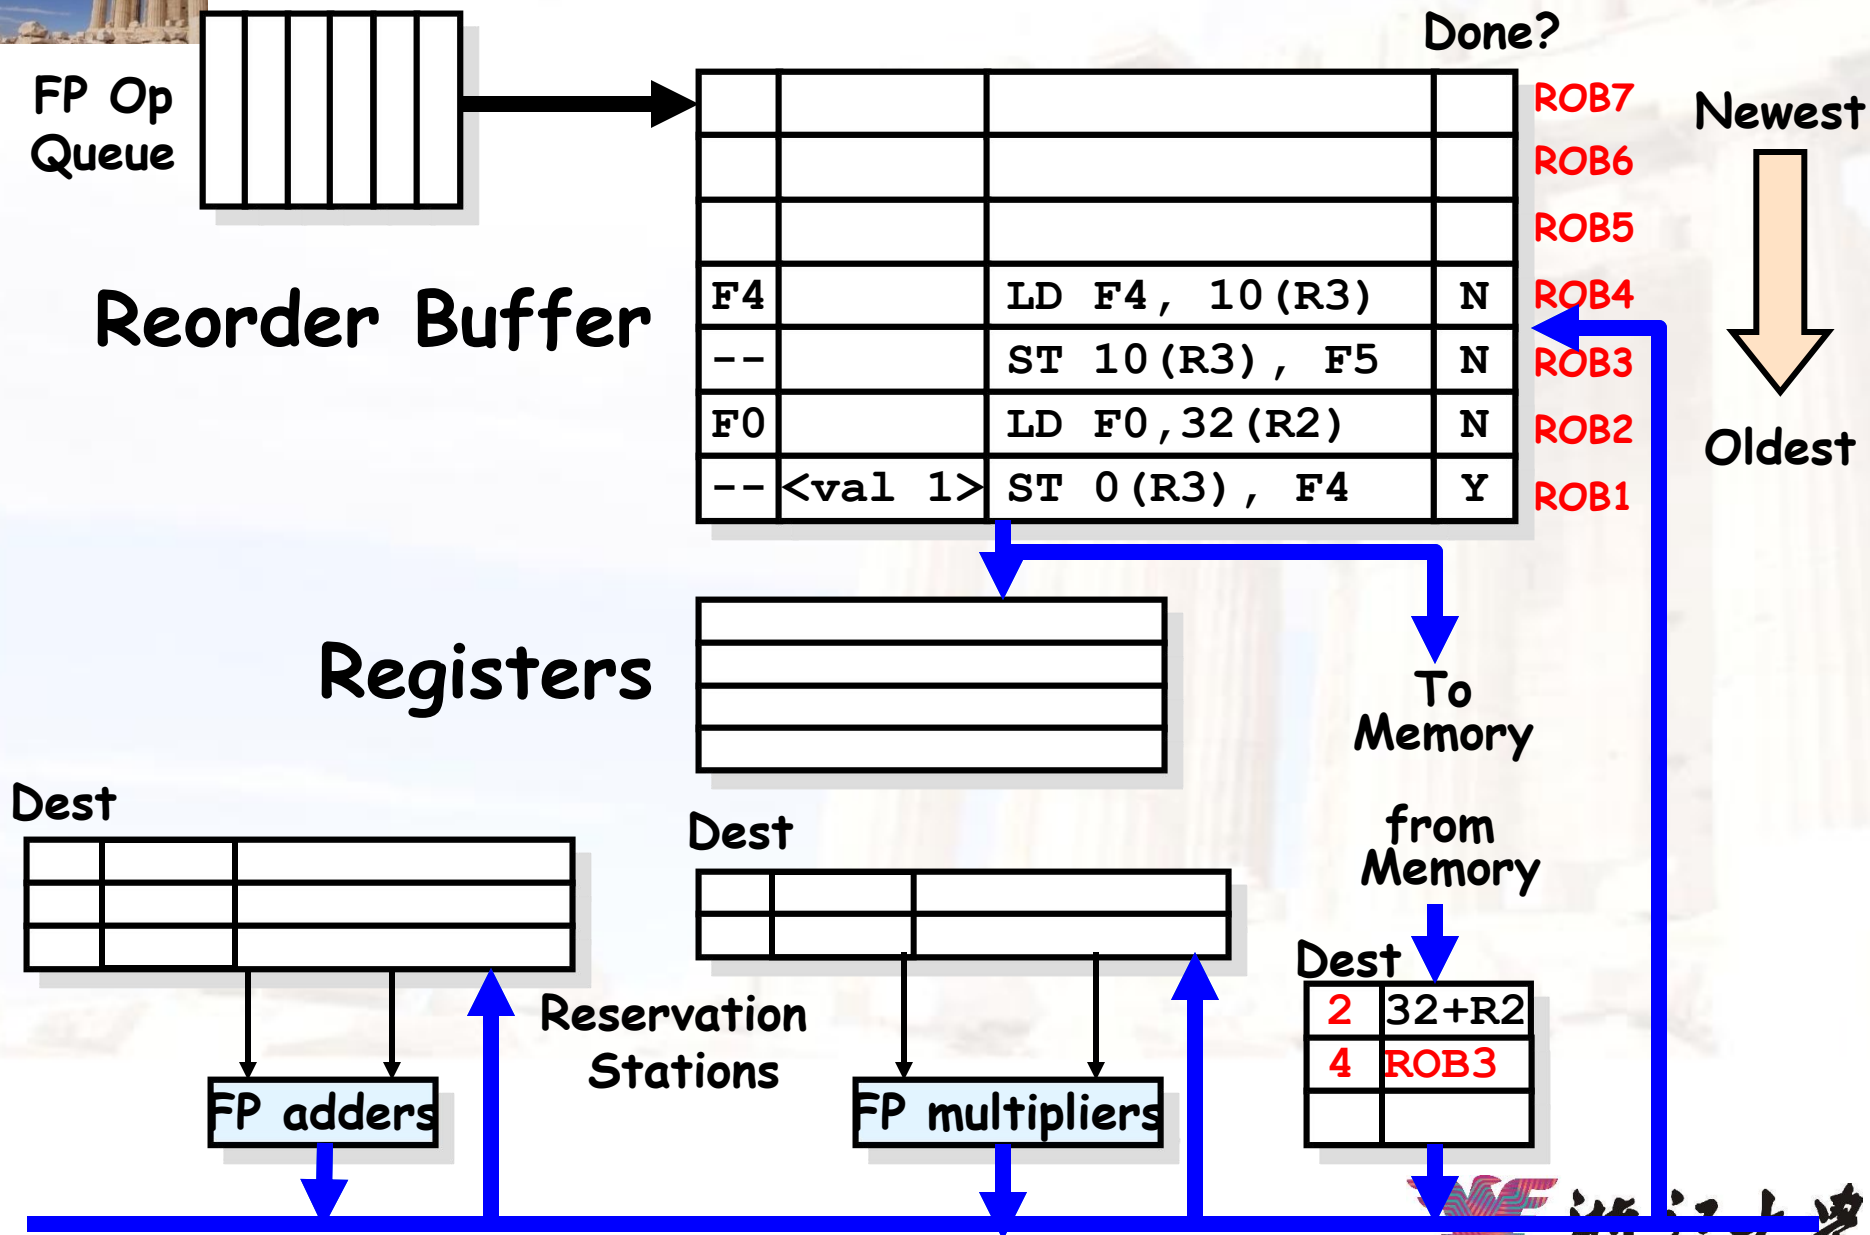


# Hardware Support for Memory Disambiguation(2)

- ❑ When have address for load, check store queue:
  - If *any* store prior to load is waiting for its address, stall load.
  - If load address matches earlier store address (associative lookup), then we have a *memory-induced RAW hazard*:
    - store value available  $\Rightarrow$  return value
    - store value not available  $\Rightarrow$  return ROB number of source
  - Otherwise, send out request to memory
- ❑ Actual stores commit in order, so no worry about WAR/WAW hazards through memory.



# Memory Disambiguation:



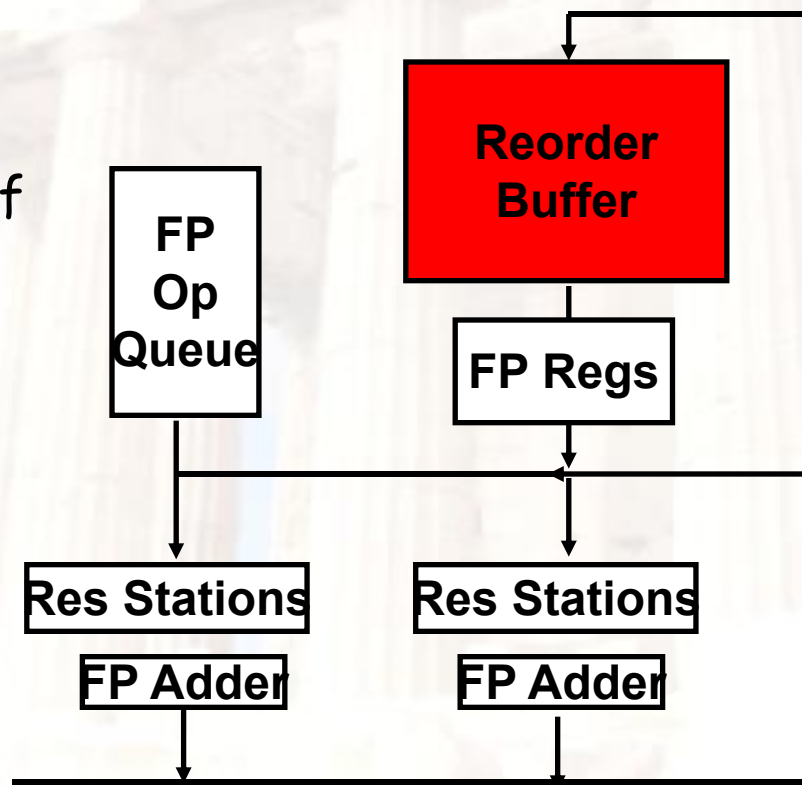


# HW support for precise interrupts

❑ Need HW buffer for results of uncommitted instructions:

## reorder buffer

- 3 fields: instr, destination, value
- Use **reorder buffer number** instead of reservation station when execution completes
- Supplies operands between execution complete & commit
- (Reorder buffer can be operand source => more registers like RS)
- Once instruction commits, result is put into register
- As a result, easy to undo speculated instructions on mispredicted branches or exceptions



	记分牌	Tomasulo	带投机的Tomasulo
主要思想	动态调度、按序发射、乱序执行、乱序完成、集中控制。	动态调度、按序发射、乱序执行、乱序完成分散控制。增加RS（保留站）,CDB。	转移预测, 动态调度、投机执行、按序完成。在RS基础上再增加ROB(reorder buffer)。
解决何种竞争	RAW数据竞争。 WAW,WAR时stall	解决数据竞争,避免WW,WAR竞争;允许硬件支持的循环展开，不局限于程序基本块。	解决数据竞争,避免WW,WAR竞争;从数据竞争扩充到控制竞争。通过预测减少控制竞争stall。
指令节拍	ISSUE：停顿于结构竞争和WAW竞争。 RO：等待两个操作数就绪, 动态解决RAW。 EX：和普通流水线一样。 WB：停顿于WAR竞争。	ISSUE：仅停顿于结构竞争（RS部件）。 EX：利用保留站（RS）进行换名，解决了WAW和WAR问题（即名字相关）。 WB：回写。	ISSUE：仅停顿于结构竞争（RS部件和ROB部件）。 EX：等待操作数就绪，执行。 WB：结果写到ROB，其它指令可以引用该临时结果，但不更新寄存器。 COMMIT：保证按序完成，保证精确中断。
关键部件	记分牌、集中控制、指令状态、功能部件状态、寄存器结果状态。	CDB(通用数据总线)：执行结果通过CDB传递，及时。 RS(保留站)：换名功能以及暂存临时数据。	ROB(REORDER BUFFER)重构序缓冲器：对移到顶部的指令进行判断。。。ROB负责寄存器重命名、暂存执行后和交付前的结果、
其它	从静态调度到动态调度的过渡算法。	一旦指令将临时结果放在保留站，则其它后续的指令都能立即在寄存器组中看见这个结果。	某指令将临时结果放在ROB后，其它指令并不能从寄存器组中看到这个结果（也就是说，寄存器组的值并没有更新），而是等待该指令提交（commit）后，寄存器组才更新。



# ARM处理器体系结构

## □ARM流水线的执行顺序：

- 取指令（Fetch）：从存储器读取指令；
- 译码（Decode）：译码以鉴别它是属于哪一条指令；
- 执行（Execute）：将操作数进行组合以得到结果或存储器地址；
- 缓冲/数据（Buffer/data）：如果需要，则访问存储器以存储数据；
- 回写：（Write-back）：将结果写回到寄存器组中；



# 基于ARMv8的鲲鹏流水线技术



Taishan coreV110 Pipeline Architecture





# 基于ARMv8的鲲鹏流水线技术

- Branch预测和取指流水线解耦设计，取指流水线每拍最多可提供32Bytes指令供译码，分支预测流水线可以不受取指流水停顿影响，超前进行预测处理；
- 定浮点流水线分开设计，解除定浮点相互反压，每拍可为后端执行部件提供4条整型微指令及3条浮点微指令；
- 整型运算单元支持每拍4条ALU运算（含2条跳转）及1条乘除运算；
- 浮点及SIMD运算单元支持每拍2条ARM Neon 128bits 浮点及SIMD运算；
- 访存单元支持每拍2条读或写访存操作，读操作最快4拍完成，每拍访存带宽为2x128bits读及1x128bits写；