



Operating Systems (操作系统 2024秋冬)

教七-306 (周3第7-8节, 周5第1-2节)

实验课: 曹西-503 (周4第11-12节)

Lecture 1: Overview

Shou Lidan
should @ zju.edu.cn

Outlines of the overview

- Previously split in two modules, known as “Principles of Operating System” and “Operating System Practices”
- Background
 - The main **purpose** of this course
 - The **contents** and **schedule** of the course
 - The pre-requisite knowledge for the course
- Chapter 1 Introduction
- Chapter 2 Operating-System Structures

Main objectives

An introductory course for Operating System

- To learn the basics and internal design of operating systems
- To look at both the history and the state-of-the-art techniques used in operating systems
- To study the design **methodologies** applied in OS
- To **prepare to practice** the techniques in your future work and
- To **enjoy hacking around!**

Main Contents

- Overview
 - Intro
 - OS structure
- Process Management
 - Processes
 - Threads
 - CPU scheduling
 - Process Synchronization
 - Deadlocks
- Memory Management
 - Main memory
 - Virtual memory
- Storage Management
 - File-system interface
 - File-system implementation
 - Mass-storage structure
 - I/O systems

Why is OS still important in ChatGPT-era?


- There's plenty of room at the Top: What will drive computer performance after Moore's law?

<https://www.science.org/doi/10.1126/science.aam9744>

- Performance improved **63000x**
- C program can be optimized **1.300x**

multiplying two 4096-by-4096 matrices

```
for i in xrange(4096):  
    for j in xrange(4096):  
        for k in xrange(4096):  
            C[i][j] += A[i][k] * B[k][j]
```




Software performance engineering, development of algorithms, and hardware streamlining at the Top can continue to make computer applications faster in the post-Moore era.

Textbook & Resources

- **Main Textbook: Operating System Concepts 7th, Abraham Silberschatz, Peter Galvin, Greg Gagne, 2005.1**
- 高等教育出版社（第7版影印版）
- **10th edition can be found on Amazon.**
- **Course webpage: <http://courses.zju.edu.cn/>**
- **Lecture slides: <http://courses.zju.edu.cn/>**

Reference Books

- 边干边学Linux内核指导



- xv6: a simple, Unix-like teaching operating system, Russ Cox, Frans Kaashoek, Robert Morris, August 31, 2020
- Latest version: <https://pdos.csail.mit.edu/6.1810/2022/xv6.html>

Online Assignments



<http://courses.zju.edu.cn/>

Teaching Assistants: 刘得志、石劲、包牧天
Please submit your assignment IN TIME!

Pre-requisite Knowledge

- Computer system architecture/organization
- Programming language C
- Data structures

Lab Projects


- 每个同学独立完成
 - 实验0 RISC-V 64 内核调试 5%
 - 实验1 内核引导；时钟和中断 15%
- 分组完成5个实验，每组2人
 - 实验2 线程调度,上下文切换 15%
 - 实验3 虚拟内存管理 15%
 - 实验4 用户模式 (shell) 20%
 - 实验5 Page fault 20%
 - 实验6 Fork 10%
- 实验7 File System Bonus 10%

Lab Projects

- 实验0、1：每个同学提交一份实验报告；
- 实验2-6：每个同学提交一份实验报告（侧重自己完成的那部分内容）；
 - 自己写的代码必须有详细的注释，每5行代码有注释
 - 每个实验必须写“讨论心得”（实验过程中遇到的问题及解决方法）这部分内容占本实验报告20%分数
- 每个实验完成后，以个人（lab0-1）/小组（lab2-3-4-5-6）为单位向老师/助教演示，以完成验收。
- 实验说明文档在 <https://zju-sec.github.io/os24fall-stu/>
- 实验所需代码 <https://github.com/ZJU-SEC/os24fall-stu>

Final Grade

- Final exam (50%)
- Assignments/homework (5%)
- In-class Quiz (5%)
- Lab Reports 实验报告 (20%)
- Lab Demos 实验验收 (20%)



My suggestions on learning OS

- Do a lot of readings before or after the lecture sessions, especially from the English textbook.
- Do NOT refer to the so-called “standard” answers to the exercises, they contain mistakes. Once you are caught using the wrong “standard” answers, penalty might be applicable.
- Hopefully more interactions during the lectures.

怎样学好系统课程

- 系统、全局的观察角度 vs 微观、细致的想象能力

王阳明《蔽月山房》

山近月远觉月小，便道此山大于月。

若有人眼大如天，当见山高月更阔。

- 理论钻研 vs 动手实践

丁肇中：在环境激变的今天，我们应该重新体会到几千年前经书里说的格物致知真正的意义。这意义有两个方面：第一，寻求真理的唯一途径是对事物客观的探索，而不是仅从内心就能领悟；第二，探索的过程不是消极的袖手旁观，而是深度思考勇于探索。

希望我们这一代对于格物和致知有新的认识和思考，使得实干精神真正地变成中国文化的一部分。

- 寻找模式

- 螺旋式上升；上下文；封装；抽象；过程式和声明式

- 寻找兴奋点

Lab0: GDB + QEMU 调试 LINUX

- 请先将os24fall-stu目录 clone 下来
- 实验说明文档在 os24fall-stu/docs/lab0.md
 - <https://zju-sec.github.io/os24fall-stu/> 可浏览器阅读文档
- 实验重点：掌握实验环境、尝试简单调试
- gdb有一个text UI可以尝试使用，按 Ctrl-x Ctrl-a 进入或者 “gdb -tui”

Chapter 1 Introduction

Outlines for Chapter ONE

- What Operating Systems Do
 - Computer-System Organization
 - Computer-System Architecture
- Overview
- Operating-System Structure
 - Operating-System Operations
- Structure
- Process Management
 - Memory Management
 - Storage Management
- Main parts
- Protection and Security
 - Distributed Systems
 - Special-Purpose Systems
 - Computing Environments

Objectives

- To provide a grand tour of the major operating systems components
- To provide coverage of basic computer system organization

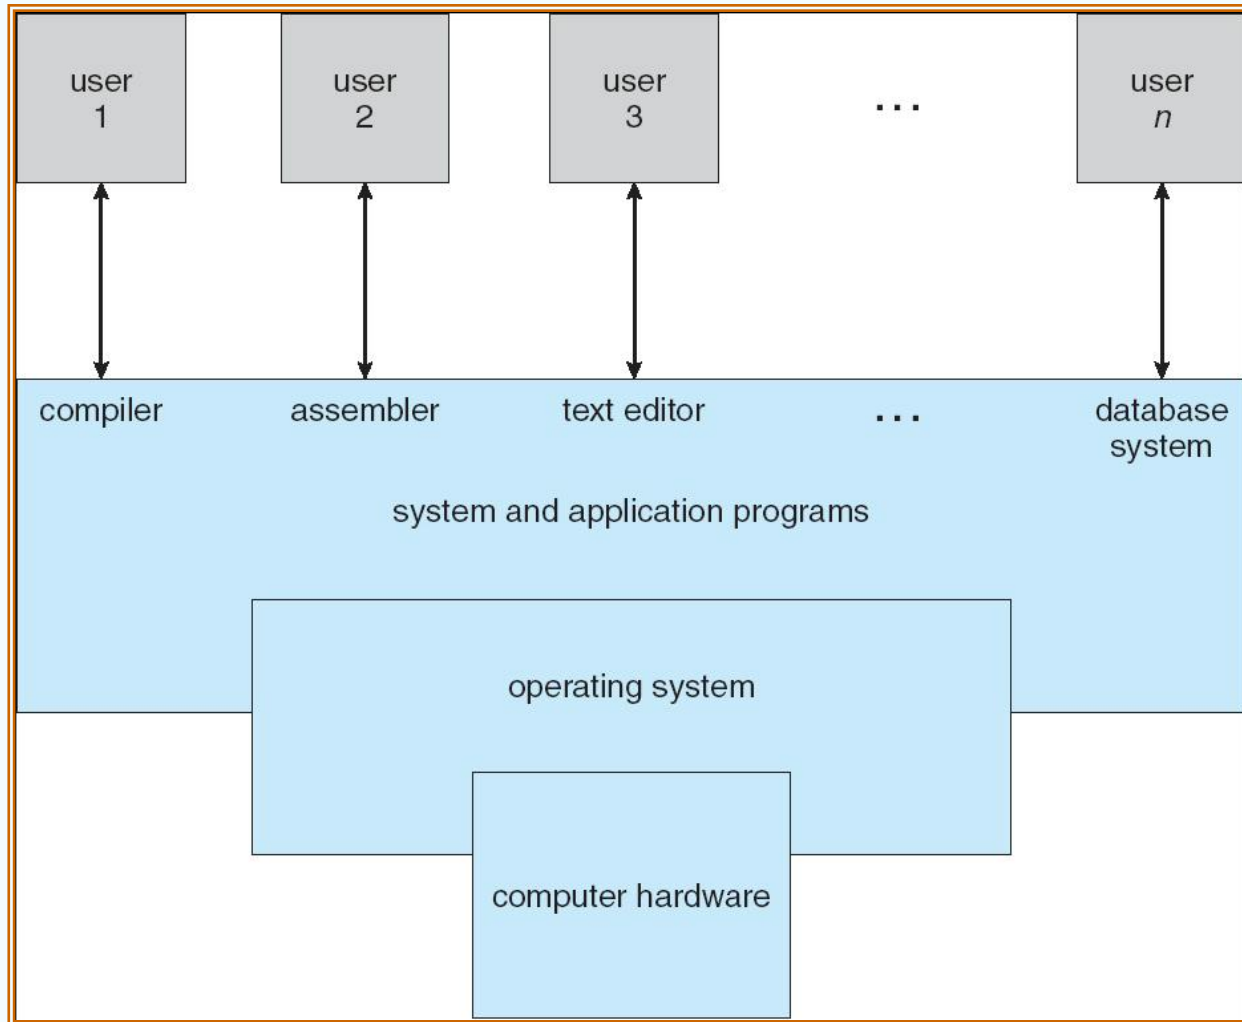
What is an Operating System?

- A program that acts as an intermediary between a user of a computer and the computer hardware.
- Operating system goals:
 - Execute user programs and make solving user problems easier.
 - Make the computer system convenient to use.
- Use the computer hardware in an efficient manner.

Computer System Structure


- Computer system can be divided into four components
 - Hardware – provides basic computing resources
 - ▶ CPU, memory, I/O devices
 - Operating system
 - ▶ Controls and coordinates use of hardware among various applications and users
 - System & application programs – define the ways in which the system resources are used to solve the computing problems of the users
 - ▶ Word processors, compilers, web browsers, database systems, video games
 - Users
 - ▶ People, machines, other computers

Four Components of a Computer System



Why To Learn Operating Systems?


- Building blocks of modern computer software



Why To Learn Operating Systems?

New system technologies involving OS:

- Cloud Computing becomes the major computing model
- Virtualization, Software-Defined Storage/Network
- Containers
- Hadoop eco-system



Operating System Definition

- OS is a **resource allocator**
 - Manages all resources
 - Decides between conflicting requests for efficient and fair resource use
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer

Operating System Definition (Cont.)

- No universally accepted definition
- “Everything a vendor ships when you order an operating system” is good approximation
 - But varies wildly
- “The one program running at all times on the computer” is the **kernel**. Everything else is either a system program (ships with the operating system) or an application program

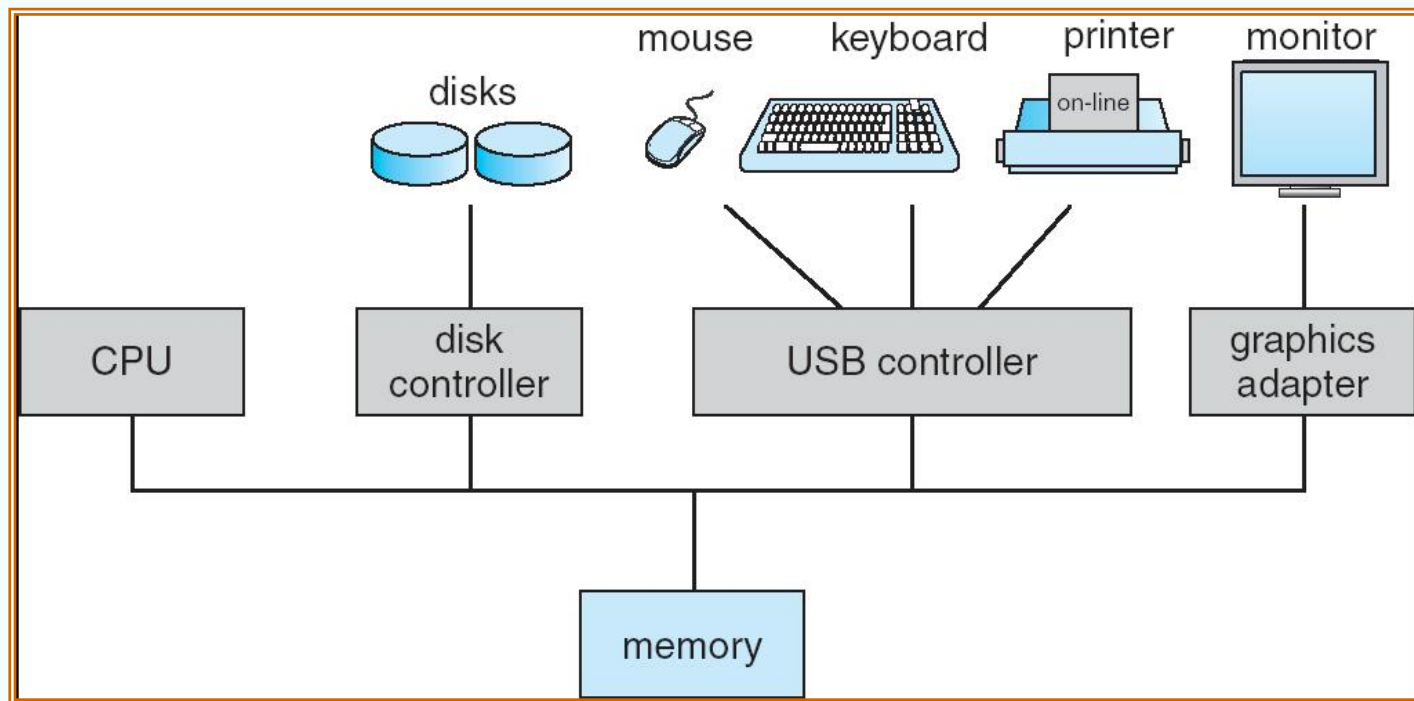
Computer Startup

- **bootstrap program** is loaded at power-up or reboot
 - Typically stored in ROM or EPROM, generally known as **firmware**
 - Initializes all aspects of system
 - Loads operating system kernel and starts execution



Computer System Organization

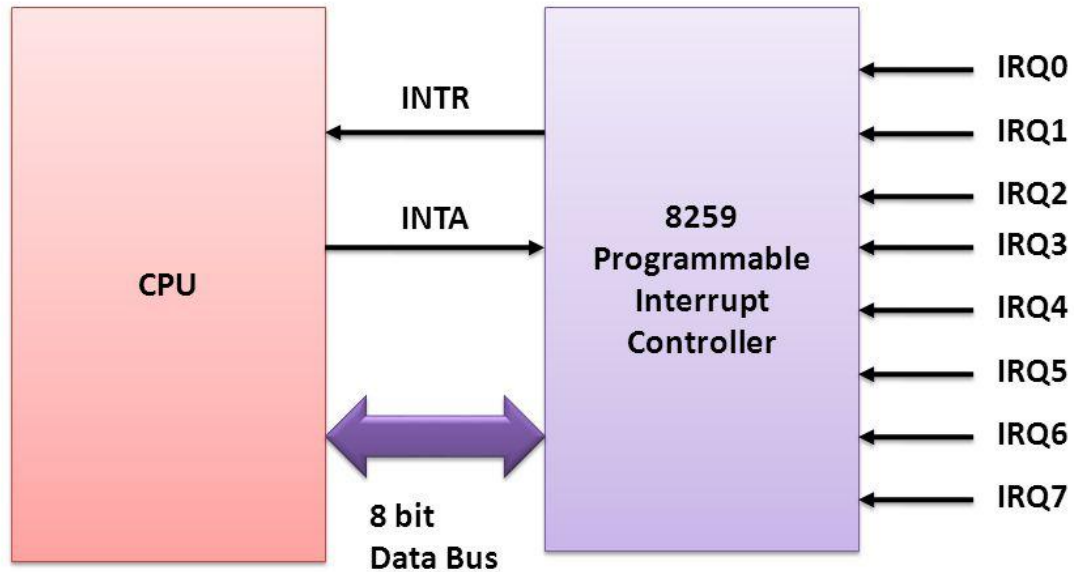
- Computer-system operation
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Concurrent execution of CPUs and devices competing for memory cycles



Computer-System Operation

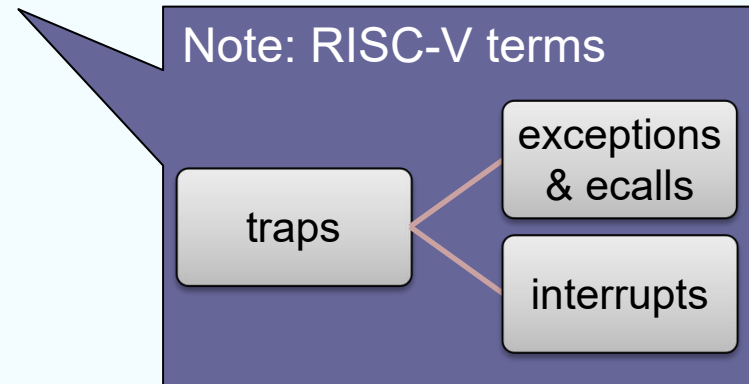
- I/O devices and the CPU can execute concurrently.
- Each device controller is in charge of a particular device type.
- Each device controller has a **local buffer**.
- CPU moves data from/to main memory to/from local buffers
- I/O is from the device to local buffer of controller.
- Device controller informs CPU that it has finished its operation by causing an *interrupt* (**via system bus**).

8259: Programmable Interrupt Controller



Common Functions of Interrupts

- Interrupt transfers control to the interrupt service routine generally, through the *interrupt vector*, which contains the addresses of all the service routines.
- Interrupt architecture must save the address of the interrupted instruction.
- Incoming interrupts are *disabled* while another interrupt is being processed to prevent a *lost interrupt*.
- A *trap* is a software-generated interrupt caused either by an *error* or a *user request* (the latter is often referred to as a *system call*). Note: names may vary across different architectures.
- An operating system is *interrupt* driven.



Interrupt Handling

- The operating system preserves the state of the CPU by storing **registers** and the **program counter**.
- Determines which type of interrupt has occurred:
 - a generic routine examines the interrupt info
 - **vector**ed interrupt system, indexed by a unique number
- Separate segments of code determine what action should be taken for each type of interrupt

Interrupt Timeline


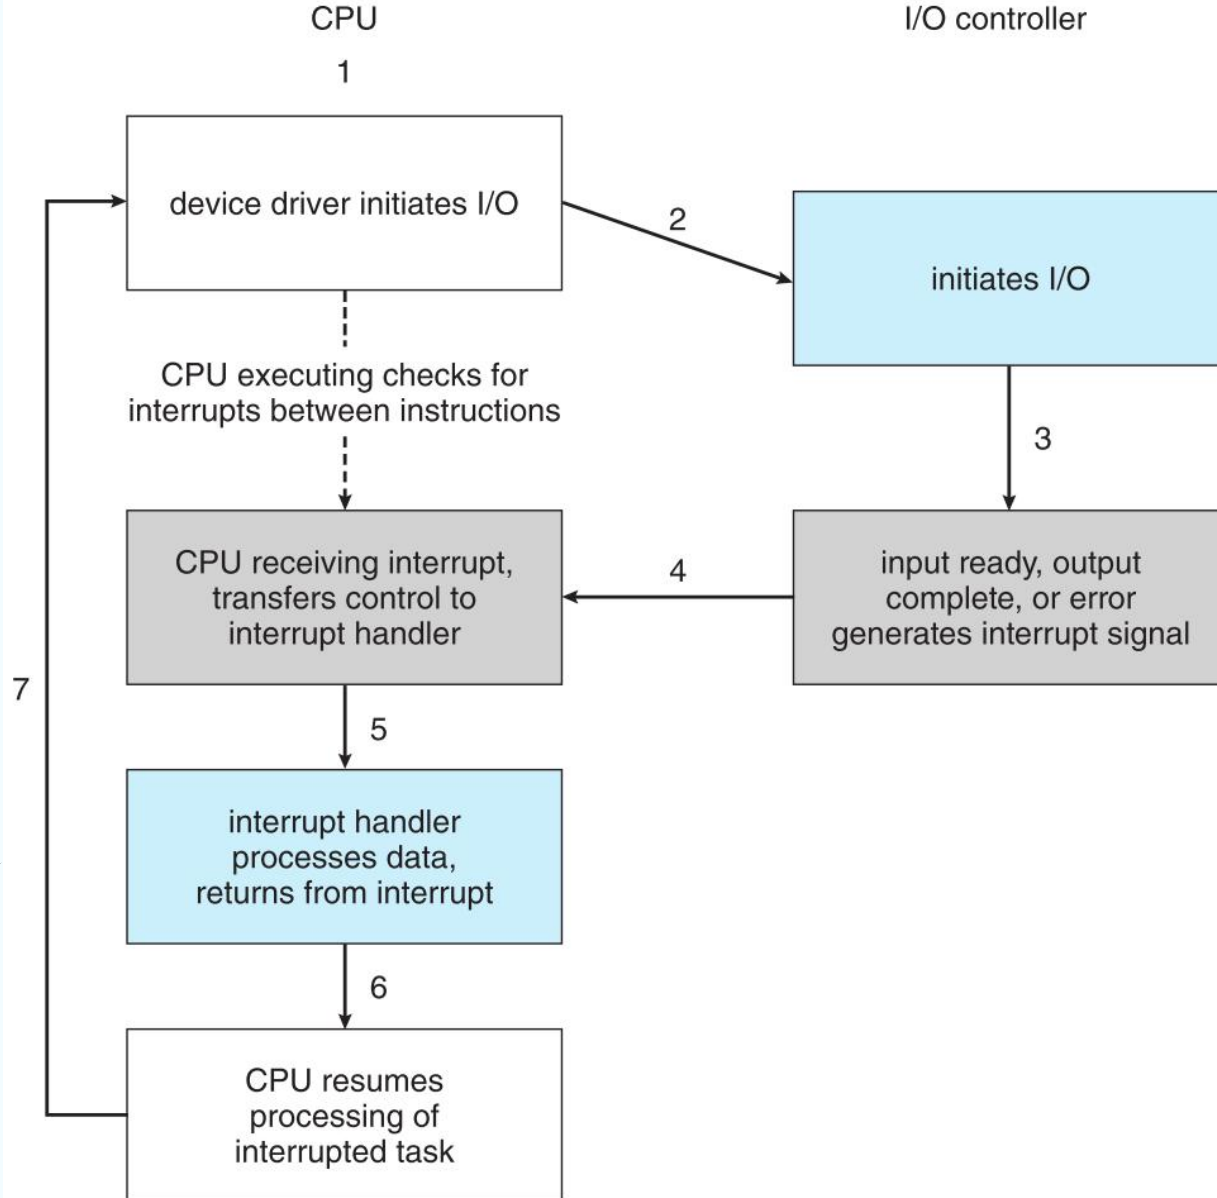


Figure 1.3 Interrupt timeline for a single program doing output.

Interrupt-Driven I/O Cycle



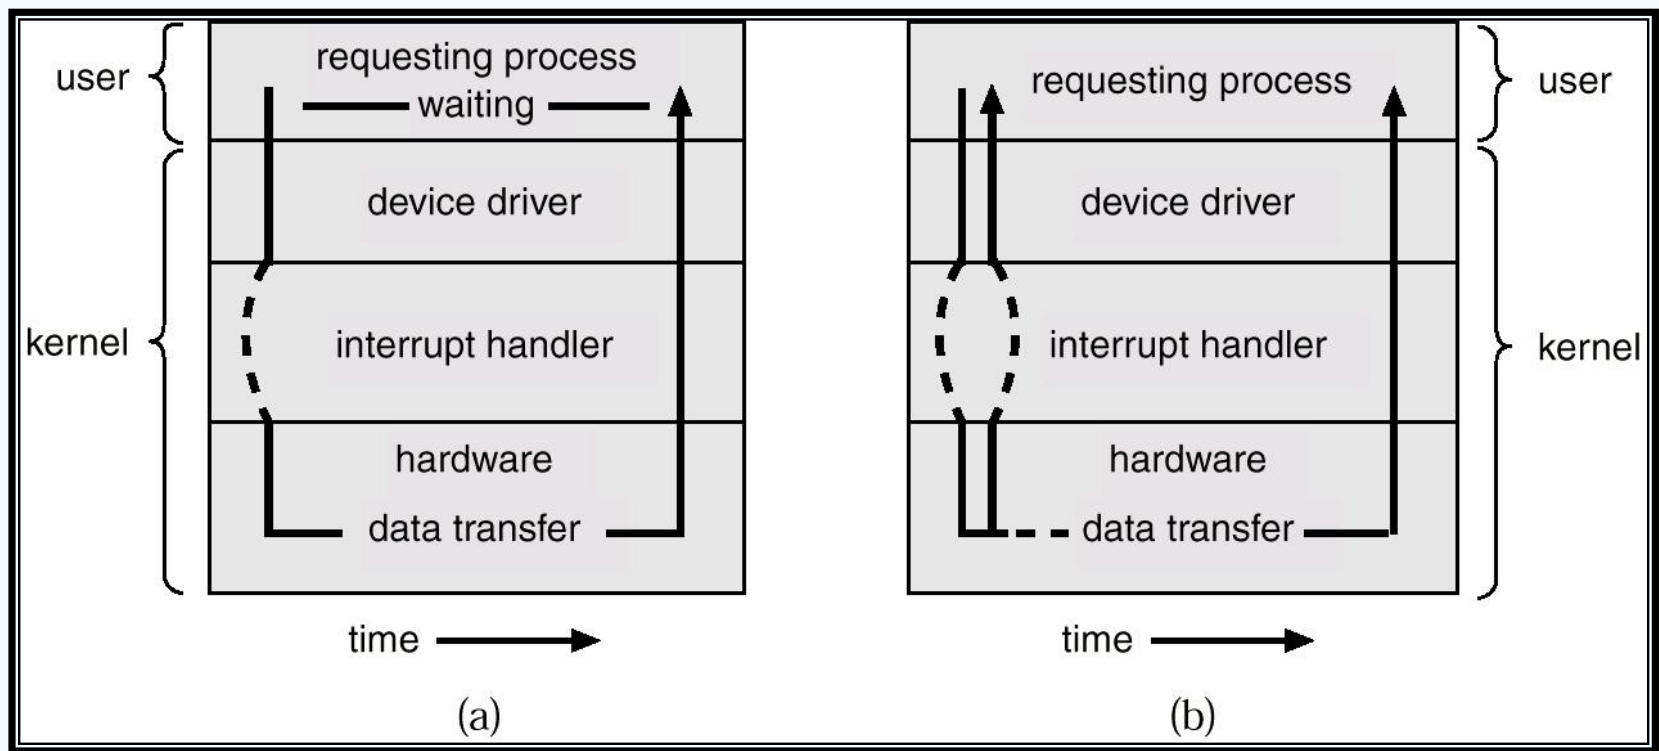
I/O Structure – Two I/O Methods

- After I/O starts, control returns to user program only upon I/O completion.
 - Wait instruction idles the CPU until the next interrupt
 - Wait loop (contention for memory access).
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing.
- After I/O starts, control returns to user program without waiting for I/O completion.
 - *System call* – request to the operating system to allow user to wait for I/O completion.
 - *Device-status table* contains entry for each I/O device indicating its type, address, and state.
 - Operating system indexes into I/O device table to determine device status and to modify table entry to include interrupt.

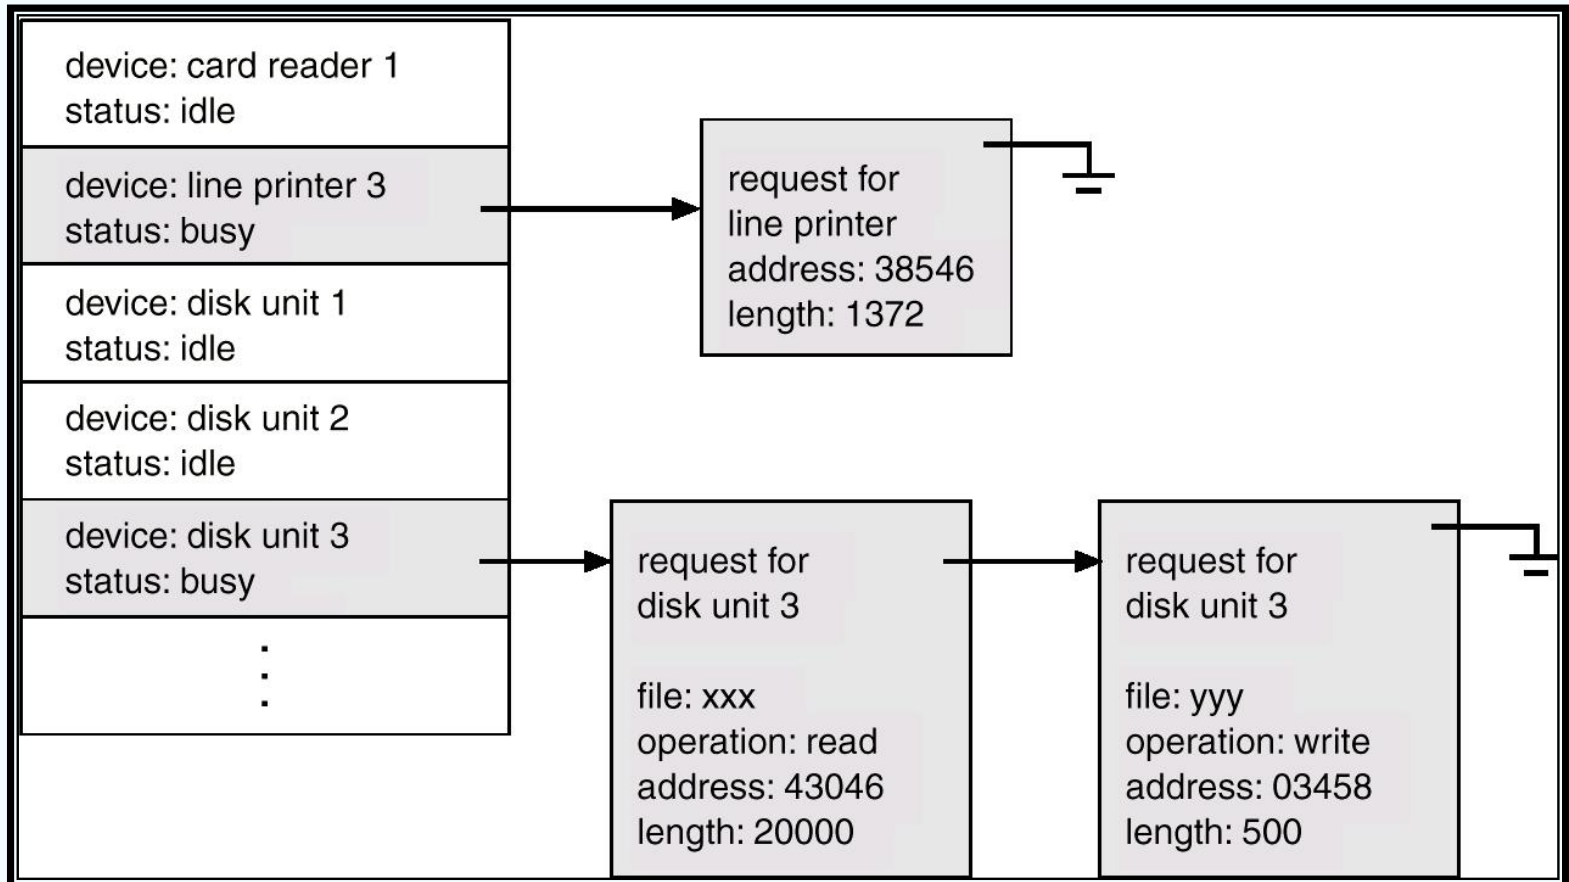
Two I/O Methods

Synchronous

Asynchronous



Device-Status Table



Direct Memory Access Structure

- Used for high-speed I/O devices able to transmit information at close to memory speeds.
- Device controller transfers blocks of data from buffer storage directly to main memory **without** CPU intervention.
- Only one interrupt is generated **per block**, rather than the one interrupt per byte.


Storage Structure

- Main memory – only large storage media that the CPU can access directly.
- Secondary storage – extension of main memory that provides large nonvolatile storage capacity.
- Magnetic disks – rigid metal or glass platters covered with magnetic recording material
 - Disk surface is logically divided into *tracks*, which are subdivided into *sectors*.
 - The *disk controller* determines the logical interaction between the device and the computer.

Storage Hierarchy

- Storage systems organized in hierarchy.
 - Speed
 - Cost
 - Volatility
- *Caching* – copying information into faster storage system; main memory can be viewed as a last *cache* for secondary storage.

Storage-Device Hierarchy



Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
- Information in use copied from slower to faster storage temporarily
 - Speed mismatch
- Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
- Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy


Performance of Various Levels of Storage

- Movement between levels of storage hierarchy can be explicit or implicit

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS SRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000.000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

Migration of Integer A from Disk to Register

- **Multitasking** environments must be careful to use most recent value, no matter where it is stored in the storage hierarchy



- **Multiprocessor** environment must provide **cache coherency** in hardware such that all CPUs have the most recent value in their cache

Multiprocessor Systems

- SMP architecture
 - Each CPU processor has its own set of registers
 - All processors share physical memory over the system bus

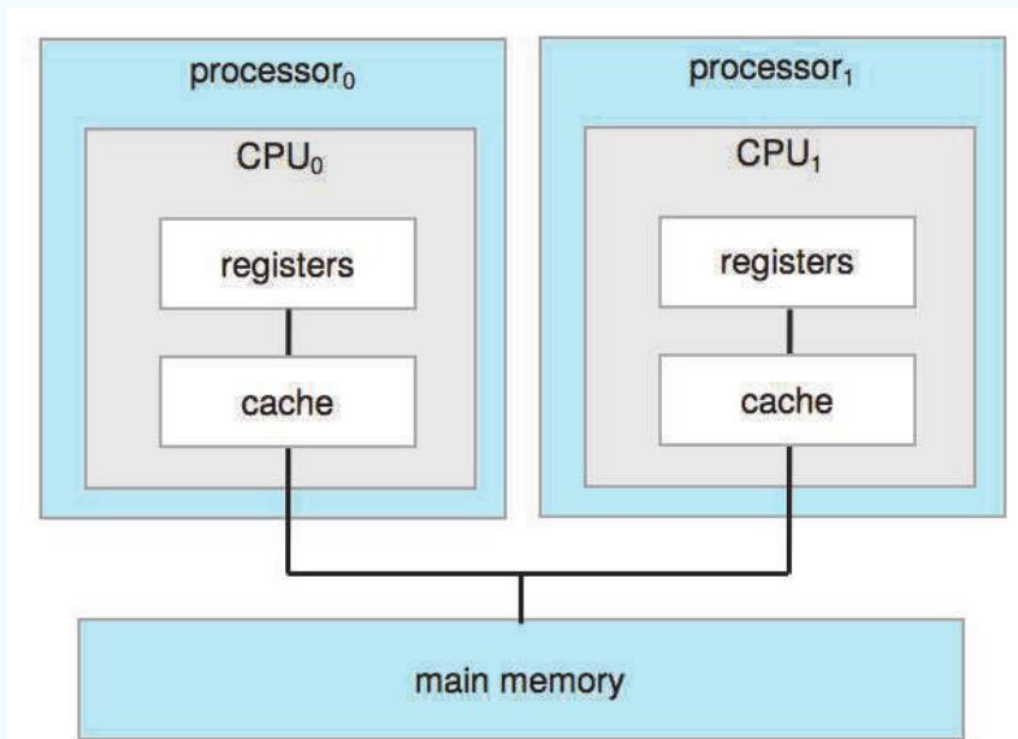



Figure 1.8 Symmetric multiprocessing architecture.


Multicore Systems

- On-chip communication is faster than between-chip communication
- Less power (good for mobile devices)



The NUMA Architecture


- The CPUs are connected by a shared system interconnect
- Scales more effectively as more processors are added
- Remote memory across the interconnect is slow
- Operating systems need careful CPU scheduling and memory management



Operating System Structure

- **Multiprogramming** needed for efficiency (**CPU utilization**)
 - Single user cannot keep CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs (code and data) so CPU always has one to execute
 - A subset of total jobs in system is kept in memory
 - One job selected and run via **job scheduling**
 - When it has to wait (for I/O for example), OS switches to another job
- **Timesharing (multitasking)** is logical extension in which CPU switches jobs so frequently that users can interact with each job while it is running, creating **interactive** computing (**interactivity**)
 - **Response time** should be < 1 second
 - Each user has at least one program executing in memory [**process**]
 - If several jobs ready to run at the same time [**CPU scheduling**]
 - If processes don't fit in memory, **swapping** moves them in and out to run
 - **Virtual memory** allows execution of processes not completely in memory

Memory Layout for Multiprogrammed System



Operating-System Operations

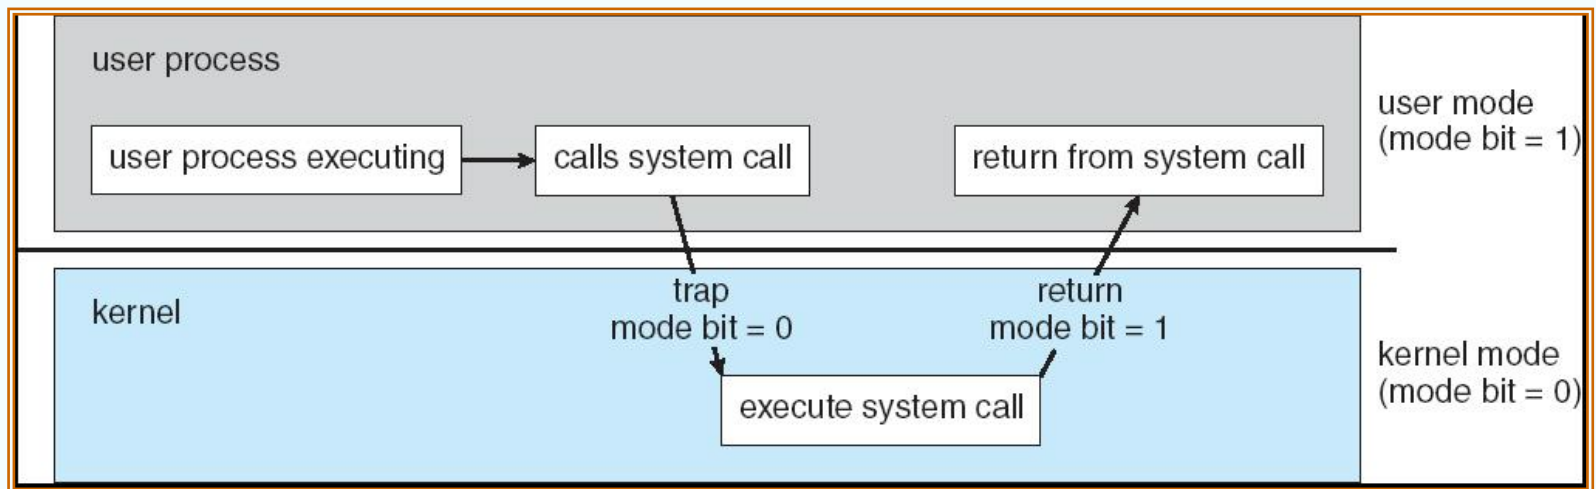
- Interrupt driven by hardware
- Software error or request creates **exception** or **trap**
 - Division by zero, request for operating system service
- Other process problems include **infinite loop**, processes **modifying** each other or the operating system

Thus we need protection:

- **Dual-mode** operation allows OS to protect itself and other system components
 - **User mode** and **kernel mode**
 - **Mode bit** provided by hardware
 - ▶ Provides ability to distinguish when system is running user code or kernel code
 - ▶ Some instructions designated as **privileged**, only executable in kernel mode
 - ▶ System call changes mode to kernel, return from call resets it to user

Transition from User to Kernel Mode

- Timer to prevent infinite loop / process hogging resources
 - Set interrupt after specific period
 - Operating system decrements counter
 - When counter zero generate an interrupt
 - Set up before scheduling process to regain control or terminate program that exceeds allotted time



Process Management

- A process is a program in execution. It is a unit of work within the system. Program is a *passive entity*, process is an *active entity*.
- Process needs resources to accomplish its task
 - CPU, memory, I/O, files
 - Initialization data
- Process termination requires reclaim of any reusable resources
- Single-threaded process has one **program counter** specifying location of next instruction to execute
 - Process executes instructions sequentially, one at a time, until completion
- Multi-threaded process has one **program counter per thread**
- Typically system has many processes, some user, some operating system running concurrently on **one or more CPUs**
 - Concurrency by **multiplexing the CPUs** among the processes / threads

Process Management Activities

The operating system is responsible for the following activities in connection with process management:

- **Creating** and **deleting** both user and system processes
- **Suspending** and **resuming** processes
- Providing mechanisms for process **synchronization**
- Providing mechanisms for process **communication**
- Providing mechanisms for **deadlock** handling

Memory Management

- All **data** must be in memory before and after processing
- All **instructions** must be in memory in order to execute
- Memory management determines what is in memory when
 - Optimizing CPU utilization and computer response to users
- Memory management activities
 - Keeping track of which parts of memory are currently being used and by whom
 - Deciding which processes (or parts thereof) and data to move into and out of memory
 - Allocating and deallocating memory space as needed

Storage Management

- OS provides uniform, logical view of information storage
 - Abstracts physical properties to logical storage unit - **file**
 - Each medium is controlled by device (i.e., disk drive, tape drive)
 - ▶ Varying properties include access speed, capacity, data-transfer rate, access method (sequential or random)
- File-System management
 - Files usually organized into **directories**
 - **Access control** on most systems to determine who can access what
 - OS activities include
 - ▶ Creating and deleting files and directories
 - ▶ Primitives to manipulate files and dirs
 - ▶ Mapping files onto secondary storage
 - ▶ Backup files onto stable (non-volatile) storage media

Mass-Storage Management

- Usually disks used to store data that does not fit in main memory or data that must be kept for a “long” period of time.
- Proper management is of central importance
- Entire **speed** of computer operation hinges on disk subsystem and its algorithms
- OS activities
 - Free-space management
 - Storage allocation
 - Disk scheduling
- Some storage needs not be fast
 - Tertiary storage includes optical storage, magnetic tape
 - Still must be managed
 - Varies between WORM (write-once, read-many-times) and RW (read-write)

I/O Subsystem

- One purpose of OS is to **hide peculiarities** of hardware devices from the user – *ease of usage & programming*
- I/O subsystem responsible for
 - Memory management of I/O including **buffering** (storing data temporarily while it is being transferred), **caching** (storing parts of data in faster storage for performance), **spooling** (the overlapping of output of one job with input of other jobs)
 - General device-driver interface
 - Drivers for specific hardware devices

OS Purposes

- Abstraction
- Multiplex
- Isolation
- Sharing
- Security
- Performance
- Range of uses

End of Chapter 1

RISC-V Registers

- Register file as follows:

- **RV32I/64I have 32 Integer Registers**
 - Optional 32 FP registers with the F and D extensions
 - RV32E reduces the register file to 16 integer registers for area constrained embedded devices
- **Width of Registers is determined by ISA**
- **RISC-V Application Binary Interface (ABI) defines standard functions for registers**
 - Allows for software interoperability
- **Development tools usually use ABI names for simplicity**

Register	ABI Name	Description	Saver
x0	zero	Hard-wired zero	-
x1	ra	Return address	Caller
x2	sp	Stack pointer	Callee
x3	gp	Global pointer	-
x4	tp	Thread pointer	-
x5-7	t0-2	Temporaries	Caller
x8	s0/fp	Saved register/Frame pointer	Callee
x9	s1	Saved register	Callee
x10-11	a0-1	Function Arguments/return values	Caller
x12-17	a2-7	Function arguments	Caller
x18-27	s2-11	Saved registers	Callee
x28-31	t3-6	Temporaries	Caller

- In addition, RISC-V has a set of Control & Status Registers (CSRs)